

Inclusus Multi-Tenancy

A Beginner's Guide

Ananda Kammampati

“The Greatest Teacher : One who makes you wanting to learn “

Dedicated to

My classmates Kannan Ramasubramanian and Rajesh Kunnath who are my favorite teachers and my source of inspiration.

Table of Contents

<i>Tasks</i>	<i>Description</i>	<i>page no.</i>
Task 00	Introduction	01
Task 01	Hardware specifications of incus server	02
Task 02	Pre-requisites	03
Task 03	Partitioning of two NVMe SSDs	08
Task 04	Creation of ZFS pools	17
Task 05	Installation and configuration of incus	23
Task 06	incus storage	28
Task 07	Installation of OVN	32
Task 08	Installation of incus-ui-canonical	33
Task 09	incus remote access setup	35
Task 10	incus web UI	37
Task 11	Sanity test with a container	41
Task 12	Creating a volume	46
Task 13	Sanity testing the volume	49
Task 14	Uploading an ISO file	51
Task 15	ISO file usage	56
Task 16	Using custom storage pools and volumes	65
Task 17	Creating Projects using web UI	87
Task 18	Verifying storage for project: Handson-lab-1	92
Task 19	default profile for project: Handson-lab-1	94
Task 20	Creating external network for project: Handson-lab-1	97
Task 21	Sanity testing of external network	106
Task 22	Creating internal network for project: Handson-lab-1	113
Task 23	Sanity testing of internal network	116
Task 24	Creating a profile for a Desktop virtual machine	120
Task 25	Creation of Desktop virtual machine	122
Task 26	Creation of project: Handson-lab-2	130
Task 27	Network tests between projects - Take #1	136
Task 28	Setting network isolation among Projects	138
Task 29	Network tests between projects - Take #2	145
Task 30	A fully built multi-tenant system with 6 tenants	148
Task 31	Credits and Thanks to my incus Gurus	155
Task 32	About the Author	156

Incus Multi-Tenancy: A Beginner's Guide

Introduction

The primary audience of this book are users who have some working knowledge on incus virtualization platform and would like to understand better on how to configure incus into a multi-tenant platform. The goal is to guide the reader in leveraging incus capabilities and features namely projects, open virtual networking, access control lists and many more to gain an overall understanding in building a minimal multi-tenant system on a single server.

This book explores only a bare minimal design needed for multi-tenancy with regards to storage and networking to gain some insights. This book is not meant to teach the basics of incus virtualization platform. Please refer to the credits section at the end of this book that provide links to the best teachers on incus.

For a complete beginner to incus virtualization platform, I would recommend to start with Trevor Sullivan's YouTube playlist on LXD containers.

The approach taken in this book is very Hands-on. It explains incus multi-tenancy by breaking it down not as chapters but a list of tasks, each task dependent on the completion of its previous task.

The recommended way to learn from this book is to use it as a Cookbook. By having a dedicated server and trying out all the tasks in the same order.

Happy learning :-)

- Ananda Kammampati

Task #01: Hardware Specifications of Incus Server

The mother board is from Supermicro (X10SDV-6C-TLN4F)

CPU	Intel XEON processor D-1528, 6-Core, 12 Threads
Memory	128GB ECC RDIMM DDR4 2133MHz
NICs	2 x 10GbE ports and 2 x 1GbE ports

Refer :

<https://www.supermicro.com/en/products/motherboard/x10sdv-tln4f>

Task #02: Pre-requisites

Here are the list of things that we are going to accomplish in this task.

1. Checking system details
2. Verifying bootable USB disk partition
3. Installation of ZFS packages
4. Removal of snap package manager
5. Adding two NVMe SSDs
6. Installation of docker (optional)

1. Checking system details

Making a bootable USB thumb drive is beyond the scope of this cookbook (there are tons of awesome online tutorials). As a pre-requisite, we have already made a 512GB USB thumb drive as bootable disk, booted the system and completed the installation.

```
# uname -a
Linux red-handsonbox 6.8.0-55-generic #57-Ubuntu SMP PREEMPT_DYNAMIC Wed Feb 12 23:42:21
UTC 2025 x86_64 x86_64 x86_64 GNU/Linux
```

```
# cat /etc/os-release
PRETTY_NAME="Ubuntu 24.04.2 LTS"
NAME="Ubuntu"
VERSION_ID="24.04"
VERSION="24.04.2 LTS (Noble Numbat)"
VERSION_CODENAME=noble
ID=ubuntu
ID_LIKE=debian
HOME_URL="https://www.ubuntu.com/"
SUPPORT_URL="https://help.ubuntu.com/"
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"
UBUNTU_CODENAME=noble
```

2. Verifying bootable USB disk partition

Let us make sure that system boots fine from the USB thumb drive and the partitions are intact

```
# lsblk
NAME      MAJ:MIN RM  SIZE    RO TYPE  MOUNTPOINTS
sda          8:0    1   460.3G  0  disk
```

```

└─sda1      8:1    1      1M      0  part
└─sda2      8:2    1    460.3G  0  part   /
# df -h
Filesystem      Size  Used Avail Use% Mounted on
tmpfs           13G  1.6M  13G  1% /run
/dev/sda2     452G  11G  419G  3% /
tmpfs           63G    0  63G  0% /dev/shm
tmpfs          5.0M    0  5.0M  0% /run/lock
tmpfs           13G  8.0K  13G  1% /run/user/0

```

3. Installing ZFS

Compared to other Linux distributions, installation of ZFS seems pretty straight forward on Ubuntu. The commands below are self-explanatory

```

# pwd
/root

# apt update -y

# apt upgrade -y

# apt install zfsutils-linux -y

# modprobe zfs

# zfs --version
zfs-2.2.2-0ubuntu9.1
zfs-kmod-2.2.2-0ubuntu9.1

# zpool --version
zfs-2.2.2-0ubuntu9.1
zfs-kmod-2.2.2-0ubuntu9.1

# modinfo zfs
filename:      /lib/modules/6.8.0-55-generic/kernel/zfs/zfs.ko.zst
version:       2.2.2-0ubuntu9.1
license:        CDDL
license:       Dual BSD/GPL
license:       Dual MIT/GPL
author:        OpenZFS
description:   ZFS

```

```

alias:      zzstd
alias:      zcommon
alias:      zunicode
alias:      znpair
alias:      zlua
alias:      icp
alias:      zavl
alias:      devname:zfs
alias:      char-major-10-249
srcversion: 307E3A0219DD7E837933412
depends:    spl
retpoline:  Y
name:       zfs
vermagic:   6.8.0-55-generic SMP preempt mod_unload modversions
sig_id:     PKCS#7
signer:     Build time autogenerated kernel key
.....
.....
.....
```

4. Removing snap package manager

I still haven't taken the time to learn. So purging it for now ;-)

```

# apt purge snapd
# apt-mark hold snapd
# rm -rf /var/cache/snapd/
# rm -rf ~/snap
# sync ; reboot
```

5. Adding two NVMe SSD's

The reason for opting 2 NVMe SSDs and the reason behind creating different partitions of equal size needs some detailed explanation, which we will look into in the following tasks. Command ‘lsblk’ shows the details of 3 storage devices. ‘sda’ is the bootable USB thumb drive. ‘nvme0n1’ and ‘nvme1n1’ are the two NVMe SSD disks of size 2GB and 1GB respectively..

```
# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    1  460.3G  0 disk
└─sda1     8:1    1     1M  0 part
└─sda2     8:2    1  460.3G  0 part /
nvme1n1   259:0   0  931.5G  0 disk
└─nvme1n1p1 259:1   0  155G  0 part
└─nvme1n1p2 259:2   0  155G  0 part
└─nvme1n1p3 259:3   0  155G  0 part
└─nvme1n1p4 259:4   0  155G  0 part
└─nvme1n1p5 259:5   0  155G  0 part
└─nvme1n1p6 259:6   0  155G  0 part
└─nvme1n1p7 259:7   0   1.5G  0 part
nvme0n1   259:8   0   1.9T  0 disk
└─nvme0n1p1 259:9   0   317G  0 part
└─nvme0n1p2 259:10  0   317G  0 part
└─nvme0n1p3 259:11  0   317G  0 part
└─nvme0n1p4 259:12  0   317G  0 part
└─nvme0n1p5 259:13  0   317G  0 part
└─nvme0n1p6 259:14  0   317G  0 part
└─nvme0n1p7 259:15  0   5.7G  0 part
```

Installation of docker

Installation of docker is purely optional at this stage. But since we are with the installation task, we might as well take care of it. The bash shell script shown below has all the required commands to fully install docker engine and start it as a service.

```
# cat install-docker.sh
-----
#!/bin/bash
clear
echo " "
echo "update...."
echo " "
apt -y update ;
echo " "
echo "upgrade...."
echo " "
apt -y upgrade ;
echo " "
echo "install apt-transport-https ca-certificates curl software-properties-common..."
```

```

echo " "
apt install apt-transport-https ca-certificates curl software-properties-common -y ;
echo " "
echo "install key..."
echo " "
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | apt-key add - ;
echo " "
echo "add repository..."
echo " "
add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
$(lsb_release -cs) stable" ;
echo " "
echo "install docker-ce..."
echo " "
apt update -y ; apt install docker-ce docker-ce-cli containerd.io docker-buildx-plugin
docker-compose-plugin -y ;
echo " "
echo "install docker-compose..."
echo " "
curl -L
https://github.com/docker/compose/releases/download/v2.32.4/docker-compose-$(uname
-s)-$(uname -m) -o /usr/local/bin/docker-compose ;
chmod +x /usr/local/bin/docker-compose ;
echo " "
docker-compose --version ;
echo " "
echo "start docker..."
echo " "
systemctl start docker ;
echo " "
echo "enable docker..."
echo " "
systemctl enable docker ;
echo " "
echo "check docker..."
echo " "
systemctl status -l docker ;
-----
# chmod 755 install-docker.sh

# ./install-docker.sh

```

Task #03: Partitioning of NVMe SSDs

Let us break down this task into following steps:

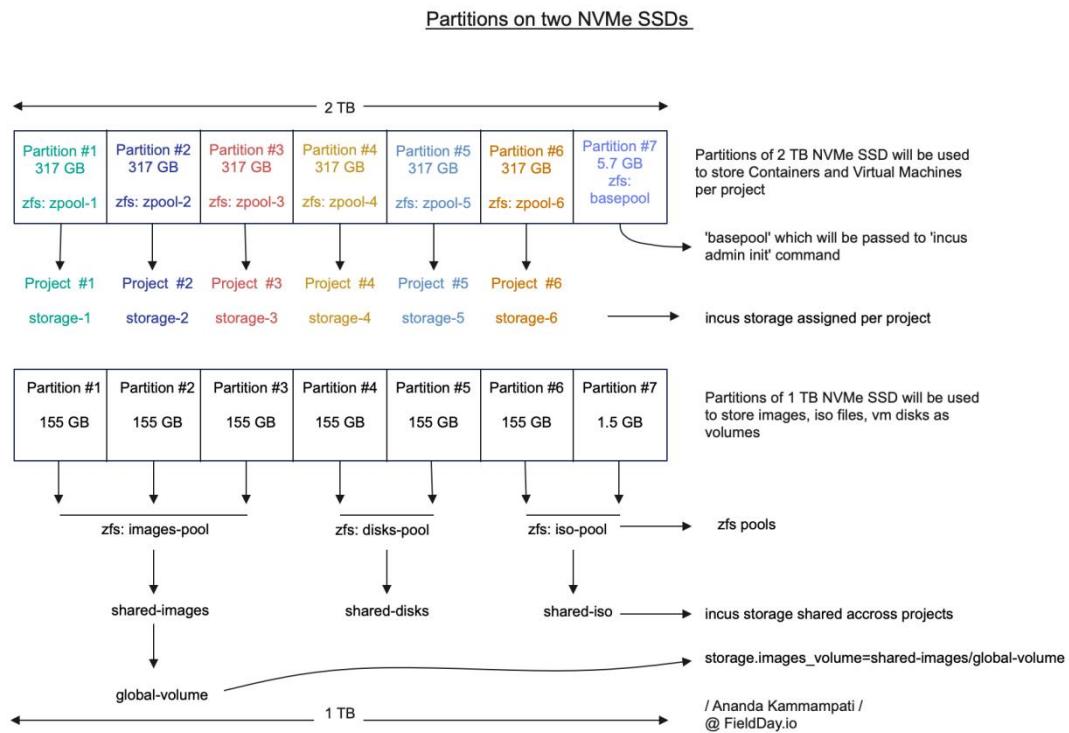
1. Understanding the anatomy of partitions
2. Partitioning of the 2TB SSD
3. Partitioning of the 1TB SSD

1. Understanding the anatomy of Partitions

In this task we will focus on creating multiple partitions of equal sizes on both 1TB and 2TB NVMe SSDs. In incus, Multi-tenancy is achieved by leveraging the construct/artifact PROJECT. The term ‘Project’ and ‘Tenant’ can be used synonymously as they both refer the same idea. In simple terms:

“ Project is a method of resource isolation to attain Multi-Tenancy “.

Our idea behind creating multiple partitions is to allocate each partitions to individual Projects. Later we will name the projects as ‘Handson-lab-1’, ‘Handson-lab-2’ and so on. Our plan is to have 6 Projects/Tenents on our system. So we will create 6 partitions on each SSDs. Later when we create individual Projects, we will allocate these individual partitions to each Project. For now we will focus on creation of Partitions.



The above diagram shows the anatomy of how we are planning to create 6 partitions on each

SSD. A single partition on the 2TB SSD will be allocated to a single Project. For example Partition #1 on the 2TB SSD will be allocated to Project #1. And this process will be repeated to all the subsequent Projects.

Note that the Partition #7 on the 2TB SSD is what we use while initializing incus with ‘incus admin init’ command. This will be clear later as we progress.

We will use the native tool ‘gdisk’ on Ubuntu24.04 which will be the primary Host OS running on our Supermicro server. To begin with let us focus on the 2TB NVMe SSD which is /dev/nvme0n1 on our system.

Explaining how to use tool gdisk is beyond the scope of this book (refer to online videos that explain better on using the tool).

We will create 6 equal sized partitions (317 GB each) on the 2TB SSD. With the remaining 5.7 GB space, we will create a separate partition. It is this Partition #7 we name it as ‘basepool’ and will use it to configure incus with the ‘incus admin init’ command.

2. Partitioning of the 2TB SSD

Let us begin configuring the 2TB SSD with gdisk tool.

```
# gdisk /dev/nvme0n1
GPT fdisk (gdisk) version 1.0.9
```

Partition table scan:

```
MBR: not present
BSD: not present
APM: not present
GPT: not present
```

Creating new GPT entries in memory.

```
Command (? for help): n
Partition number (1-128, default 1):
First sector (34-4000797326, default = 2048) or {+-}size{KMGTP}:
Last sector (2048-4000797326, default = 4000796671) or {+-}size{KMGTP}: +317G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n
Partition number (2-128, default 2):
First sector (34-4000797326, default = 664799232) or {+-}size{KMGTP}:
```

```
Last sector (664799232-4000797326, default = 4000796671) or {+-}size{KMGTP}: +317G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n
Partition number (3-128, default 3):
First sector (34-4000797326, default = 1329596416) or {+-}size{KMGTP}:
Last sector (1329596416-4000797326, default = 4000796671) or {+-}size{KMGTP}: +317G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n
Partition number (4-128, default 4):
First sector (34-4000797326, default = 1994393600) or {+-}size{KMGTP}:
Last sector (1994393600-4000797326, default = 4000796671) or {+-}size{KMGTP}: +317G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n
Partition number (5-128, default 5):
First sector (34-4000797326, default = 2659190784) or {+-}size{KMGTP}:
Last sector (2659190784-4000797326, default = 4000796671) or {+-}size{KMGTP}: +317G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n
Partition number (6-128, default 6):
First sector (34-4000797326, default = 3323987968) or {+-}size{KMGTP}:
Last sector (3323987968-4000797326, default = 4000796671) or {+-}size{KMGTP}: +317G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): p
Disk /dev/nvme0n1: 4000797360 sectors, 1.9 TiB
Model: TEAM TM8FP6002T
Sector size (logical/physical): 512/512 bytes
Disk identifier (GUID): A7E37AAC-57CC-4D4F-B995-0919B01794EA
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
```

```
First usable sector is 34, last usable sector is 4000797326
Partitions will be aligned on 2048-sector boundaries
Total free space is 12014189 sectors (5.7 GiB)
```

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	664799231	317.0 GiB	8300	Linux filesystem
2	664799232	1329596415	317.0 GiB	8300	Linux filesystem
3	1329596416	1994393599	317.0 GiB	8300	Linux filesystem
4	1994393600	2659190783	317.0 GiB	8300	Linux filesystem
5	2659190784	3323987967	317.0 GiB	8300	Linux filesystem
6	3323987968	3988785151	317.0 GiB	8300	Linux filesystem

```
Command (? for help): n
```

```
Partition number (7-128, default 7):
```

```
First sector (34-4000797326, default = 3988785152) or {+-}size{KMGT}P:
```

```
Last sector (3988785152-4000797326, default = 4000796671) or {+-}size{KMGT}P:
```

```
Current type is 8300 (Linux filesystem)
```

```
Hex code or GUID (L to show codes, Enter = 8300):
```

```
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): p
```

```
Disk /dev/nvme0n1: 4000797360 sectors, 1.9 TiB
```

```
Model: TEAM TM8FP6002T
```

```
Sector size (logical/physical): 512/512 bytes
```

```
Disk identifier (GUID): A7E37AAC-57CC-4D4F-B995-0919B01794EA
```

```
Partition table holds up to 128 entries
```

```
Main partition table begins at sector 2 and ends at sector 33
```

```
First usable sector is 34, last usable sector is 4000797326
```

```
Partitions will be aligned on 2048-sector boundaries
```

```
Total free space is 2669 sectors (1.3 MiB)
```

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	664799231	317.0 GiB	8300	Linux filesystem
2	664799232	1329596415	317.0 GiB	8300	Linux filesystem
3	1329596416	1994393599	317.0 GiB	8300	Linux filesystem
4	1994393600	2659190783	317.0 GiB	8300	Linux filesystem
5	2659190784	3323987967	317.0 GiB	8300	Linux filesystem
6	3323987968	3988785151	317.0 GiB	8300	Linux filesystem
7	3988785152	4000796671	5.7 GiB	8300	Linux filesystem

```
Command (? for help): w
```

```
Final checks complete. About to write GPT data. THIS WILL OVERWRITE EXISTING
PARTITIONS!!
```

```
Do you want to proceed? (Y/N): Y
OK; writing new GUID partition table (GPT) to /dev/nvme0n1.
The operation has completed successfully.
```

```
# gdisk -l /dev/nvme0n1
GPT fdisk (gdisk) version 1.0.9
```

Partition table scan:

```
  MBR: protective
  BSD: not present
  APM: not present
  GPT: present
```

```
Found valid GPT with protective MBR; using GPT.
Disk /dev/nvme0n1: 4000797360 sectors, 1.9 TiB
Model: TEAM TM8FP6002T
Sector size (logical/physical): 512/512 bytes
Disk identifier (GUID): A7E37AAC-57CC-4D4F-B995-0919B01794EA
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
First usable sector is 34, last usable sector is 4000797326
Partitions will be aligned on 2048-sector boundaries
Total free space is 2669 sectors (1.3 MiB)
```

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	664799231	317.0 GiB	8300	Linux filesystem
2	664799232	1329596415	317.0 GiB	8300	Linux filesystem
3	1329596416	1994393599	317.0 GiB	8300	Linux filesystem
4	1994393600	2659190783	317.0 GiB	8300	Linux filesystem
5	2659190784	3323987967	317.0 GiB	8300	Linux filesystem
6	3323987968	3988785151	317.0 GiB	8300	Linux filesystem
7	3988785152	4000796671	5.7 GiB	8300	Linux filesystem

```
# fdisk -l
Disk /dev/nvme0n1: 1.86 TiB, 2048408248320 bytes, 4000797360 sectors
Disk model: TEAM TM8FP6002T
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: A7E37AAC-57CC-4D4F-B995-0919B01794EA
```

Device	Start	End	Sectors	Size	Type
--------	-------	-----	---------	------	------

```

/dev/nvme0n1p1      2048 664799231 664797184 317G Linux filesystem
/dev/nvme0n1p2  664799232 1329596415 664797184 317G Linux filesystem
/dev/nvme0n1p3 1329596416 1994393599 664797184 317G Linux filesystem
/dev/nvme0n1p4 1994393600 2659190783 664797184 317G Linux filesystem
/dev/nvme0n1p5 2659190784 3323987967 664797184 317G Linux filesystem
/dev/nvme0n1p6 3323987968 3988785151 664797184 317G Linux filesystem
/dev/nvme0n1p7 3988785152 4000796671 12011520 5.7G Linux filesystem

```

Now that we have completed creating 7 partitions on 2TB SSD, let us move onto the 1TB SSD and repeat the same procedure. The size of each partition will be 155 GB. ie, 1 TB SSD divided into 6 equal partition size. We will create a separate partition with the remaining 1.5 GB space which is Partition #7.

3. Partitioning of the 1TB SSD

```

# gdisk /dev/nvme1n1
GPT fdisk (gdisk) version 1.0.9

```

Partition table scan:

```

MBR: protective
BSD: not present
APM: not present
GPT: present

```

Found valid GPT with protective MBR; using GPT.

```

Command (? for help): n
Partition number (1-128, default 1):
First sector (34-1953525134, default = 2048) or {+-}size{KMGTP}:
Last sector (2048-1953525134, default = 1953523711) or {+-}size{KMGTP}: +155G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'

```

```

Command (? for help): n
Partition number (2-128, default 2):
First sector (34-1953525134, default = 325060608) or {+-}size{KMGTP}:
Last sector (325060608-1953525134, default = 1953523711) or {+-}size{KMGTP}: +155G
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'

```

```

Command (? for help): n
Partition number (3-128, default 3):

```

```
First sector (34-1953525134, default = 650119168) or {+-}size{KMGTP}:  
Last sector (650119168-1953525134, default = 1953523711) or {+-}size{KMGTP}: +155G  
Current type is 8300 (Linux filesystem)  
Hex code or GUID (L to show codes, Enter = 8300):  
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n  
Partition number (4-128, default 4):  
First sector (34-1953525134, default = 975177728) or {+-}size{KMGTP}:  
Last sector (975177728-1953525134, default = 1953523711) or {+-}size{KMGTP}: +155G  
Current type is 8300 (Linux filesystem)  
Hex code or GUID (L to show codes, Enter = 8300):  
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n  
Partition number (5-128, default 5):  
First sector (34-1953525134, default = 1300236288) or {+-}size{KMGTP}:  
Last sector (1300236288-1953525134, default = 1953523711) or {+-}size{KMGTP}: +155G  
Current type is 8300 (Linux filesystem)  
Hex code or GUID (L to show codes, Enter = 8300):  
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n  
Partition number (6-128, default 6):  
First sector (34-1953525134, default = 1625294848) or {+-}size{KMGTP}:  
Last sector (1625294848-1953525134, default = 1953523711) or {+-}size{KMGTP}: +155G  
Current type is 8300 (Linux filesystem)  
Hex code or GUID (L to show codes, Enter = 8300):  
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): p  
Disk /dev/nvme1n1: 1953525168 sectors, 931.5 GiB  
Model: WDC WDS100T2B0C-00PXH0  
Sector size (logical/physical): 512/512 bytes  
Disk identifier (GUID): D5557850-4398-F748-8720-008C4186AEAC  
Partition table holds up to 128 entries  
Main partition table begins at sector 2 and ends at sector 33  
First usable sector is 34, last usable sector is 1953525134  
Partitions will be aligned on 2048-sector boundaries  
Total free space is 3173741 sectors (1.5 GiB)
```

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	325060607	155.0 GiB	8300	Linux filesystem
2	325060608	650119167	155.0 GiB	8300	Linux filesystem

3	650119168	975177727	155.0 GiB	8300	Linux filesystem
4	975177728	1300236287	155.0 GiB	8300	Linux filesystem
5	1300236288	1625294847	155.0 GiB	8300	Linux filesystem
6	1625294848	1950353407	155.0 GiB	8300	Linux filesystem

Command (? for help): n

Partition number (7-128, default 7):

First sector (34-1953525134, default = 1950353408) or {+-}size{KMGTP}:

Last sector (1950353408-1953525134, default = 1953523711) or {+-}size{KMGTP}:

Current type is 8300 (Linux filesystem)

Hex code or GUID (L to show codes, Enter = 8300):

Changed type of partition to 'Linux filesystem'

Command (? for help): p

Disk /dev/nvme1n1: 1953525168 sectors, 931.5 GiB

Model: WDC WDS100T2B0C-00PXH0

Sector size (logical/physical): 512/512 bytes

Disk identifier (GUID): D5557850-4398-F748-8720-008C4186AEAC

Partition table holds up to 128 entries

Main partition table begins at sector 2 and ends at sector 33

First usable sector is 34, last usable sector is 1953525134

Partitions will be aligned on 2048-sector boundaries

Total free space is 3437 sectors (1.7 MiB)

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	325060607	155.0 GiB	8300	Linux filesystem
2	325060608	650119167	155.0 GiB	8300	Linux filesystem
3	650119168	975177727	155.0 GiB	8300	Linux filesystem
4	975177728	1300236287	155.0 GiB	8300	Linux filesystem
5	1300236288	1625294847	155.0 GiB	8300	Linux filesystem
6	1625294848	1950353407	155.0 GiB	8300	Linux filesystem
7	1950353408	1953523711	1.5 GiB	8300	Linux filesystem

Command (? for help): w

Final checks complete. About to write GPT data. THIS WILL OVERWRITE EXISTING PARTITIONS!!

Do you want to proceed? (Y/N): Y

OK; writing new GUID partition table (GPT) to /dev/nvme1n1.

The operation has completed successfully.

root@red-handsonbox:~#

root@red-handsonbox:~# gdisk -l /dev/nvme1n1

GPT fdisk (gdisk) version 1.0.9

Partition table scan:

MBR: protective
BSD: not present
APM: not present
GPT: present

Found valid GPT with protective MBR; using GPT.
Disk /dev/nvme1n1: 1953525168 sectors, 931.5 GiB
Model: WDC WDS100T2B0C-00PXH0
Sector size (logical/physical): 512/512 bytes
Disk identifier (GUID): D5557850-4398-F748-8720-008C4186AEAC
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
First usable sector is 34, last usable sector is 1953525134
Partitions will be aligned on 2048-sector boundaries
Total free space is 3437 sectors (1.7 MiB)

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	325060607	155.0 GiB	8300	Linux filesystem
2	325060608	650119167	155.0 GiB	8300	Linux filesystem
3	650119168	975177727	155.0 GiB	8300	Linux filesystem
4	975177728	1300236287	155.0 GiB	8300	Linux filesystem
5	1300236288	1625294847	155.0 GiB	8300	Linux filesystem
6	1625294848	1950353407	155.0 GiB	8300	Linux filesystem
7	1950353408	1953523711	1.5 GiB	8300	Linux filesystem

fdisk -l /dev/nvme1n1
Disk /dev/nvme1n1: 931.51 GiB, 1000204886016 bytes, 1953525168 sectors
Disk model: WDC WDS100T2B0C-00PXH0
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: D5557850-4398-F748-8720-008C4186AEAC

Device	Start	End	Sectors	Size	Type
/dev/nvme1n1p1	2048	325060607	325058560	155G	Linux filesystem
/dev/nvme1n1p2	325060608	650119167	325058560	155G	Linux filesystem
/dev/nvme1n1p3	650119168	975177727	325058560	155G	Linux filesystem
/dev/nvme1n1p4	975177728	1300236287	325058560	155G	Linux filesystem
/dev/nvme1n1p5	1300236288	1625294847	325058560	155G	Linux filesystem
/dev/nvme1n1p6	1625294848	1950353407	325058560	155G	Linux filesystem
/dev/nvme1n1p7	1950353408	1953523711	3170304	1.5G	Linux filesystem

Task #04: Creation of ZFS pools

In this task we will focus on creating ZFS pools with the partitions that we have already created with the previous task. Let us break down the procedures into following steps.

1. Verification of zfs
2. Creation of zfs pool per partition on 2TB SSD
3. Creation of zfs pool per partition on 1TB SSD

Please note that the technical know-how on how ZFS filesystem work is beyond the scope of this book. We will primarily focus on using it. ie, creating multiple ZFS storage pools so that they each can be assigned to different incus projects

1. Verification of zfs

The commands we use to manage zfs filesystem are ‘zpool’ and ‘zfs’. But first we need to verify that the zfs kernel module is indeed loaded and this step is very crucial. Without this kernel module loaded, zfs will not work, even if zfs related commands are installed. Luckily, the latest versions of Ubuntu (24.04 as of this writing) come with ZFS support by default which makes life much easier.

If the output of ‘lsmod | grep zfs’ doesn’t list anything, then we need to resolve loading zfs kernel module first before proceeding further. If one is not planning on using ZFS, it is still possible to build a Multi-tenancy incus system with other filesystems (example ext4, btrfs). But in our case, we will primarily focus on using ZFS.

```
# lsmod | grep zfs
zfs 6602752 6
spl 180224 1 zfs

# zfs --version
zfs-2.2.2-0ubuntu9.2
zfs-kmod-2.2.2-0ubuntu9.1

# zpool --version
zfs-2.2.2-0ubuntu9.2
zfs-kmod-2.2.2-0ubuntu9.1
```

2. Creation of zfs pool per partition on 2TB SSD

The command ‘lsblk’ lists all the partitions that we have created so far on both the SSDs.

Let us take a look at all the partitions of all the disks, including the bootable USD disk

```
# fdisk -l

Disk /dev/nvme0n1: 1.86 TiB, 2048408248320 bytes, 4000797360 sectors
Disk model: TEAM TM8FP6002T
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: A7E37AAC-57CC-4D4F-B995-0919B01794EA
```

Device	Start	End	Sectors	Size	Type
/dev/nvme0n1p1	2048	664799231	664797184	317G	Linux filesystem
/dev/nvme0n1p2	664799232	1329596415	664797184	317G	Linux filesystem
/dev/nvme0n1p3	1329596416	1994393599	664797184	317G	Linux filesystem
/dev/nvme0n1p4	1994393600	2659190783	664797184	317G	Linux filesystem
/dev/nvme0n1p5	2659190784	3323987967	664797184	317G	Linux filesystem
/dev/nvme0n1p6	3323987968	3988785151	664797184	317G	Linux filesystem
/dev/nvme0n1p7	3988785152	4000796671	12011520	5.7G	Linux filesystem

```
Disk /dev/nvme1n1: 931.51 GiB, 1000204886016 bytes, 1953525168 sectors
Disk model: WDC WDS100T2B0C-00PXH0
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: D5557850-4398-F748-8720-008C4186AEAC
```

Device	Start	End	Sectors	Size	Type
/dev/nvme1n1p1	2048	325060607	325058560	155G	Linux filesystem
/dev/nvme1n1p2	325060608	650119167	325058560	155G	Linux filesystem
/dev/nvme1n1p3	650119168	975177727	325058560	155G	Linux filesystem
/dev/nvme1n1p4	975177728	1300236287	325058560	155G	Linux filesystem
/dev/nvme1n1p5	1300236288	1625294847	325058560	155G	Linux filesystem
/dev/nvme1n1p6	1625294848	1950353407	325058560	155G	Linux filesystem
/dev/nvme1n1p7	1950353408	1953523711	3170304	1.5G	Linux filesystem

```
Disk /dev/sda: 460.27 GiB, 494206451712 bytes, 965246976 sectors
Disk model: SanDisk 3.2Gen1
```

```

Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: 1DDC2624-2581-4560-AE31-44AAC5E3A796

```

Device	Start	End	Sectors	Size	Type
/dev/sda1	2048	4095	2048	1M	BIOS boot
/dev/sda2	4096	965244927	965240832	460.3G	Linux filesystem

Let us understand how we are planning on creating multiple zfs pools on each partitions. We will start with the 2 TB SSD first

```
# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    1 460.3G 0 disk
└─sda1     8:1    1   1M 0 part
└─sda2     8:2    1 460.3G 0 part / ---> Root on USB thumb drive
nvme0n1   259:0   0  1.9T 0 disk  ----> 2 TB NVMe SSD
└─nvme0n1p1 259:1  0 317G 0 part
└─nvme0n1p2 259:2  0 317G 0 part
└─nvme0n1p3 259:3  0 317G 0 part
└─nvme0n1p4 259:4  0 317G 0 part
└─nvme0n1p5 259:5  0 317G 0 part
└─nvme0n1p6 259:6  0 317G 0 part
└─nvme0n1p7 259:7  0  5.7G 0 part
nvme1n1   259:8   0 931.5G 0 disk  ---> 1 TB NVMe SSD
└─nvme1n1p1 259:9  0 155G 0 part
└─nvme1n1p2 259:10 0 155G 0 part
└─nvme1n1p3 259:11 0 155G 0 part
└─nvme1n1p4 259:12 0 155G 0 part
└─nvme1n1p5 259:13 0 155G 0 part
└─nvme1n1p6 259:14 0 155G 0 part
└─nvme1n1p7 259:15 0  1.5G 0 part
```

```
# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    1 460.3G 0 disk
└─sda1     8:1    1   1M 0 part
└─sda2     8:2    1 460.3G 0 part / --> Root partition on USB thumb drive
nvme0n1   259:0   0  1.9T 0 disk  --> 2 TB NVMe SSD
└─nvme0n1p1 259:1  0 317G 0 part  --> Create zfs zpool-1
└─nvme0n1p2 259:2  0 317G 0 part  --> Create zfs zpool-2
└─nvme0n1p3 259:3  0 317G 0 part  --> Create zfs zpool-3
```

```

└─nvme0n1p4 259:4    0  317G  0 part  --> Create zfs zpool-4
└─nvme0n1p5 259:5    0  317G  0 part  --> Create zfs zpool-5
└─nvme0n1p6 259:6    0  317G  0 part  --> Create zfs zpool-6
└─nvme0n1p7 259:7    0   5.7G  0 part  --> Create zfs basepool
nvme1n1    259:8    0 931.5G  0 disk  --> 1 TB NVMe SSD
└─nvme1n1p1 259:9    0  155G  0 part
└─nvme1n1p2 259:10   0  155G  0 part
└─nvme1n1p3 259:11   0  155G  0 part
└─nvme1n1p4 259:12   0  155G  0 part
└─nvme1n1p5 259:13   0  155G  0 part
└─nvme1n1p6 259:14   0  155G  0 part
└─nvme1n1p7 259:15   0   1.5G  0 part

```

Once again, we will take a look at the output of ‘lsblk’ command and see how we are planning on creating multiple storage pools and how each are assigned to/consumed by individual projects.

```

# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0     1 460.3G 0 disk
└─sda1     8:1     1   1M 0 part
└─sda2     8:2     1 460.3G 0 part /
nvme0n1   259:0   0  1.9T 0 disk
└─nvme0n1p1 259:1  0  317G 0 part  --> Assign to Project-1
└─nvme0n1p2 259:2  0  317G 0 part  --> Assign to Project-2
└─nvme0n1p3 259:3  0  317G 0 part  --> Assign to Project-3
└─nvme0n1p4 259:4  0  317G 0 part  --> Assign to Project-4
└─nvme0n1p5 259:5  0  317G 0 part  --> Assign to Project-5
└─nvme0n1p6 259:6  0  317G 0 part  --> Assign to Project-6
└─nvme0n1p7 259:7  0   5.7G 0 part  --> Pass to ‘incus admin init’
nvme1n1   259:8   0 931.5G 0 disk
└─nvme1n1p1 259:9  0  155G 0 part
└─nvme1n1p2 259:10 0  155G 0 part
└─nvme1n1p3 259:11 0  155G 0 part
└─nvme1n1p4 259:12 0  155G 0 part
└─nvme1n1p5 259:13 0  155G 0 part
└─nvme1n1p6 259:14 0  155G 0 part
└─nvme1n1p7 259:15 0   1.5G 0 part

```

Now that we have a better understanding on what we plan on doing with separate partitions, let us go ahead and run the zfs commands to create separate zfs pools on each partition.

```

# zpool create zpool-1 /dev/nvme0n1p1
# zpool create zpool-2 /dev/nvme0n1p2

```

```
# zpool create zpool-3 /dev/nvme0n1p3
# zpool create zpool-4 /dev/nvme0n1p4
# zpool create zpool-5 /dev/nvme0n1p5
# zpool create zpool-6 /dev/nvme0n1p6
# zpool create basepool /dev/nvme0n1p7
```

As simple as that.

# zpool list											
NAME	SIZE	ALLOC	FREE	CKPOINT	EXPANDSZ	FRAG	CAP	DEDUP	HEALTH	ALTROOT	
basepool	5.50G	132K	5.50G	-	-	0%	0%	1.00x	ONLINE	-	
zpool-1	316G	110K	316G	-	-	0%	0%	1.00x	ONLINE	-	
zpool-2	316G	108K	316G	-	-	0%	0%	1.00x	ONLINE	-	
zpool-3	316G	110K	316G	-	-	0%	0%	1.00x	ONLINE	-	
zpool-4	316G	110K	316G	-	-	0%	0%	1.00x	ONLINE	-	
zpool-5	316G	110K	316G	-	-	0%	0%	1.00x	ONLINE	-	
zpool-6	316G	110K	316G	-	-	0%	0%	1.00x	ONLINE	-	

3. Creation of zfs pools on 1TB SSD

Let us repeat a similar exercise but with a slight twist with regards to the 1TB SSD

- Let us create a zfs pool comprising of partition 1, 2, 3 and name it as ‘images-pool’
- Let us create a zfs pool comprising of partitions 4 and 5 and name it as ‘disks-pool’
- Let us create a zfs pool comprising with partitions 6 and 7 and name it as ‘iso-pool’

For clarity, let us mark down the name of the zfs pools that we plan on creating on the existing partitions on the 1TB SSD

```
# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    1 460.3G 0 disk
└─sda1     8:1    1   1M 0 part
└─sda2     8:2    1 460.3G 0 part / --> Root partition on USB thumb drive
nvme0n1   259:0   0  1.9T 0 disk  --> 2 TB NVMe SSD
└─nvme0n1p1 259:1  0 317G 0 part
└─nvme0n1p2 259:2  0 317G 0 part
└─nvme0n1p3 259:3  0 317G 0 part
└─nvme0n1p4 259:4  0 317G 0 part
└─nvme0n1p5 259:5  0 317G 0 part
└─nvme0n1p6 259:6  0 317G 0 part
└─nvme0n1p7 259:7  0  5.7G 0 part
nvme1n1   259:8   0 931.5G 0 disk  --> 1 TB NVMe SSD
```

```

└─nvme1n1p1 259:9  0  155G 0 part  --> Part of images-pool
└─nvme1n1p2 259:10 0  155G 0 part  --> Part of images-pool
└─nvme1n1p3 259:11 0  155G 0 part  --> Part of imgpool-pool
└─nvme1n1p4 259:12 0  155G 0 part  --> Part of disks-pool
└─nvme1n1p5 259:13 0  155G 0 part  --> Part of disks-pool
└─nvme1n1p6 259:14 0  155G 0 part  --> Part of iso-pool
└─nvme1n1p7 259:15 0  1.5G 0 part  --> Part of iso-pool

```

Finally let us go ahead and create 3 zfs pools as planned.

```

# zpool create images-pool /dev/nvme1n1p1 /dev/nvme1n1p2 /dev/nvme1n1p3
# zpool create disks-pool /dev/nvme1n1p4 /dev/nvme1n1p5
# zpool create iso-pool /dev/nvme1n1p6 /dev/nvme1n1p7

```

Let us list all the zfs pools that we have created on both the SSDs

```

# zpool list


| NAME        | SIZE  | ALLOC | FREE  | CKPOINT | EXPANDSZ | FRAG | CAP | DEDUP | HEALTH | ALTROOT |
|-------------|-------|-------|-------|---------|----------|------|-----|-------|--------|---------|
| basepool    | 5.50G | 5.62M | 5.49G | -       | -        | 2%   | 0%  | 1.00x | ONLINE | -       |
| disks-pool  | 308G  | 694K  | 308G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| images-pool | 462G  | 797K  | 462G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| iso-pool    | 154G  | 702K  | 154G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| zpool-1     | 316G  | 916K  | 316G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| zpool-2     | 316G  | 636K  | 316G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| zpool-3     | 316G  | 646K  | 316G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| zpool-4     | 316G  | 634K  | 316G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| zpool-5     | 316G  | 632K  | 316G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |
| zpool-6     | 316G  | 634K  | 316G  | -       | -        | 0%   | 0%  | 1.00x | ONLINE | -       |


```

Task #05: Installation and Configuration of incus

Here are the list of things that we are going to accomplish in this task

1. Installiation of incus
2. Installation verification

1. Installation of incus

The commands to install incus is pretty straight forward. Though there are multiple ways to install incus, let us install from the repository maintained by zabbly.com - the sponsors, developers and maitainers of this open source project incus.

```
# uname -a
Linux red-handsonbox 6.8.0-55-generic #57-Ubuntu SMP PREEMPT_DYNAMIC Wed Feb 12 23:42:21
UTC 2025 x86_64 x86_64 x86_64 GNU/Linux

# cat /etc/os-release
-----
PRETTY_NAME="Ubuntu 24.04.2 LTS"
NAME="Ubuntu"
VERSION_ID="24.04"
VERSION="24.04.2 LTS (Noble Numbat)"
VERSION_CODENAME=noble
ID=ubuntu
ID_LIKE=debian
HOME_URL="https://www.ubuntu.com/"
SUPPORT_URL="https://help.ubuntu.com/"
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"
UBUNTU_CODENAME=noble
LOGO=ubuntu-logo
-----

# curl -fsSL https://pkgs.zabbly.com/key.asc | gpg --show-keys --fingerprint
-----
gpg: directory '/root/.gnupg' created
gpg: keybox '/root/.gnupg/pubring.kbx' created
pub    rsa3072 2023-08-23 [SC] [expires: 2025-08-22]
        4EFC 5906 96CB 15B8 7C73 A3AD 82CC 8797 C838 DCFD
uid            Zabbly Kernel Builds <info@zabbly.com>
sub    rsa3072 2023-08-23 [E] [expires: 2025-08-22]
-----
```

```

# mkdir -p /etc/apt/keyrings/

# curl -fsSL https://pkgs.zabbly.com/key.asc -o /etc/apt/keyrings/zabbly.asc

# sh -c 'cat <<EOF > /etc/apt/sources.list.d/zabbly-incus-stable.sources
Enabled: yes
Types: deb
URIs: https://pkgs.zabbly.com/incus/stable
Suites: $(. /etc/os-release && echo ${VERSION_CODENAME})
Components: main
Architectures: $(dpkg --print-architecture)
Signed-By: /etc/apt/keyrings/zabbly.asc

EOF'

# apt-get update

# apt install net-tools bridge-utils -y

# apt-get install incus -y

# systemctl restart incus.service

# systemctl status -l incus.service

# systemctl enable incus.service

```

The parameters passed to the command ‘incus admin init’ is important to understand and here is the breakup.

- 1) We are telling the command/installer not to configure any Clustering setup
- 2) We are planning to configure a default storage pool
- 3) We are passing the name ‘basepool’ to be used as the default storage. As you may recollect we have already created a zfs pool named ‘basepool’ in the previous chapter where it was created on Partition #7 of the 2 GB SSD
- 4) We intended to use zfs filesystem instead of other filesystems
- 5) We are telling the command/installer not to create a new zfs pool but use an existing zfs pool based storage pool named ‘basepool’ that we already pre-created and is ready for use
- 6) With regards to network configuration, we are going with the suggested name ‘incusbr0’ to be used to create a network bridge
- 7) We are assigning a static ip address of our choice (10.1.1.253/24) to the incus bridge (incusbr0)

The rest of the parameters passed are self explanatory

```
# incus admin init
-----
Would you like to use clustering? (yes/no) [default=no]:
Do you want to configure a new storage pool? (yes/no) [default=yes]:
Name of the new storage pool [default=default]: basepool
Name of the storage backend to use (zfs, btrfs, dir, lvm, lvmcluster) [default=zfs]:
Create a new ZFS pool? (yes/no) [default=yes]: no
Name of the existing ZFS pool or dataset: basepool
Would you like to create a new local network bridge? (yes/no) [default=yes]:
What should the new bridge be called? [default=incusbr0]:
What IPv4 address should be used? (CIDR subnet notation, "auto" or "none") [default=auto]:
10.1.1.253/24
Would you like to NAT IPv4 traffic on your bridge? [default=yes]:
What IPv6 address should be used? (CIDR subnet notation, "auto" or "none") [default=auto]:
none
Would you like the server to be available over the network? (yes/no) [default=no]: yes
Address to bind to (not including port) [default=all]:
Port to bind to [default=8443]:
Would you like stale cached images to be updated automatically? (yes/no) [default=yes]:
Would you like a YAML "init" preseed to be printed? (yes/no) [default=no]: yes
config:
  core.https_address: '[::]:8443'
networks:
- config:
    ipv4.address: 10.1.1.253/24
    ipv4.nat: "true"
    ipv6.address: none
    description: ""
    name: incusbr0
    type: ""
    project: default
storage_pools:
- config:
    source: basepool
    description: ""
    name: basepool
    driver: zfs
storage_volumes: []
profiles:
- config: {}
  description: ""
  devices:
```

```

eth0:
  name: eth0
  network: incusbr0
  type: nic
root:
  path: /
  pool: basepool
  type: disk
name: default
project: default
projects: []
cluster: null
-----

```

2. Installation verification

```

# ps -ef | grep dns
incus      21299  21299  0 04:14 ?        00:00:00 dnsmasq --keep-in-foreground
--strict-order --bind-interfaces --except-interface=lo --pid-file= --no-ping
--interface=incusbr0 --dhcp-rapid-commit --no-negcache --quiet-dhcp --quiet-dhcp6
--quiet-ra --listen-address=10.1.1.253 --dhcp-no-override --dhcp-authoritative
--dhcp-leasefile=/var/lib/incus/networks/incusbr0/dnsmasq.leases
--dhcp-hostsfile=/var/lib/incus/networks/incusbr0/dnsmasq.hosts --dhcp-range
10.1.1.2,10.1.1.254,1h -s incus --interface-name _gateway.incus,incusbr0 -S /incus/
--conf-file=/var/lib/incus/networks/incusbr0/dnsmasq.raw -u incus -g incus

```

```

# brctl show incusbr0
bridge name bridge id      STP enabled interfaces
incusbr0      8000.00163e1b7238  no

```

```

# incus network show incusbr0
config:
  ipv4.address: 10.1.1.253/24
  ipv4.nat: "true"
  ipv6.address: none
description: ""
name: incusbr0
type: bridge
used_by:
- /1.0/profiles/default
managed: true
status: Created
locations:

```

- none

project: default

```
# incus network list -f compact
```

NAME	TYPE	MANAGED	IPV4	IPV6	DESCRIPTION	USED BY	STATE
br-int	bridge	NO				0	
docker0	bridge	NO				0	
eno1	physical	NO				0	
eno2	physical	NO				0	
eno3	physical	NO				0	
eno4	physical	NO				0	
incusbr0	bridge	YES	10.1.1.253/24	none		1	CREATED
lo	loopback	NO				0	

Task #06: Incus Storage

Here are the list of things that we are going to accomplish in this task.

1. incus managing storage pools on 2 TB SSD
2. incus managing storage pools on 1 TB SSD

1. incus managing storage pools of 2 TB SSD

Let us take a quick look the output of ‘lsblk’ command and see how we have partitioned the SSDs and what zfs pools reside on each partition.

```
# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    1 460.3G 0 disk
└─sda1     8:1    1   1M 0 part
└─sda2     8:2    1 460.3G 0 part / --> Root partition on USB thumb drive
nvme0n1   259:0   0  1.9T 0 disk  --> 2 TB NVMe SSD
└─nvme0n1p1 259:1   0 317G 0 part  --> zfs pool named zpool-1
└─nvme0n1p2 259:2   0 317G 0 part  --> zfs pool named zpool-2
└─nvme0n1p3 259:3   0 317G 0 part  --> zfs pool named zpool-3
└─nvme0n1p4 259:4   0 317G 0 part  --> zfs pool named zpool-4
└─nvme0n1p5 259:5   0 317G 0 part  --> zfs pool named zpool-5
└─nvme0n1p6 259:6   0 317G 0 part  --> zfs pool named zpool-6
└─nvme0n1p7 259:7   0  5.7G 0 part  --> zfs pool named basepool
nvme1n1   259:8   0 931.5G 0 disk  --> 1 TB NVMe SSD
└─nvme1n1p1 259:9   0 155G 0 part
└─nvme1n1p2 259:10  0 155G 0 part
└─nvme1n1p3 259:11  0 155G 0 part
└─nvme1n1p4 259:12  0 155G 0 part
└─nvme1n1p5 259:13  0 155G 0 part
└─nvme1n1p6 259:14  0 155G 0 part
└─nvme1n1p7 259:15  0  1.5G 0 part
```

Let us list all the zfs pools that we have created so far on each partition of both the SSDs.

```
# zpool list
NAME      SIZE  ALLOC   FREE  CKPOINT  EXPANDSZ   FRAG    CAP  DEDUP   HEALTH  ALTROOT
basepool  5.50G  6.49M  5.49G          -          -   6%    0%  1.00x  ONLINE  - --> on 2TB SSD
disks-pool 308G  694K  308G          -          -   0%    0%  1.00x  ONLINE  -
images-pool 462G  681K  462G          -          -   0%    0%  1.00x  ONLINE  -
iso-pool   154G  657K  154G          -          -   0%    0%  1.00x  ONLINE  -
zpool-1   316G  746K  316G          -          -   0%    0%  1.00x  ONLINE  - --> on 2TB SSD
```

zpool-2	316G	636K	316G	-	-	0%	0%	1.00x	ONLINE	- -> on 2TB SSD
zpool-3	316G	646K	316G	-	-	0%	0%	1.00x	ONLINE	- -> on 2TB SSD
zpool-4	316G	634K	316G	-	-	0%	0%	1.00x	ONLINE	- -> on 2TB SSD
zpool-5	316G	632K	316G	-	-	0%	0%	1.00x	ONLINE	- -> on 2TB SSD
zpool-6	316G	634K	316G	-	-	0%	0%	1.00x	ONLINE	- -> on 2TB SSD

We will now go ahead and hand over the zfs pools that we have created on 2TB SSDs to incus for managing storage resources.

Note that the name ‘Hands-on-lab’ is being used in the description. The reason being, soon we will be creating Projects with the same name (Handson-lab). That way it will be easier to identify which storage pool is assigned to which Project. This will be much clearer with following tasks.

```
# incus storage create storage-1 zfs source=zpool-1 --description="Storage for Containers,
VMs of Handson-lab-1"
# incus storage create storage-2 zfs source=zpool-2 --description="Storage for Containers,
VMs of Handson-lab-2"
# incus storage create storage-3 zfs source=zpool-3 --description="Storage for Containers,
VMs of Handson-lab-3"
# incus storage create storage-4 zfs source=zpool-4 --description="Storage for Containers,
VMs of Handson-lab-4"
# incus storage create storage-5 zfs source=zpool-5 --description="Storage for Containers,
VMs of Handson-lab-5"
# incus storage create storage-6 zfs source=zpool-6 --description="Storage for Containers,
VMs of Handson-lab-6"
```

You may recollect that ‘basepool’ is already under the control of incus as that is what we passed to ‘incus admin init’ command at initial setup. Though this step is optional, let us manually edit the configuration of basepool , modify its description so that it will be verbatim

```
# export EDITOR=vi

# incus storage edit basepool
-----
config:
  source: basepool
  volatile.initial_source: basepool
  zfs.pool_name: basepool
  description: passed to incus admin init
  name: basepool
  driver: zfs
  used_by:
    - /1.0/profiles/default
```

```
status: Created
locations:
- none
-----
```

Let us now list all the storage pools managed by incus

```
# incus storage list -f compact
```

NAME	DRIVER	DESCRIPTION	USED BY	STATE
basepool	zfs	passed to incus admin init	1	CREATED
storage-1	zfs	Storage for Containers, VMs of Handson-lab-1	0	CREATED
storage-2	zfs	Storage for Containers, VMs of Handson-lab-2	0	CREATED
storage-3	zfs	Storage for Containers, VMs of Handson-lab-3	0	CREATED
storage-4	zfs	Storage for Containers, VMs of Handson-lab-4	0	CREATED
storage-5	zfs	Storage for Containers, VMs of Handson-lab-5	0	CREATED
storage-6	zfs	Storage for Containers, VMs of Handson-lab-6	0	CREATED

From now on, we will refer to the storage pools with their assigned names storage-1, storage-2, storage-3 and so on.

2. incus managing storage pools of 1 TB SSD

Here is the current status of the 1TB SSD

```
# lsblk
```

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINTS
sda	8:0	1	460.3G	0	disk	
└─sda1	8:1	1	1M	0	part	
└─sda2	8:2	1	460.3G	0	part	/ --> Root partition on USB thumb drive
nvme0n1	259:0	0	1.9T	0	disk	--> 2 TB NVMe SSD
└─nvme0n1p1	259:1	0	317G	0	part	
└─nvme0n1p2	259:2	0	317G	0	part	
└─nvme0n1p3	259:3	0	317G	0	part	
└─nvme0n1p4	259:4	0	317G	0	part	
└─nvme0n1p5	259:5	0	317G	0	part	
└─nvme0n1p6	259:6	0	317G	0	part	
└─nvme0n1p7	259:7	0	5.7G	0	part	
nvme1n1	259:8	0	931.5G	0	disk	--> 1 TB NVMe SSD
└─nvme1n1p1	259:9	0	155G	0	part	--> part of images-pool
└─nvme1n1p2	259:10	0	155G	0	part	--> part of images-pool
└─nvme1n1p3	259:11	0	155G	0	part	--> part of images-pool
└─nvme1n1p4	259:12	0	155G	0	part	--> part of disks-pool
└─nvme1n1p5	259:13	0	155G	0	part	--> part of disks-pool
└─nvme1n1p6	259:14	0	155G	0	part	--> part of iso-pool

```
└─nvme1n1p7 259:15 0 1.5G 0 part --> part of iso-pool
```

Let us list all the storage pools that reside on both the SSDs

```
# zpool list
```

NAME	SIZE	ALLOC	FREE	CKPOINT	EXPANDSZ	FRAG	CAP	DEDUP	HEALTH	ALTROOT
basepool	5.50G	6.49M	5.49G	-	-	6%	0%	1.00x	ONLINE	-
disks-pool	308G	694K	308G	-	-	0%	0%	1.00x	ONLINE	- --> On 1TB SSD
images-pool	462G	681K	462G	-	-	0%	0%	1.00x	ONLINE	- --> On 1TB SSD
iso-pool	154G	657K	154G	-	-	0%	0%	1.00x	ONLINE	- --> On 1TB SSD
zpool-1	316G	746K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-2	316G	636K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-3	316G	646K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-4	316G	634K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-5	316G	632K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-6	316G	634K	316G	-	-	0%	0%	1.00x	ONLINE	-

Let us now hand over the zfs pools that we have created on 1 TB SSDs to incus for managing storage resources.

```
# incus storage create shared-images zfs source=images-pool --description="shared images accross Projects"
```

```
# incus storage create shared-disks zfs source=disks-pool --description="shared disks accross Projects"
```

```
# incus storage create shared-iso zfs source=iso-pool --description="shared ISOs accross Projects"
```

Here is the final list of all storage pools managed by incus

```
# incus storage list -f compact
```

NAME	DRIVER	DESCRIPTION	USED BY	STATE
basepool	zfs	passed to incus admin init	1	CREATED
shared-disks	zfs	shared disks accross Projects	0	CREATED
shared-images	zfs	shared images accross Projects	0	CREATED
shared-iso	zfs	shared ISOs accross Projects	0	CREATED
storage-1	zfs	Storage for Containers, VMs of Handson-lab-1	0	CREATED
storage-2	zfs	Storage for Containers, VMs of Handson-lab-2	0	CREATED
storage-3	zfs	Storage for Containers, VMs of Handson-lab-3	0	CREATED
storage-4	zfs	Storage for Containers, VMs of Handson-lab-4	0	CREATED
storage-5	zfs	Storage for Containers, VMs of Handson-lab-5	0	CREATED
storage-6	zfs	Storage for Containers, VMs of Handson-lab-6	0	CREATED

Task # 07: OVN Installation

OVN stands for Open Virtual Networking. The primary reason to install OVN networking with incus in our setup is to reuse the same ip addresses and ip address ranges assigned to all the virtual machines and containers across projects - all at the same time without any conflicts. This is an advanced concept and is easy to get confused. So we will take it slowly to understand it better in a way we can build them in a methodical way in the upcoming tasks.

In this task we will focus on installing ovn networking packages, which is pretty straight forward.

```
# apt install ovn-host ovn-central -y

# ovs-vsctl set open_vswitch . \
    external_ids:ovn-remote=unix:/run/ovn/ovnsb_db.sock \
    external_ids:ovn-encap-type=geneve \
    external_ids:ovn-encap-ip=127.0.0.1

# incus network list -f compact
  NAME      TYPE    MANAGED     IPV4      IPV6  DESCRIPTION  USED BY   STATE
  br-int    bridge   NO          0
  docker0   bridge   NO          0
  eno1      physical NO          0
  eno2      physical NO          0
  eno3      physical NO          0
  eno4      physical NO          0
  incusbr0  bridge   YES         10.1.1.253/24 none        1       CREATED
  lo        loopback NO          0
  ovs-system unknown  NO          0
```

Here are the details:

- a. ‘eno3’ is the physical network interface where host ip address (192.168.1.10) is configured.
This is the only uplink on the host
- b. eno1, eno2 and eno4 are interfaces where none of them are configured
- c. ‘docker0’ is the docker bridge that was created while installing docker
- d. ‘incusbr0’ is the incus bridge that was created while installing incus
- e. ‘ovn-system’ is what we installed just now

Task #08: Installation of incus-ui-canonical

Apart from the command line, incus platform can be administered from a browser as well. The package providing web UI is straight forward, which is shown below.

There is yet another open source project named LXConsole to manage incus via browser. For now we will focus and use incus-ui-canonical.

```
# apt install incus-ui-canonical -y
-----
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  incus-ui-canonical
0 upgraded, 1 newly installed, 0 to remove and 2 not upgraded.
Need to get 3,728 kB of archives.
After this operation, 20.1 MB of additional disk space will be used.
Get:1 https://pkgs.zabbly.com/incus/stable noble/main amd64 incus-ui-canonical amd64
1:6.10.1-ubuntu24.04-202503210251 [3,728 kB]
Fetched 3,728 kB in 2s (1,655 kB/s)
Selecting previously unselected package incus-ui-canonical.
(Reading database ... 90200 files and directories currently installed.)
Preparing to unpack .../incus-ui-canonical_1%3a6.10.1-ubuntu24.04-202503210251_amd64.deb ...
Unpacking incus-ui-canonical (1:6.10.1-ubuntu24.04-202503210251) ...
Setting up incus-ui-canonical (1:6.10.1-ubuntu24.04-202503210251) ...
Scanning processes...
Scanning processor microcode...
Scanning linux images...
```

Running kernel seems to be up-to-date.

The processor microcode seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.

```
-----  
# incus config set core.https_address :8443
```

```
# incus config show
```

```
-----
```

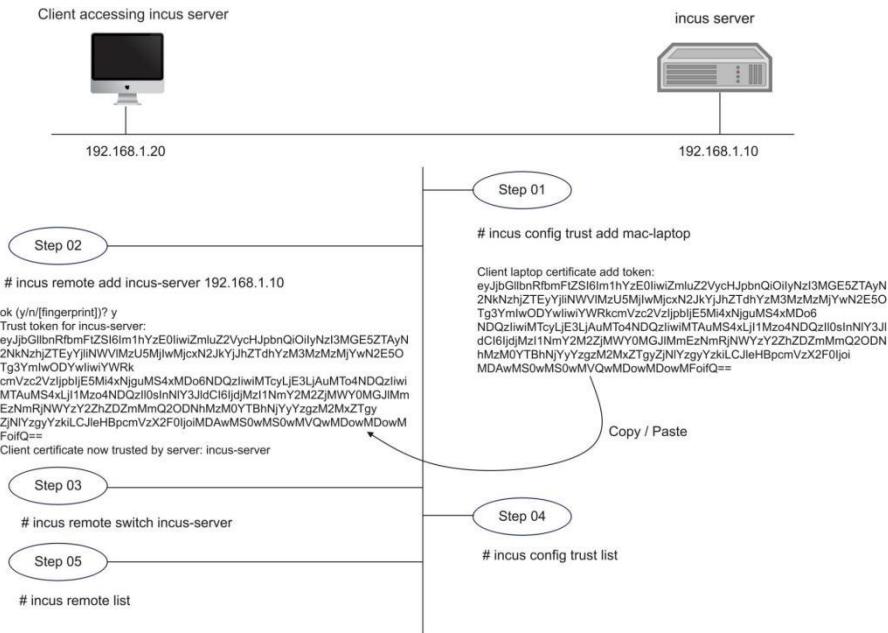
```
config:
```

```
  core.https_address: :8443
```

```
-----
```

Task #09: Incus Remote access setup

Once the incus server is up and running, configuring a client machine to remotely access the incus server needs a simple setup. Shown below is a workflow explaining the required steps on both the incus server and the client machine.



In the commands shown below, hostname of the incus server is ‘red-handsonbox’ and the name of the client machine which will eventually access the incus server remotely is ‘mac’.

Step 01:

On the incus server, add the name of the client machine to its trust list. Entering the command shown below will generate a token

```
# incus config trust add mac
```

Client mac certificate add token:

```
eyJjbGllbnRfbmFtZSI6Im1hYzE0IiwiZmluZ2VycHJpbnQi0iIyNzI3MGE5ZTAyN2NkNzhjZTEyYjliNWVlMzU  
5MjIwMjcxN2JkYjJhZTdhYzM3MzMzMjYwN2E50Tg3YmIwODYwIiwiYWRkcmVzc2VzIjpBIjE5Mi4xNjguMS4xMD  
o6NDQzIiwiMTcyLjE3LjAuMTo4NDQzIiwiMTAuMS4xLjI1Mzo4NDQzI10sInN1Y3JldCI6IjdjMzI1NmY2M2ZjM  
WY0MGJlMmEzNmRjNwYzY2ZhZDZmMmQ2ODNhMzM0YTbhNjYyYzgzM2MxZTgyZjN1YzgyYzkiLCJleHBpcmVzX2F0  
IjoiNDAwMS0wMS0wMVQwMDowMDowMFoifQ==
```

Step 02:

On the client (mac) laptop, run the command shown below, followed by the generated token on the server

```

mac$ incus remote add red-handsonbox 192.168.1.10
ok (y/n/[fingerprint])? y
Trust token for red-handsonbox:
eyJjbGllbnRfbmFtZSI6Im1hYzE0IiwiZmluZ2VycHJpbnQi0iIyNzI3MGE5ZTAyN2NkNzhjZTEyYjliNWVlMzU
5MjIwMjcxN2JkYjJhZTdhYzM3MzMjYwN2E50Tg3YmIw0DYwIiwiYWRkcmVzc2VzIjpbIjE5Mi4xNjguMS4xMD
o6NDQzIiwiMTcyLjE3LjAuMTo4NDQzIiwiMTAuMS4xLjI1Mzo4NDQzIl0sInNlY3JldCI6IjdjMzI1NmY2M2ZjM
WY0MGJlMmEzNmRjNWYzY2ZhZDZmMmQ20DNhMzM0YTbhNjYyYzgzM2MxZTgyZjNlYzgyYzkiLCJleHBpcmVzX2F0
IjoiNDAwMS0wMS0wMVQwMDowMDowMFoifQ==
Client certificate now trusted by server: red-handsonbox

```

Step 03:

On the client (mac) laptop, run the command to switch context to the newly added remote incus server

```
mac$ incus remote switch red-handsonbox
```

Step 04:

On the client, verify that incus server is sucessfully added to its remote list

```

mac$ incus remote list -f compact
NAME URL PROTOCOL AUTH TYPE PUBLIC STATIC GLOBAL
images https://images.linuxcontainers.org simplestreams none YES NO NO
local unix:// incus file access NO YES NO
red-handsonbox (current) https://192.168.1.10:8443 incus tls NO NO NO

```

Step 05:

On the incus serer, verify that the client ‘mac’ is added successfully to its trust list

```
# incus config trust list -f compact
NAME      TYPE    DESCRIPTION FINGERPRINT          EXPIRY DATE
mac       client      c88b6a7dd7e8 2035/01/04 21:59 UTC
```

Now that the trust for mac laptop is established on the incus server (red-handonbox), we can issue commands directly from a shell terminal on Mac laptop (represented as ‘mac%’)

Task #10: Incus web UI

In this task we will take a quick look at incus web UI.

Now that we have established trust between the mac client and the incus server, let us run the following command on the mac machine, which will open up a browser and point to incus web UI

```
mac$ incus webui
```

Web server running at:

```
http://127.0.0.1:53100/ui?auth_token=eb79cec0-eef6-480c-b6e2-3178c62e91e5
```

This would be the landing page. It currently shows we do not have any instances of containers or virtual machines running.

The screenshot shows the Incus web UI interface. On the left, a sidebar titled 'Incus UI' contains a 'Project' dropdown set to 'default', and a tree view with nodes like 'Instances', 'Profiles', 'Networking', 'Storage', 'Images', 'Configuration', 'Operations', 'Warnings', and 'Settings'. Below the sidebar are links for 'c88b6a7dd7e0fa5f...', 'Documentation', 'Discussion', and 'Report a bug'. The main content area has a heading 'Instances Ø' and a message 'No instances found' with the subtext 'There are no instances in this project. Spin up your first instance!'. A blue 'Create instance' button is prominently displayed. At the bottom of the page, a footer bar includes a double-left arrow icon and the text 'Version 6.12-ui-0.15'.

Let us take a look at the Networks

The screenshot shows the Incus UI interface with the 'Networks' section selected in the sidebar. The main table lists the following network interfaces:

NAME	TYPE	MANAGED	IPV4	IPV6	DESCRIPTION	FORWARDS	USED BY	STATE
incusbr0	bridge	Yes	10.1.1.253/24	none		0	1	Created
eno1	physical	No				-	0	
eno2	physical	No				-	0	
eno3	physical	No				-	0	
eno4	physical	No				-	0	
br-int	bridge	No				-	0	
docker0	bridge	No				-	0	

At the bottom left, there are links for 'c88b6a7dd7e8fa5f...', 'Documentation', 'Discussion', and 'Report a bug'. At the bottom right, it says 'Version 6.12-ui-0.15'.

Will go ahead and check storage pools

The screenshot shows the Incus UI interface with the 'Pools' section selected in the sidebar. The main table lists the following storage pools:

NAME	DRIVER	SIZE	VOLUMES (THIS PROJECT)	VOLUMES (ALL PROJECTS)	STATUS
basepool	zfs	5.6 MiB of 5.3 GiB used	0	0	Created
shared-disks	zfs	1.9 MiB of 298.4 GiB used	0	0	Created
shared-images	zfs	1.1 MiB of 447.7 GiB used	1	1	Created
shared-iso	zfs	831.0 KiB of 149.2 GiB used	0	0	Created
storage-1	zfs	1.0 MiB of 306.2 GiB used	0	0	Created
storage-2	zfs	814.5 KiB of 306.2 GiB used	0	0	Created
storage-3	zfs	723.0 KiB of 306.2 GiB used	0	0	Created
storage-4	zfs	717.0 KiB of 306.2 GiB used	0	0	Created
storage-5	zfs	739.5 KiB of 306.2 GiB used	0	0	Created
storage-6	zfs	748.5 KiB of 306.2 GiB used	0	0	Created

At the bottom left, there are links for 'c88b6a7dd7e8fa5f...', 'Documentation', 'Discussion', and 'Report a bug'. At the bottom right, it says 'Version 6.12-ui-0.15'.

Finally we will take a look at the details of the default Profile

Profiles Search

Showing 1 out of 1 profile

NAME	DESCRIPTION	USED BY
default	Default Incus profile	0 instances 0 in all projects

[+ Create profile](#)

Version 6.12-ui-0.15

Profiles / default

[Delete](#)

Overview Configuration

General

	Name	Description
	default	Default Incus profile

Networks

NAME	INTERFACE	TYPE	MANAGED
incusbr0	eth0	bridge	Yes

Devices

NAME	TYPE	DETAILS
root	disk (root)	Pool (basepool)

Limits

CPU	Memory
-	-

Usage (0)

Version 6.12-ui-0.15

Incus UI

Profiles / default

Delete

Project: default

Instances

Profiles

Networking

ACLs

Storage

Pools

Volumes

Custom ISOs

Images

Configuration

Operations

Warnings

Settings

c88b6a7dd7e8fa5f...

Documentation

Discussion

Report a bug

Overview Configuration

Main configuration

Devices

Disk

Network

GPU

Proxy

Other

Resource limits

Security policies

Snapshots

Migration

Boot

Cloud init

Root storage

CONFIGURATION INHERITED OVERRIDE

Root storage ~~x~~

basepool ~~l~~

unlimited ~~l~~

+ Attach disk device

YAML Configuration

<< Version 6.12-ui-0.15

The screenshot shows the Incus UI interface for managing profiles. The left sidebar contains a navigation tree with sections like Project, Instances, Profiles, Networking, ACLs, Storage, Images, Configuration, Operations, Warnings, and Settings. Below the sidebar, there are links for Documentation, Discussion, and Report a bug. The main content area is titled 'Profiles / default' and has tabs for Overview and Configuration, with Configuration being active. Under Configuration, there's a 'Main configuration' section with a 'Devices' dropdown. The 'Disk' option is selected, showing configuration for 'Root storage'. It lists 'CONFIGURATION', 'INHERITED', and 'OVERRIDE' columns. The 'Root storage' row shows 'basepool' in the OVERRIDE column with a delete icon. A button '+ Attach disk device' is available. At the bottom, there's a 'YAML Configuration' toggle and a footer note about the version.

Task #11: Sanity test with a Container

Let us do a sanity testing by launching an ubuntu container. By analyzing the container we see where things stand with regards to what we have built so far.

```
# incus launch images:ubuntu/24.04 c1

# incus list
+-----+-----+-----+-----+
| NAME | STATE | IPV4 | IPV6 | TYPE | SNAPSHOTS |
+-----+-----+-----+-----+
| c1 | RUNNING | 10.1.1.57 (eth0) | | CONTAINER | 0 |
+-----+-----+-----+-----+
```

Let us get inside the newly spawned container (c1) and ping google dns server, which confirms the network is working

```
# incus shell c1
root@c1:~# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=115 time=22.0 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=115 time=23.1 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=115 time=22.4 ms
```

Let us go ahead and update, upgrade ubuntu distro installed in the container (c1)

```
root@c1:~# apt -y update
Hit:1 http://archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:3 http://archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:4 http://archive.ubuntu.com/ubuntu noble-updates/main amd64 Packages [922 kB]
Fetched 1,174 kB in 1s (876 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
All packages are up to date.
```

```
root@c1:~# apt -y upgrade
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
Calculating upgrade... Done
0 upgraded, 0 newly installed, 0 to remove and 0 not upgraded.
```

```
root@c1:~#
```

Now let us install package for ‘traceroute’

```
root@c1:~# apt install -y traceroute
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  traceroute
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 60.5 kB of archives.
After this operation, 162 kB of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu noble/universe amd64 traceroute amd64 1:2.1.5-1
[60.5 kB]
Fetched 60.5 kB in 1s (107 kB/s)
Selecting previously unselected package traceroute.
(Reading database ... 16181 files and directories currently installed.)
Preparing to unpack .../traceroute_1%3a2.1.5-1_amd64.deb ...
Unpacking traceroute (1:2.1.5-1) ...
Setting up traceroute (1:2.1.5-1) ...
update-alternatives: using /usr/bin/traceroute.db to provide /usr/bin/traceroute
(traceroute) in auto mode
update-alternatives: using /usr/bin/traceroute6.db to provide /usr/bin/traceroute6
(traceroute6) in auto mode
update-alternatives: using /usr/bin/lft.db to provide /usr/bin/lft (lft) in auto mode
update-alternatives: using /usr/bin/traceproto.db to provide /usr/bin/traceproto
(traceproto) in auto mode
update-alternatives: using /usr/sbin/tcptraceroute.db to provide /usr/sbin/tcptraceroute
(tcptraceroute) in auto mode
root@c1:~#
```

Let us confirm that the network traffic does go through newly created incus bridge (incusbr0) and then via its uplink physical nic - en03 in this case.

```
root@c1:~# traceroute zabbly.com
traceroute to zabbly.com (45.45.148.7), 30 hops max, 60 byte packets
 1 _gateway.incus (10.1.1.253)  0.059 ms  0.023 ms  0.017 ms
 2 _gateway (192.168.1.1)  0.336 ms  0.872 ms  0.832 ms
 3 10.61.209.3 (10.61.209.3)  10.462 ms  16.893 ms  10.61.209.2 (10.61.209.2)  16.817 ms
 4 * * *
.....
.....
.....
```

```
root@c1:~# exit
logout
```

Let us take a look at the configuration of the newly spawned container (c1)

```
# incus config show c1
architecture: x86_64
config:
  image.architecture: amd64
  image.description: Ubuntu noble amd64 (20250503_07:42)
  image.os: Ubuntu
  image.release: noble
  image.requirements.cgroup: v2
  image.serial: "20250503_07:42"
  image.type: squashfs
  image.variant: default
  volatile.base_image: b3f31aa5264c50ba425e769160f586b43a07248c41f1b1680b1b634fd0b1eac
  volatile.cloud-init.instance-id: 1be64a01-047e-43fd-8157-5e72176d5a95
  volatile.eth0.host_name: veth41f59eb7
  volatile.eth0.hwaddr: 10:66:6a:67:88:e0
  volatile.idmap.base: "0"
  volatile.idmap.current:
'[{{"Isuid":true,"Isgid":false,"Hostid":1000000,"Nsid":0,"Maprange":1000000000},{"Isuid":false,"Isgid":true,"Hostid":1000000,"Nsid":0,"Maprange":1000000000}]'
  volatile.idmap.next:
'[{{"Isuid":true,"Isgid":false,"Hostid":1000000,"Nsid":0,"Maprange":1000000000},{"Isuid":false,"Isgid":true,"Hostid":1000000,"Nsid":0,"Maprange":1000000000}]'
  volatile.last_state.idmap: '[]'
  volatile.last_state.power: RUNNING
  volatile.uuid: f7bb3e51-b902-4218-ab45-018f66d2feeb
  volatile.uuid.generation: f7bb3e51-b902-4218-ab45-018f66d2feeb
devices: {}
ephemeral: false
profiles:
- default
stateful: false
description: ""
```

Let us look at the values of 2 configuration parameters namely ‘volatile.base_image’ and ‘profile’ and see what they reveal

# incus image list -f compact	ALIAS	FINGERPRINT	PUBLIC	DESCRIPTION	ARCHITECTURE	TYPE
SIZE	UPLOAD	DATE				

```
b3f31aaf5264 no      Ubuntu noble amd64 (20250503_07:42) x86_64      CONTAINER
129.19MiB 2025/05/04 05:10 UTC
```

The finger print value reveals from which image the container is spawned (which is the image Id of ubuntu24.04 from the public repository “images”). The following command shows the same:

```
# incus image list images: b3f31aaf5264 -f compact
          ALIAS      FINGERPRINT      PUBLIC      DESCRIPTION
ARCHITECTURE   TYPE      SIZE      UPLOAD DATE
    ubuntu/noble (7 more) b3f31aaf5264 yes      Ubuntu noble amd64 (20250503_07:42) x86_64
CONTAINER 129.19MiB 2025/05/03 00:00 UTC
```

```
# incus image show b3f31aaf5264
auto_update: true
properties:
  architecture: amd64
  description: Ubuntu noble amd64 (20250503_07:42)
  os: Ubuntu
  release: noble
  requirements.cgroup: v2
  serial: "20250503_07:42"
  type: squashfs
  variant: default
public: false
expires_at: 1970-01-01T00:00:00Z
profiles:
- default
```

Now let us look at the value of profile (which is ‘default’) shown by the following command:

```
# incus profile list
+-----+-----+-----+
| NAME |      DESCRIPTION      | USED BY |
+-----+-----+-----+
| default | Default Incus profile | 1       |
+-----+-----+-----+
```

```
# incus profile show default
config: {}
description: Default Incus profile
devices:
eth0:
  name: eth0
  network: incusbr0
```

```

type: nic
root:
  path: /
  pool: basepool
  type: disk
name: default
used_by:
- /1.0/instances/c1
project: default

```

The value of ‘pool’ shows ‘basepool’ which means the root disk of this container (c1) is located in the storage named basepool. We can display the configuration of the storage basepool with the command:

```

# incus storage show basepool
config:
  source: basepool
  volatile.initial_source: basepool
  zfs.pool_name: basepool
description: passed to incus admin init
name: basepool
driver: zfs
used_by:
- /1.0/images/b3f31aaf5264c50ba425e769160f586b43a07248c41f1b1680b1b634fd0b1eac
- /1.0/instances/c1
- /1.0/profiles/default
status: Created
locations:
- none

```

The values listed under ‘used_by’ reveals exactly what we have seen so far:

- a) the name of the container (**c1**)
- b) the image id used to spawn the container (**b3f31aaf5264**)
- c) the profile (**default**) used.

Finally let us do cleanup. Let us stop and delete the container

```
# incus stop c1 ; incus delete c1
```

Task #12: Creating a Volume

In this task we will work on creating a new storage volume. We will create this volume on a specific incus storage pool (shared-images) that we have already created.

This storage volume can be used as a local repository where we can store all the images that can be downloaded from public image servers that are available on the internet. Having this kind of storage volume will be handy when we want to build solutions that can be air-gaped. With this approach, it can save us time and network bandwidth on downloading images from public images servers.

As you may recall, we have already created an incus storage pool named ‘shared-images’ (backed by zfs pool named ‘images-pool’). It is on this storage pool we will create a new volume for the purpose explained above.

```
# incus storage show shared-images
config:
  source: images-pool
  volatile.initial_source: images-pool
  zfs.pool_name: images-pool
description: shared images accross Projects
name: shared-images
driver: zfs
used_by: []
status: Created
locations:
- none
```

Let us create a volume on it with the name ‘global-volume’.

```
# incus storage volume create shared-images global-volume
```

Let us take a look at the same storage pool again. Note that the newly created ‘global-volume’ is now listed under ‘used_by’:

```
# incus storage show shared-images
config:
  source: images-pool
  volatile.initial_source: images-pool
  zfs.pool_name: images-pool
description: shared images accross Projects
name: shared-images
driver: zfs
used_by:
- global-volume
```

```

used_by:
- /1.0/storage-pools/shared-images/volumes/custom/global-volume
status: Created
locations:
- none

# incus storage volume show shared-images global-volume
config: {}
description: ""
name: global-volume
type: custom
used_by: []
location: none
content_type: filesystem
project: default
created_at: 2025-04-30T01:52:52.264481907Z

```

For clarity we will modify the description field to reflect its purpose

```
# incus storage volume edit shared-images global-volume --> add the description
```

Please note the syntax. We need to supply both the name of the volume as well as the storage pool on which it resides

```

# incus storage volume show shared-images global-volume
config: {}
description: Global Volume to store all images, shared across all projects
name: global-volume
type: custom
used_by: []
location: none
content_type: filesystem
project: default
created_at: 2025-04-30T01:52:52.264481907Z

```

Now that we have created the volume, let us assign it to a global key which is primarily meant for the same purpose. i.e, defining the location of the volume that contains images.

```

# incus config edit
config:
  core.https_address: :8443
  storage.images_volume: shared-images/global-volume <-- add this line

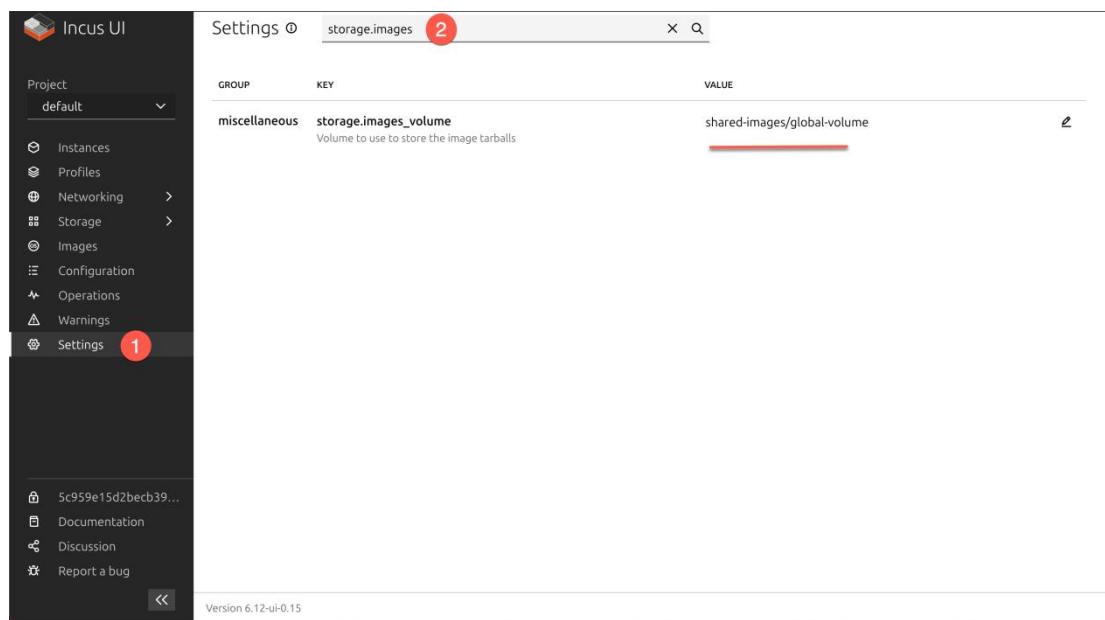
```

```
# incus config show
config:
  core.https_address: :8443
  storage.images_volume: shared-images/global-volume
```

It is also worth knowing the location of this volume in the file system so that we can verify manually where all the images will eventually be stored. And here is that location:

```
# ls -l /var/lib/incus/storage-pools/shared-images/custom/default_global-volume/
```

This setting can also be cross-checked from the web UI by selecting Settings and searching for the key `storage.images_volume`



The screenshot shows the Incus UI settings page. A red circle labeled '1' is on the 'Settings' button in the sidebar. A red circle labeled '2' is on the search bar which contains the text 'storage.images'. The main table shows one entry:

GROUP	KEY	VALUE
miscellaneous	<code>storage.images_volume</code>	shared-images/global-volume

Task #13: Sanity testing the volume

In this task let us do a sanity test on the volume that we created to hold all the images, making sure it is working as expected. This is the location where all the images are expected to be stored. Initially the directory will be empty.

```
# ls -l /var/lib/incus/storage-pools/shared-images/custom/default_global-volume/
```

Let us launch a rockylinux 9 container

```
# incus image copy images:rockylinux/9 local:  
Image copied successfully!
```

```
# incus image list
```

ALIAS	FINGERPRINT	PUBLIC	DESCRIPTION	ARCHITECTURE	TYPE	SIZE	UPLOAD DATE
	e80fa8f0662e	no	Rockylinux 9 amd64 (20250506_02:06)	x86_64	CONTAINER	112.16MiB	2025/05/06 05:50 UTC

We will take a look at the same directory again.

```
# ls -l /var/lib/incus/storage-pools/shared-images/custom/default_global-volume/  
total 114930  
-rw-r--r-- 1 root root 932 May  6 05:50 e80fa8f0662efdc1781749cf3cb5f58a95a745327d125108410b0c7ad61ea299  
-rw-r--r-- 1 root root 117612544 May  6 05:50 e80fa8f0662efdc1781749cf3cb5f58a95a745327d125108410b0c7ad61ea299.rootfs
```

We can see that the directory is now populated with the image. We can also confirm the name of image based on its fingerprint. Let us launch a new container (c1) using the fingerprint of the image

```
# incus launch local:e80fa8f0662e c1  
Launching c1
```

```
# incus list
```

NAME	STATE	IPV4	IPV6	TYPE	SNAPSHOTS
c1	RUNNING	10.1.1.151 (eth0)		CONTAINER	0

```
# incus shell c1  
[root@c1 ~]# ping 8.8.8.8
```

```

PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=115 time=21.5 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=115 time=21.2 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=115 time=19.8 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=115 time=12.6 ms
^C
--- 8.8.8.8 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4007ms
rtt min/avg/max/mdev = 12.647/17.726/21.464/3.855 ms
[root@c1 ~]#
[root@c1 ~]# exit
logout

```

Now that we know the container is working as expected, let us stop it and delete it.

```

# incus stop c1 ; incus delete c1

# incus list
+-----+-----+-----+-----+
| NAME | STATE | IPV4 | IPV6 | TYPE | SNAPSHOTS |
+-----+-----+-----+-----+

```

We will go ahead and remove the image as well

```

# incus image delete e80fa8f0662e

# incus image list
+-----+-----+-----+-----+-----+-----+
| ALIAS | FINGERPRINT | PUBLIC | DESCRIPTION | ARCHITECTURE | TYPE | SIZE | UPLOAD DATE |
+-----+-----+-----+-----+-----+-----+

```

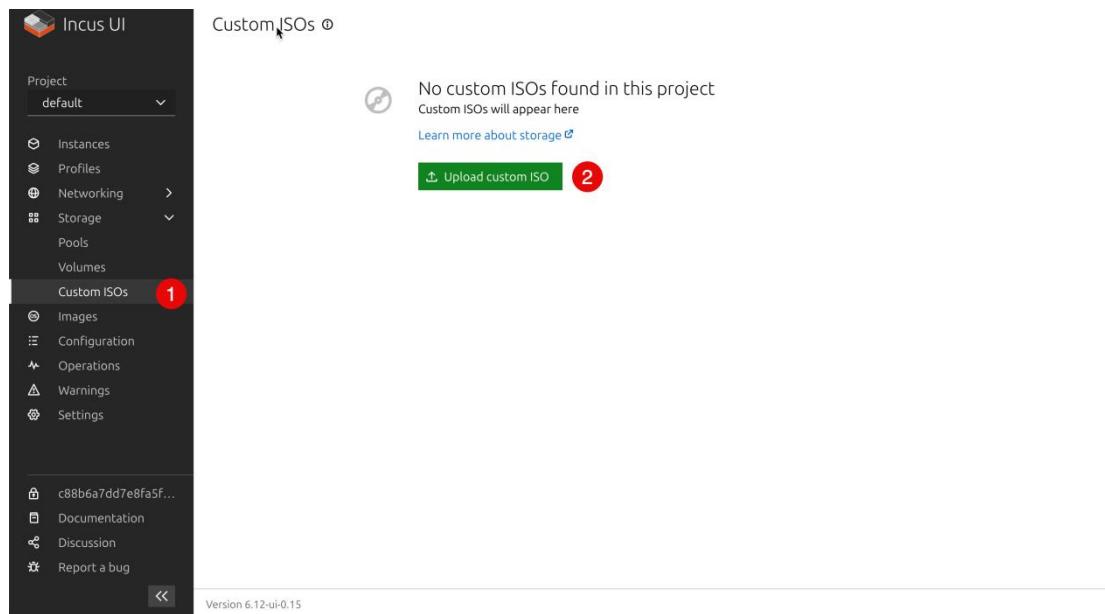
Finally we will take a look at the same directory for one last time and notice that the directory is now empty as we deleted the image.

```
# ls -l /var/lib/incus/storage-pools/shared-images/custom/default_global-volume/
total 0
```

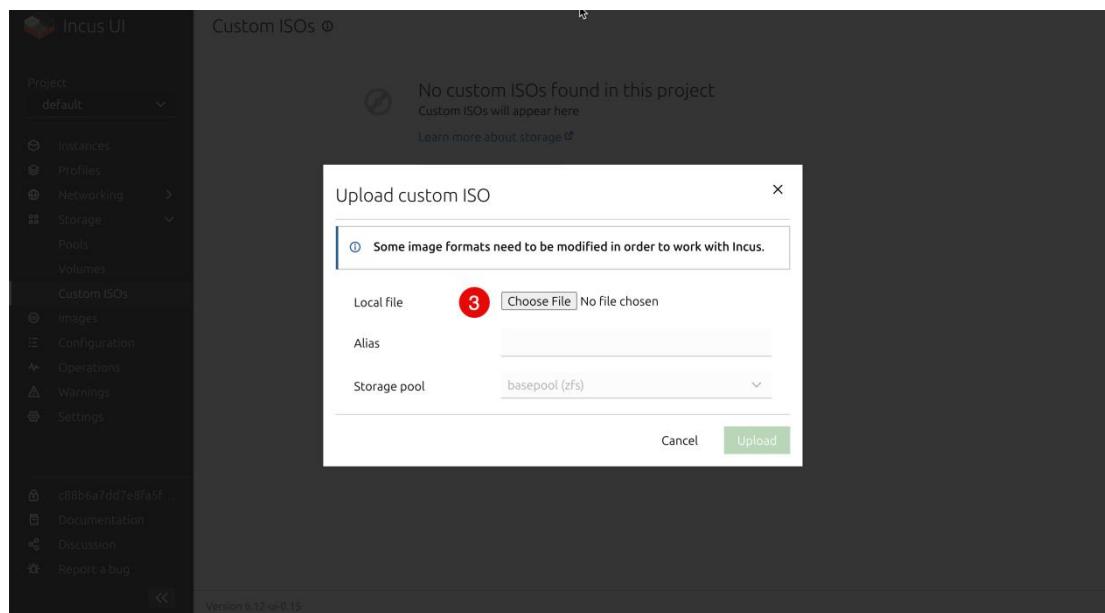
This proves us that the image is stored in the volume (global-volume) that we created before.

Task #14: Uploading an ISO file

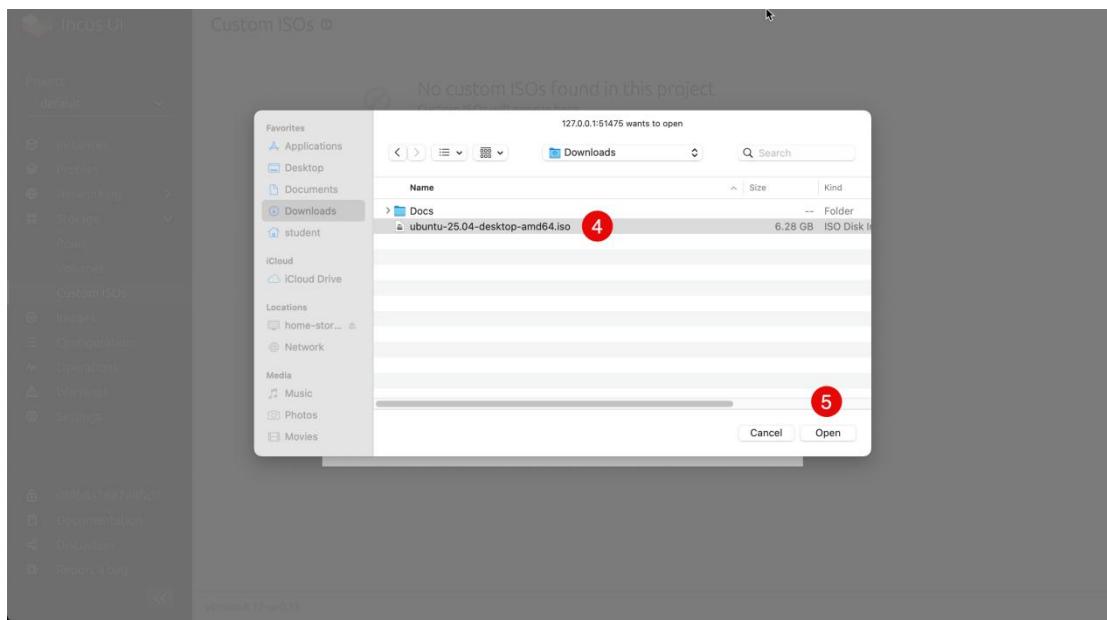
In this task let us look at how to use the incus storage pool ‘shared-iso’ (backed by zfs pool ‘iso-pool’) that we created exclusively to store all the ISO files. Let us upload an ISO file to that pool, which is straight forward using web UI.



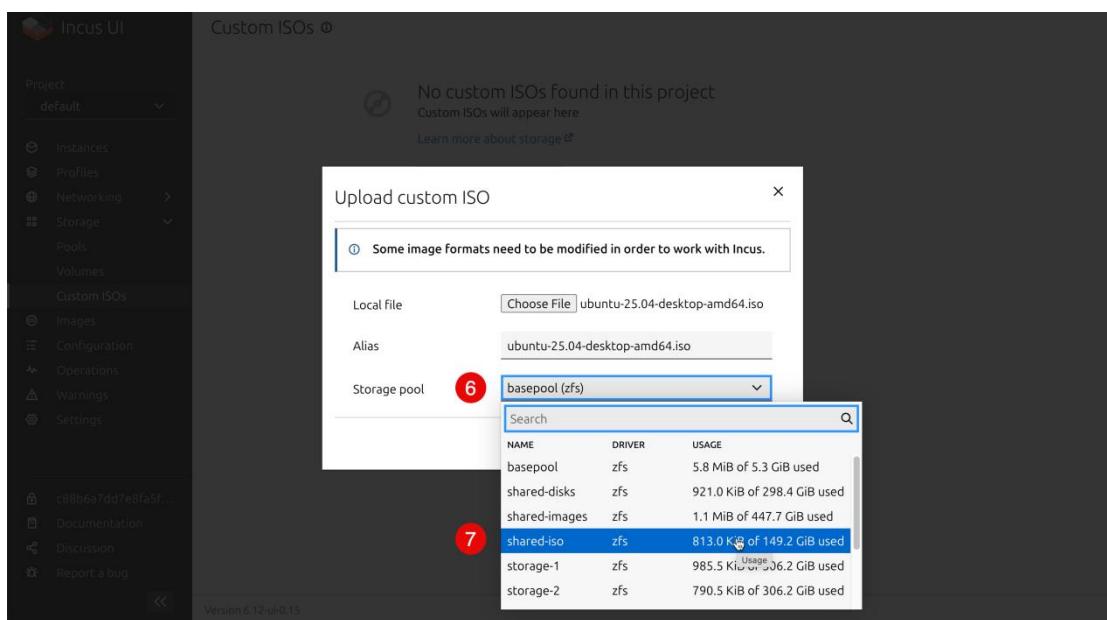
Select the ISO file by clicking on the upload button



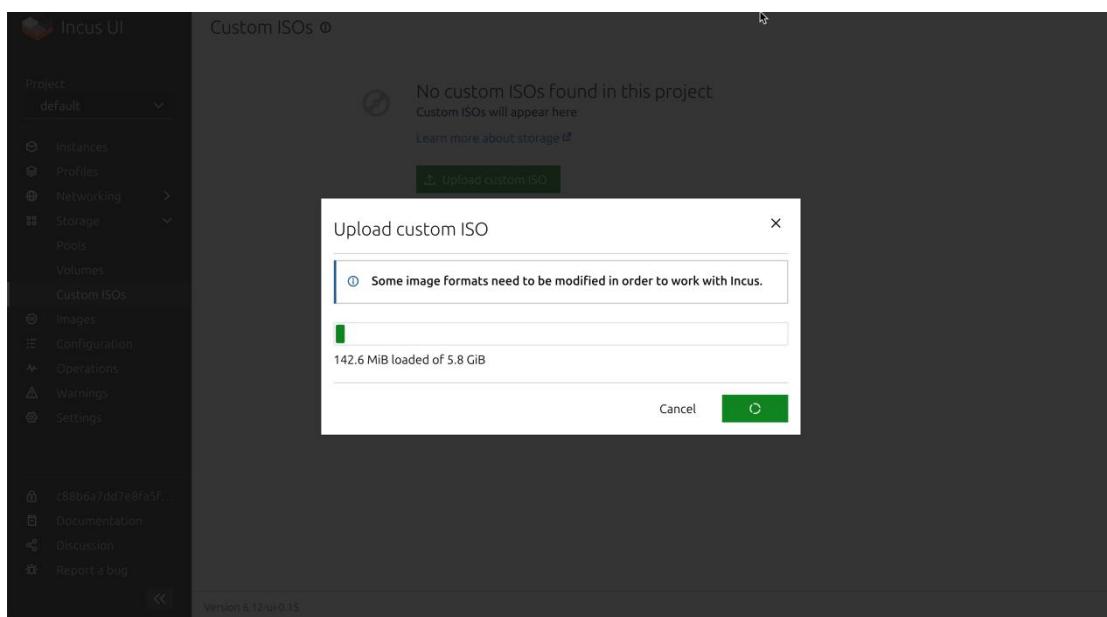
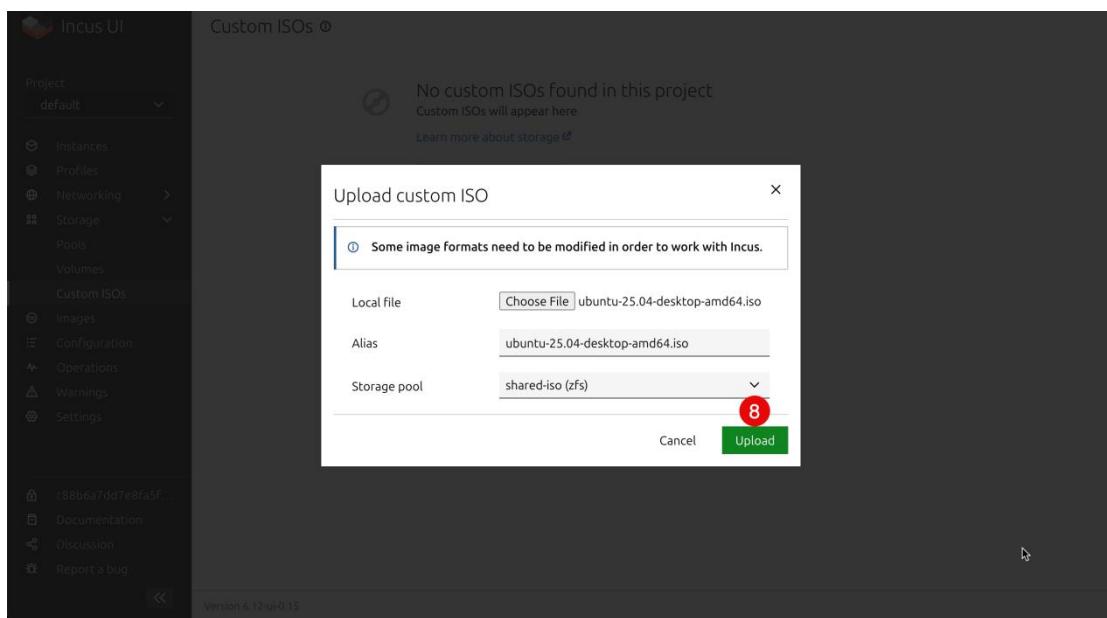
Click Open



Select ‘shared-iso’ from the dropdown menu under Storage pool



Click Upload



After uploading, the iso file will be listed under ‘shared-iso’ storage pool

The screenshot shows the Incus UI interface at the URL 127.0.0.1:51475/ui/project/default/storage/custom-isos. The left sidebar is open, showing the 'Storage' section with 'Custom ISOs' selected. The main content area displays a table titled 'Custom ISOs' with one entry:

NAME	STORAGE POOL	UPLOAD DATE	SIZE	USED BY
ubuntu-25.04-desktop-amd64.iso	shared-iso 9	May 7, 2025, 08:00 PM	5.8 GiB	0

Below the table, there are navigation links: 'c88b6a7dd7e0fa5...', 'Documentation', 'Discussion', and 'Report a bug'. At the bottom right, it says 'Version 6.12-ui-0.15'.

Clicking on 'shared-iso' will show where it is located. ie, on the zfs pool named 'iso-pool'.

The screenshot shows the Incus UI interface at the URL 127.0.0.1:51475/ui/project/default/storage/pools/shared-iso. The left sidebar is open, showing the 'Storage' section with 'Pools' selected. The main content area displays the 'Overview' tab for the 'shared-iso' pool:

General	Name	shared-iso
Status	Created	
Size	5.9 GiB of 149.2 GiB used	
Source	iso-pool 10	
Description	shared ISOs across Projects	
Driver	zfs	

Below the table, under 'Used by (1)', it lists 'Instances (0)', 'Profiles (0)', 'Images (0)', 'Snapshots (0)', and 'Custom volumes (1)'. The 'Custom volumes (1)' row has a link to 'ubuntu-25.04...' with a red circle containing the number '11'.

It is important to understand that the uploaded iso file (ubuntu25.04-desktop-amd64.iso) file will be listed as a custom volume.

Incus UI

Storage volumes / ubuntu-25.04-desktop-a...

Migrate Duplicate Delete

Project: default

- Instances
- Profiles
- Networking >
- Storage >
- Pools
- Volumes
- Custom ISOs
- Images
- Configuration
- Operations
- Warnings
- Settings

c88b6a7dd7e8fa5f...

Documentation

Discussion

Report a bug

Version 6.12-ui-0.15

General

Name	ubuntu-25.04-desktop-amd64.iso
Type	Custom
Content type	ISO
Description	-
Cluster member	-
Pool	shared-iso 12
Date created	May 7, 2025, 08:00 PM
Size	5.9 GiB
Custom config	size 6278520832

Used by (0)

Instances (0)
Profiles (0)
Images (0)
Snapshots (0)
Custom volumes (0)

Incus UI

Volumes **13** Search and filter **Create volume**

Project: default

- Instances
- Profiles
- Networking >
- Storage >
- Pools
- Volumes **13**
- Custom ISOs
- Images
- Configuration
- Operations
- Warnings
- Settings

c88b6a7dd7e8fa5f...

Documentation

Discussion

Report a bug

Version 6.12-ui-0.15

Showing all 2 volumes

NAME	POOL	TYPE	CONTENT TYPE	CREATED AT	SIZE	USED BY	SNAPSHOTS
global-volume	shared-images	Custom	Filesystem	Apr 29, 2025, 06:52 PM	34.0 KiB	0	0
ubuntu-25.04-desktop...	shared-iso	Custom	ISO	May 7, 2025, 08:00 PM	5.9 GiB	0	0

Task #15: ISO file usage

In this task let us work on installing an operating system on a blank virtual machine using the ISO image stored in the storage pool ‘shared-iso’.

To begin with, let us create a blank virtual machine and install Ubuntu25.04 on it with the uploaded ISO image.

There are a couple of ways to accomplish this. One way is to use the web UI. The other way is to use a combination of command lines as well as a software named ‘virt-viewer’. Both of them can be used from the Mac laptop as we have already taken care of its authentication in task #09.

We have seen using WeUI before. So this time, let us take a look at incus commands and using ‘virt-viewer’.

‘virt-viewer’ is a separate stand alone free software which is a pre-requisite that needs to be installed on (my) Mac laptop. This software is used to access the console of the virtual machine on which we are planning to install ubuntu25.04 on it

We will begin with using incus command lines. Once the virtual machine is started, we will then go ahead continuing the rest of the installation using virt-viewer.

Let us keep the installation screenshots as minimal as we can or else this task will end up becoming nothing more than a bunch of screenshots.

All these commands are issued on the Mac Laptop.

First step is to create a blank virtual machine with 16GB of RAM, 2 x CPU’s, and its root disk of size 10GiB

```
mac% incus init vm-ubuntu25 --empty -c limits.memory=16GiB -c limits.cpu=2 -d root,size=10GiB  
--vm  
  
% incus list  
+-----+-----+-----+-----+-----+  
| NAME | STATE | IPV4 | IPV6 | TYPE | SNAPSHOTS |  
+-----+-----+-----+-----+-----+  
| vm-ubuntu25 | STOPPED | | VIRTUAL-MACHINE | 0 |  
+-----+-----+-----+-----+
```

Let us configure the newly created virtual machine named vm-ubuntu25 with the ISO image

```
mac% incus config device add vm-ubuntu25 iso disk pool=shared-iso  
source=ubuntu-25.04-desktop-amd64.iso boot.priority=10
```

The above command line in plain English:

- 1) Configure the virtual machine by attaching the ISO image located in the storage pool ‘shared-iso’ as a new device
- 2) Name of the ISO file being ubuntu-25.04-desktop-amd64.iso
- 3) Set the boot priority to 10
- 4) boot priority is about setting the boot order when powered on. It ensures the virtual machine will boot using the ISO image first

Let us go ahead and start the virtual machine

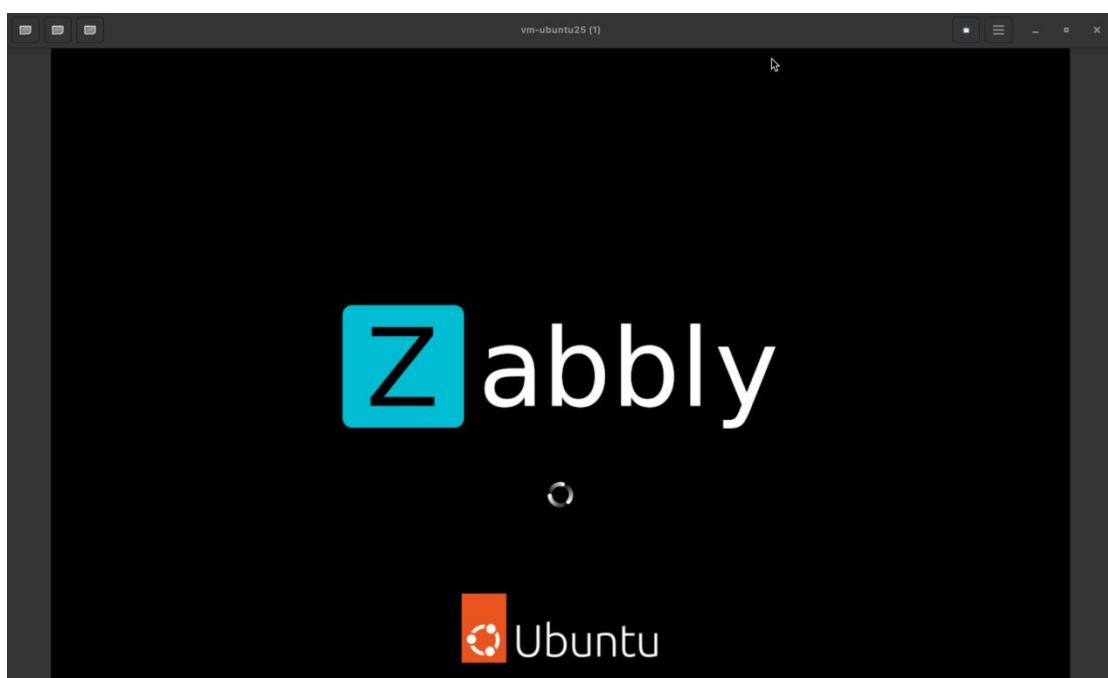
```
mac% incus start vm-ubuntu25
```

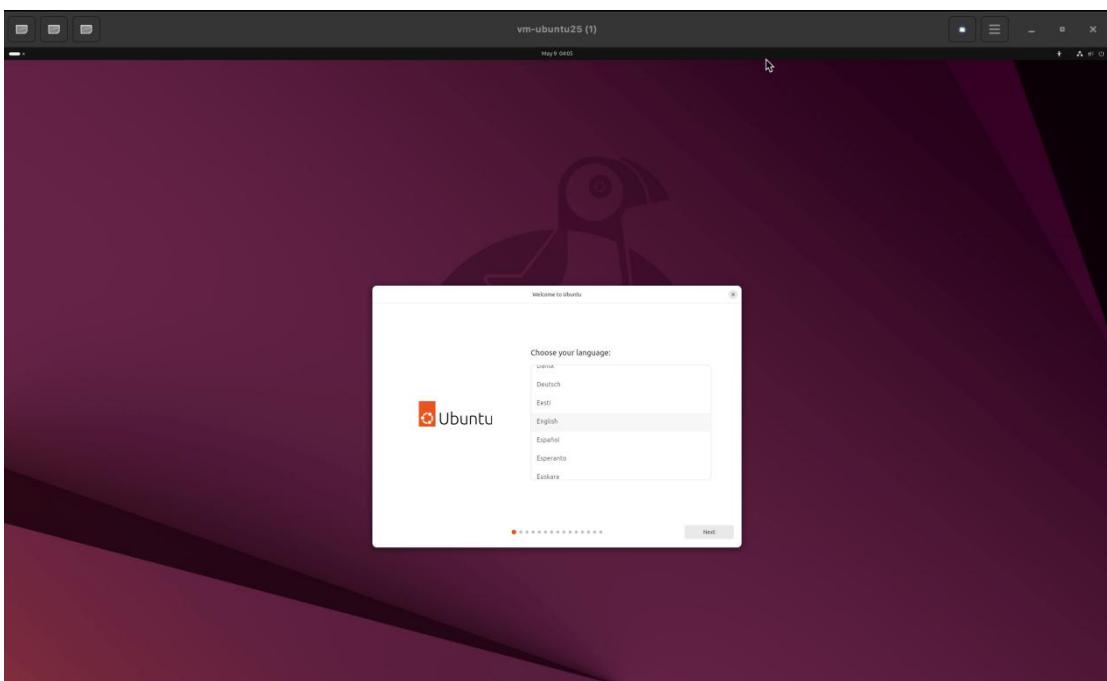
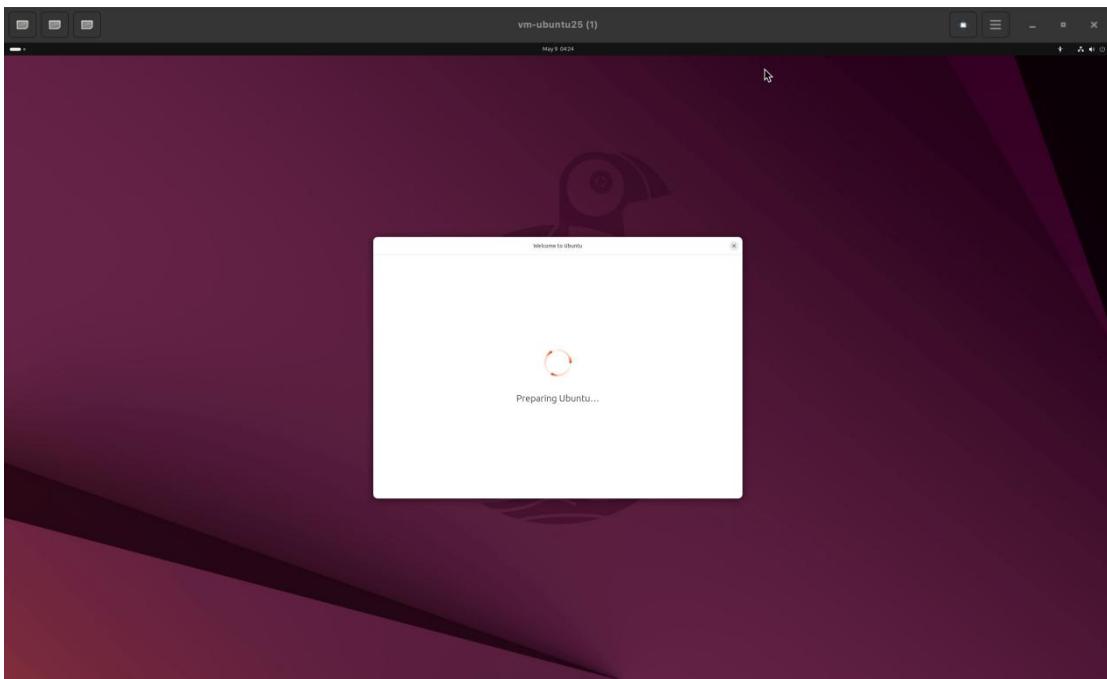
Now that the virtual machine is started, we need to access its console from our Mac Laptop. And here is where the role of virt-viewer comes in. Running the command line below will automatically invoke virt-viewer and gets us the console of the virtual machine (we don’t have to run the virt-viewer by ourselves manually)

```
mac% incus console vm-ubuntu25 --type=vga
```

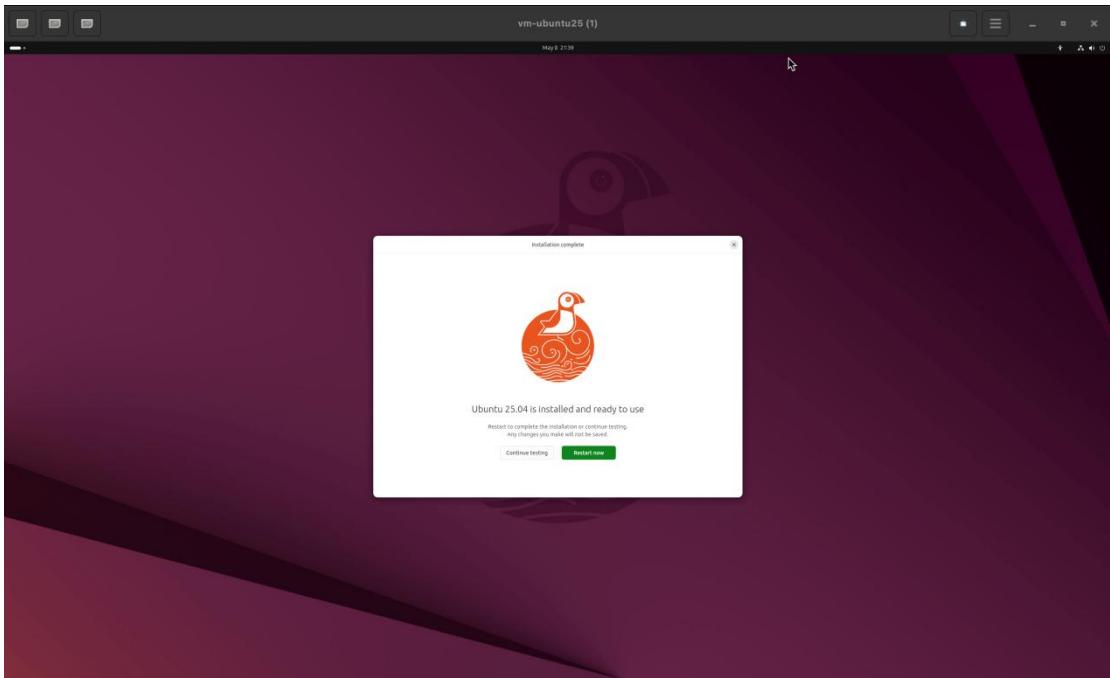
You might see warning messages thrown on the terminal but you will also see virt-viewer instantiated on your laptop.

To begin with you will see the splash screen with the company logo of Zabbly (maintainer of project Incus)

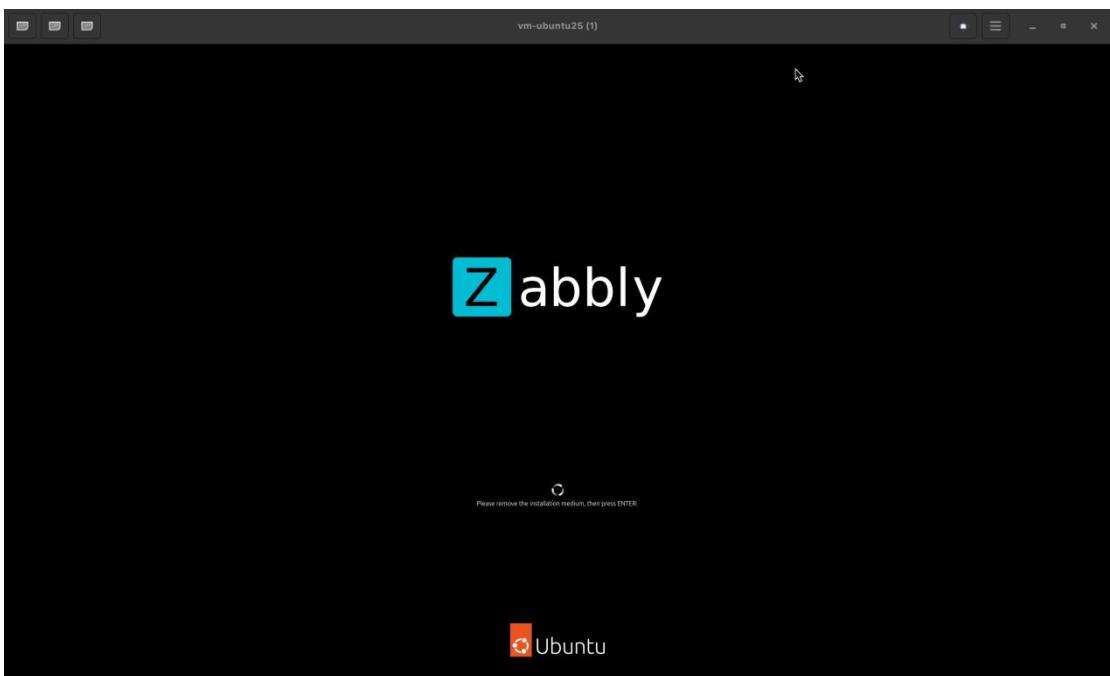




The rest of the screenshots are well known. so let us skip all that and go towards the end



On clicking on the Reboot now button, the virtual machines power cycles and reaches this spot



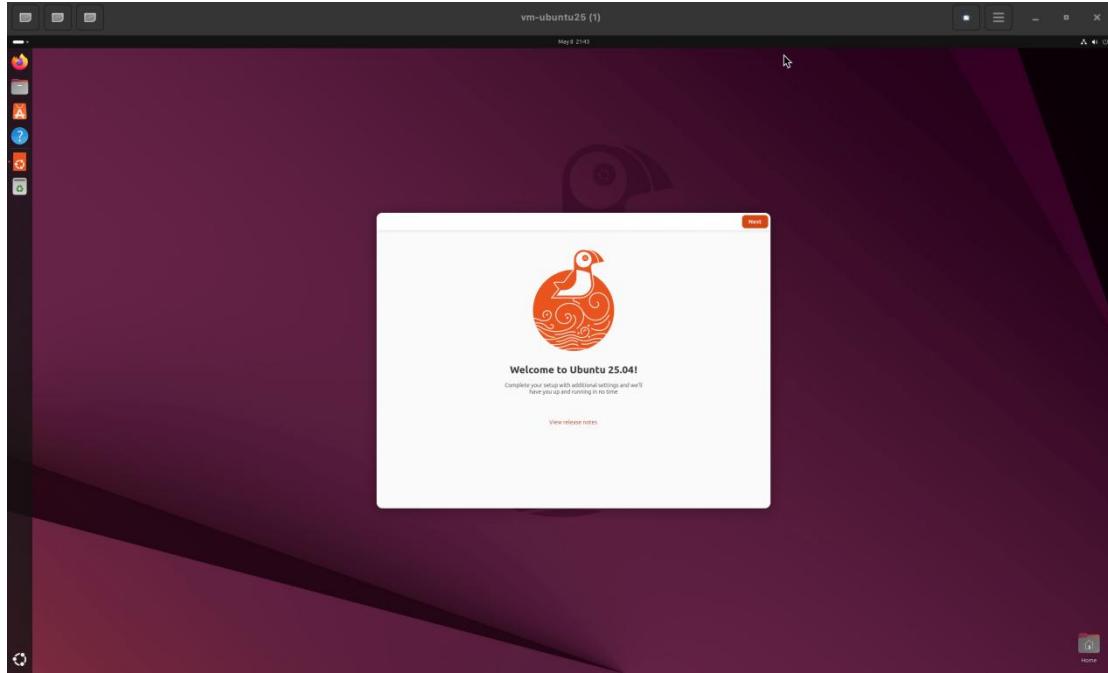
You will notice the message saying “Please remove the installation medium, then press ENTER.

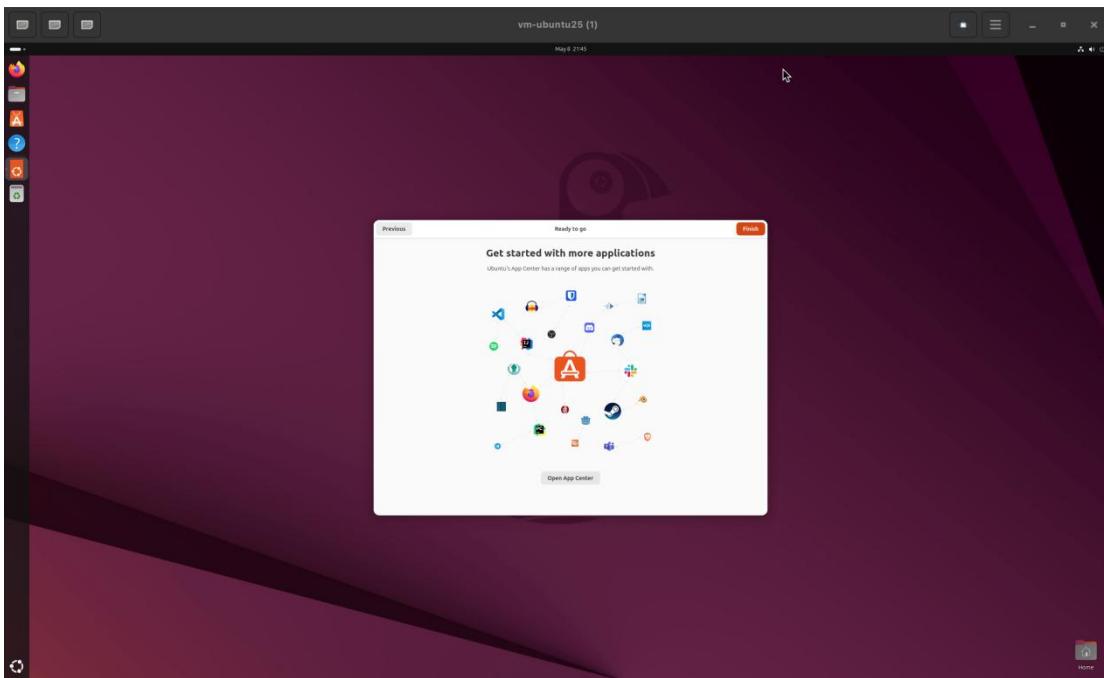
And before pressing the ENTER, run the following incus commands on the Mac laptop

```
mac% incus config device remove vm-ubuntu25 iso
```

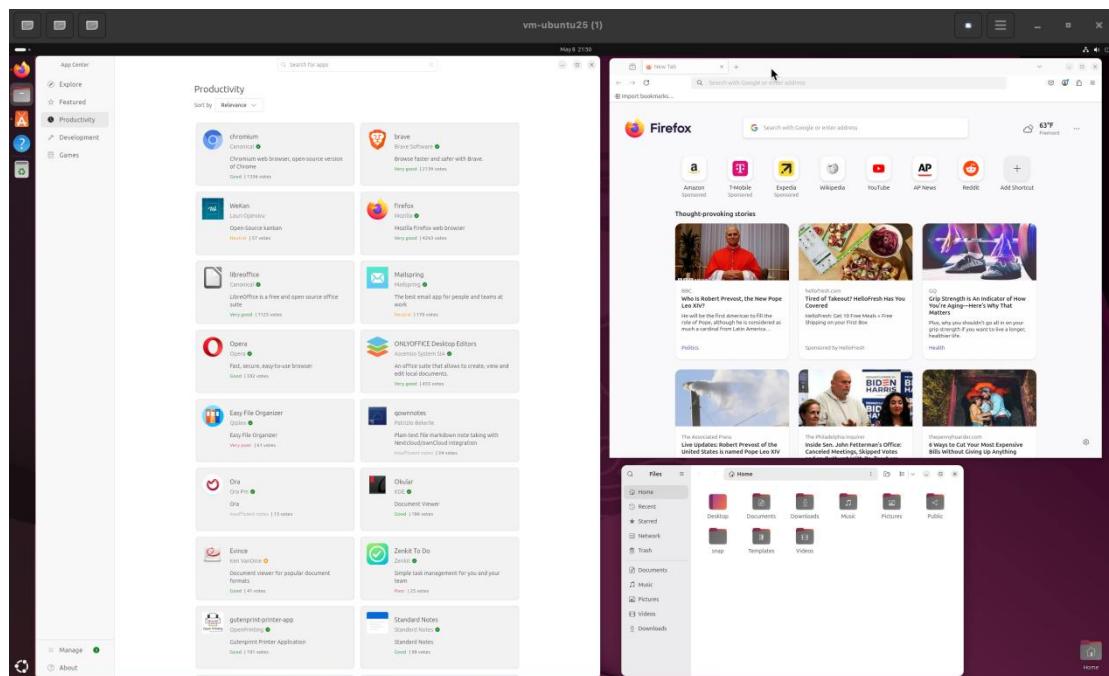
```
mac% incus start vm-ubuntu25
```

You will now see the virtual machine getting power cycled boots from its root disk

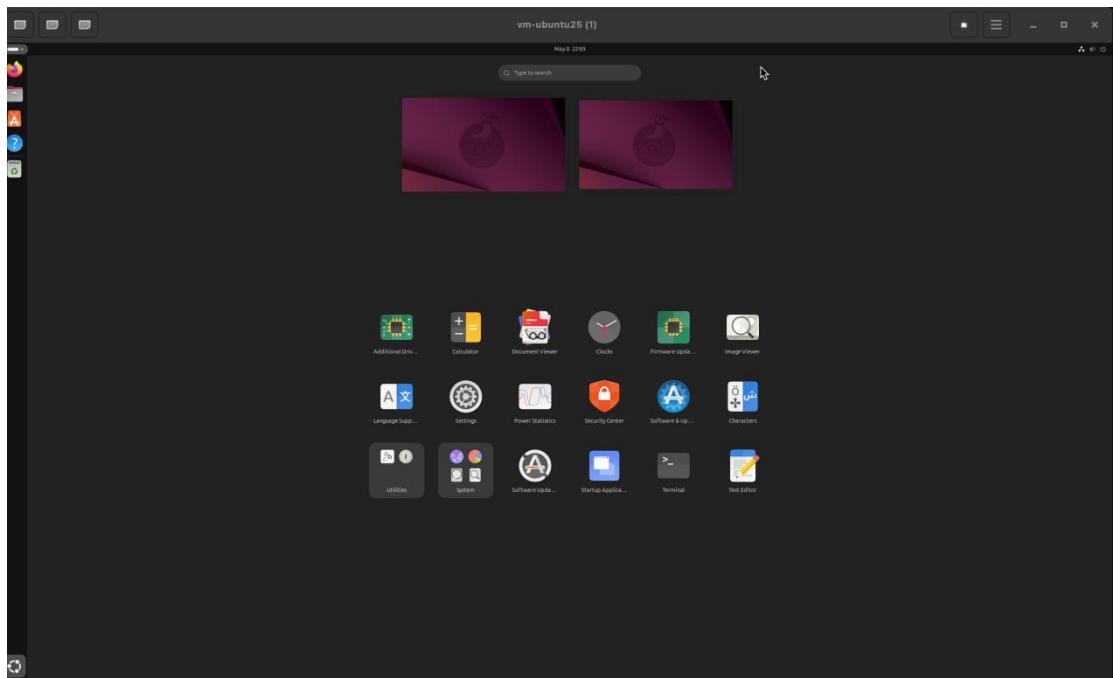




The App store, Firefox browser and File Manager



Desktop



Now that we have access the console of the virual machine via virt-viewer, let us take a quick look at the same virtual machine using web UI

A screenshot of the Incus UI web interface, accessible at 127.0.0.1:54863/ui/project/default/instances. The left sidebar shows a navigation tree with "Project: default" selected, followed by "Instances", "Profiles", "Networking", "Storage", "Images", "Configuration", "Operations", "Warnings", and "Settings". Below this are links for "c88b6a7dd7e8fa5f...", "Documentation", "Discussion", and "Report a bug". The main content area shows a table titled "Instances" with one entry: "vm-ubuntu25" (Type: VM, Description: "", IPv4: 10.1.1.240, Status: Running). There are buttons for Start, Restart, Freeze, Stop, and Delete. The table includes a header row with columns for NAME, TYPE, DESCRIPTION, IPV4, and STATUS, and a footer indicating 1 instance is selected and a page size of 50 per page.

127.0.0.1:54863/ui/project/default/instance/vm-ubuntu25

Incus UI

Instances / vm-ubuntu25

Running | Create Image | Duplicate | Export | Delete

Project: default

Instances

Profiles

Networking

Storage

Pools

Volumes

Custom ISOs

Images

Configuration

Operations

Warnings

Settings

c88b6a7dd7e8fa5f...

Documentation

Discussion

Report a bug

Overview Configuration Snapshots Terminal Console Logs

Type: VM
IPv4: 10.1.1.240 (eth0)
IPv6:
MAC address: 10:66:6a:c8:76:a9 (eth0)
Architecture: x86_64
Location:
PID: 11732
Date created: May 8, 2025, 09:21 PM
Last used: May 8, 2025, 10:08 PM

Preview

Version 6.12-ui-0.15

127.0.0.1:54863/ui/project/default/instance/vm-ubuntu25/console

Incus UI

Instances / vm-ubuntu25

Running | Create Image | Duplicate | Export | Delete

Project: default

Instances

Profiles

Networking

Storage

Pools

Volumes

Custom ISOs

Images

Configuration

Operations

Warnings

Settings

c88b6a7dd7e8fa5f...

Documentation

Discussion

Report a bug

Console

Logs

Graphic (selected) Text console

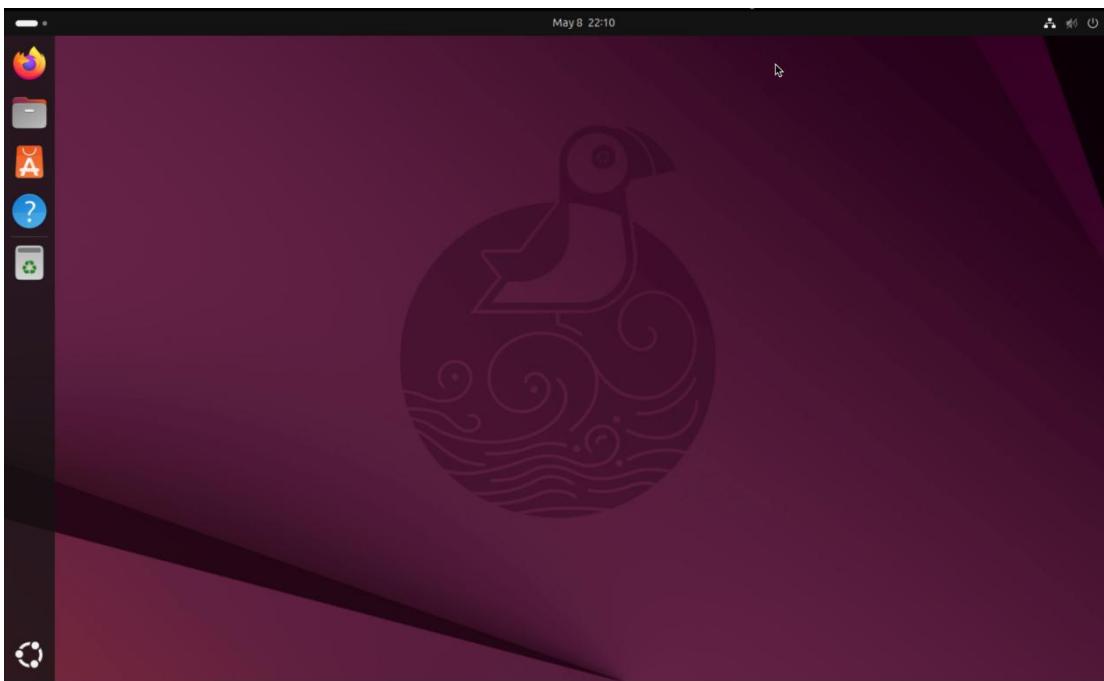
Attach ISO Fullscreen Shortcuts

May 8 22:09

Home

1 operation in progress...

Version 6.12-ui-0.15



Invenio UI

Instances / vm-ubuntu25

Running | | | |

+ Create Image | | |

Project: default

Instances

Profiles

Networking >

Storage

Pools

Volumes

Custom ISOs

Images

Configuration

Operations

Warnings

Settings

c88b6a7dd7e8fa5f...

Documentation

Discussion

Report a bug

Overview Configuration Snapshots Terminal Console Logs

Usage

	CPU Time(s)	
Memory	0 B of 0 B memory used	
Disk	4.0 GiB of 10.0 GiB disk used	

Networks

NAME	INTERFACE	TYPE	MANAGED
incusbr0	eth0	bridge	Yes

Devices

NAME	TYPE	DETAILS
root	disk (root)	Pool: basepool

Profiles

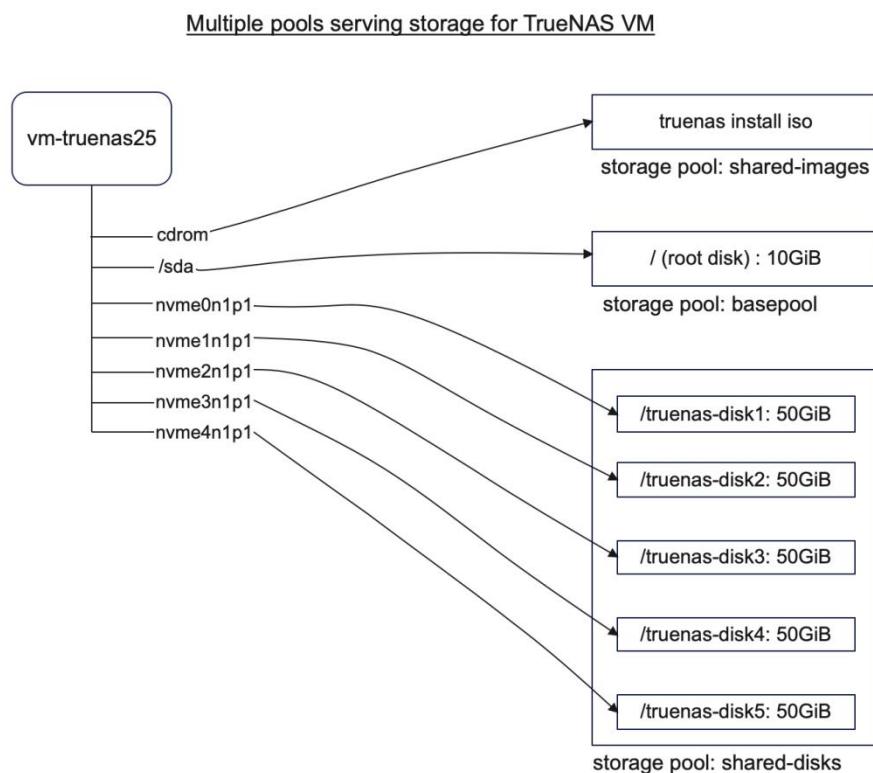
NAME	DESCRIPTION
default	Default Invenio profile

Version 6.12-ui-0.15

Task #16: Using custom storage pools and volumes

To understand how to use multiple storage pools, let us create a single virtual machine that spans across multiple storage pools. For this purpose, let us create a blank virtual machine with its root disk stored in one storage pool (basepool), all its data disks gets stored in a second storage pool (shared-disks), and finally installing TrueNAS on it by attaching an ISO file to it, which will be stored in the third storage pool (shared-iso).

This task is based on [a video posted by Stephane Graber](#) (incus lead) while evaluating TrueNAS Community Edition.



To begin with, let us import TrueNAS iso file to the storage pool that we specifically created to store all iso files (ie, shared-iso).

These commands can be carried out from the Mac Laptop

```
mac% incus storage volume import shared-iso ~/Downloads/TrueNAS-SCALE-25.04.0.iso  
truenas-install --type=iso
```

Let us check make sure it did land in the storage pool 'shared-iso'. You can see it below listed as truenas-install (under 'used_by')

```

mac% incus storage show shared-iso
config:
  source: iso-pool
  volatile.initial_source: iso-pool
  zfs.pool_name: iso-pool
description: shared ISOs accross Projects
name: shared-iso
driver: zfs
used_by:
- /1.0/storage-pools/shared-iso/volumes/custom/truenas-install
- /1.0/storage-pools/shared-iso/volumes/custom/ubuntu-25.04-desktop-amd64.iso
status: Created
locations:
- none

```

Create an empty virtual machine named vm-truenas25

```

mac% incus create vm-truenas25 --vm --empty -c limits.cpu=4 -c limits.memory=8GiB -c
security.secureboot=false

```

```

mac% incus list
+-----+-----+-----+-----+-----+
|   NAME    | STATE | IPV4 | IPV6 |      TYPE      | SNAPSHOTS |
+-----+-----+-----+-----+-----+
| vm-truenas25 | STOPPED |     |     | VIRTUAL-MACHINE | 0          |
+-----+-----+-----+-----+-----+

```

The root disk of this newly created virtual machine will be stored in storage pool ‘basepool’

```

mac% incus storage show basepool
config:
  source: basepool
  volatile.initial_source: basepool
  zfs.pool_name: basepool
description: passed to incus admin init
name: basepool
driver: zfs
used_by:
- /1.0/instances/vm-truenas25
- /1.0/profiles/default
status: Created
locations:
- none

```

Let us create 5 volumes (of block type and size 50GiB each) on the incus storage pool 'shared-disks'. These will be attached as data disks to the newly created virtual machine. Later we use all of them to create a RAIDZ1 volume combining all of them.

```
mac% incus storage show shared-disks
config:
  source: disks-pool
  volatile.initial_source: disks-pool
  zfs.pool_name: disks-pool
description: shared disks across Projects
name: shared-disks
driver: zfs
used_by: []
status: Created
locations:
- none

mac% incus storage volume create shared-disks truenas-disk1 size=50GiB --type=block

mac% incus storage volume create shared-disks truenas-disk2 size=50GiB --type=block

mac% incus storage volume create shared-disks truenas-disk3 size=50GiB --type=block

mac% incus storage volume create shared-disks truenas-disk4 size=50GiB --type=block

mac% incus storage volume create shared-disks truenas-disk5 size=50GiB --type=block

mac% incus storage show shared-disks
config:
  source: disks-pool
  volatile.initial_source: disks-pool
  zfs.pool_name: disks-pool
description: shared disks across Projects
name: shared-disks
driver: zfs
used_by:
- /1.0/storage-pools/shared-disks/volumes/custom/truenas-disk1
- /1.0/storage-pools/shared-disks/volumes/custom/truenas-disk2
- /1.0/storage-pools/shared-disks/volumes/custom/truenas-disk3
- /1.0/storage-pools/shared-disks/volumes/custom/truenas-disk4
- /1.0/storage-pools/shared-disks/volumes/custom/truenas-disk5
status: Created
locations:
- none
```

Let us configure the virtual machine vm-truenas25 with all the volumes that we created so far, starting with volume created out of the tuenas iso file that resides in storage pool 'shared-iso'

```
mac% incus config device add vm-truenas25 install disk pool=shared-iso  
source=truenas-install boot.priority=10
```

Now let us add all the 5 disks that reside in storage pool 'shared-disks'

```
mac% incus config device add vm-truenas25 disk1 disk pool=shared-disks source=truenas-disk1  
io.bus=nvme
```

```
mac% incus config device add vm-truenas25 disk2 disk pool=shared-disks source=truenas-disk2  
io.bus=nvme
```

```
mac% incus config device add vm-truenas25 disk3 disk pool=shared-disks source=truenas-disk3  
io.bus=nvme
```

```
mac% incus config device add vm-truenas25 disk4 disk pool=shared-disks source=truenas-disk4  
io.bus=nvme
```

```
mac% incus config device add vm-truenas25 disk5 disk pool=shared-disks source=truenas-disk5  
io.bus=nvme
```

```
mac% incus config show --expanded vm-truenas25  
architecture: x86_64  
config:  
    limits.cpu: "4"  
    limits.memory: 8GiB  
    security.secureboot: "false"  
    volatile.apply_template: create  
    volatile.cloud-init.instance-id: 298a1379-bb19-4687-b96d-c782a5bfaf7  
    volatile.eth0.hwaddr: 10:66:6a:bd:0a:cb  
    volatile.uuid: 4f4476c1-5c37-4054-a498-7969cc3ea8f1  
    volatile.uuid.generation: 4f4476c1-5c37-4054-a498-7969cc3ea8f1  
devices:  
    disk1:  
        io.bus: nvme  
        pool: shared-disks  
        source: truenas-disk1  
        type: disk  
    disk2:  
        io.bus: nvme  
        pool: shared-disks
```

```

source: truenas-disk2
type: disk
disk3:
  io.bus: nvme
  pool: shared-disks
  source: truenas-disk3
  type: disk
disk4:
  io.bus: nvme
  pool: shared-disks
  source: truenas-disk4
  type: disk
disk5:
  io.bus: nvme
  pool: shared-disks
  source: truenas-disk5
  type: disk
eth0:
  name: eth0
  network: incusbr0
  type: nic
install:
  boot.priority: "10"
  pool: shared-iso
  source: truenas-install
  type: disk
root:
  path: /
  pool: basepool
  type: disk
ephemeral: false
profiles:
- default
stateful: false
description: ""

mac% incus list
+-----+-----+-----+-----+-----+
| NAME | STATE | IPV4 | IPV6 | TYPE | SNAPSHOTS |
+-----+-----+-----+-----+-----+
| vm-truenas25 | RUNNING | 10.1.1.115 (enp5s0) | | VIRTUAL-MACHINE | 0 |
+-----+-----+-----+-----+-----+

```

Now that the virtual machine is up and running, let us access its shell and run a few commands

as a sanity check. And also verify the networking setup by installing traceroute command and reach out to zabbly.com

```
mac% incus shell vm-truenas25
root@truenas[~]# uname -a
Linux truenas 6.12.15-production+truenas #1 SMP PREEMPT_DYNAMIC Tue Apr 15 20:07:13 UTC 2025
x86_64 GNU/Linux
root@truenas[~]#
root@truenas[~]# cat /etc/os-release
PRETTY_NAME="Debian GNU/Linux 12 (bookworm)"
NAME="Debian GNU/Linux"
VERSION_ID="12"
VERSION="12 (bookworm)"
VERSION_CODENAME=bookworm
ID=debian
HOME_URL="https://www.debian.org/"
SUPPORT_URL="https://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"

root@truenas[~]# lsblk
NAME   MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda     8:0    0  10G  0 disk
`-sda1  8:1    0   1M  0 part
`-sda2  8:2    0 512M  0 part
`-sda3  8:3    0  9.5G  0 part
nvme0n1 259:0  0   50G  0 disk
nvme1n1 259:1  0   50G  0 disk
nvme2n1 259:2  0   50G  0 disk
nvme3n1 259:3  0   50G  0 disk
nvme4n1 259:4  0   50G  0 disk
root@truenas[~]#
root@truenas[~]# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
        inet6 ::1/128 scope host noprefixroute
            valid_lft forever preferred_lft forever
2: enp5s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 10:66:6a:bd:0a:cb brd ff:ff:ff:ff:ff:ff
        inet 10.1.1.115/24 brd 10.1.1.255 scope global dynamic enp5s0
            valid_lft 2859sec preferred_lft 2859sec
        inet6 fe80::1266:6aff:febd:acb/64 scope link
```

```

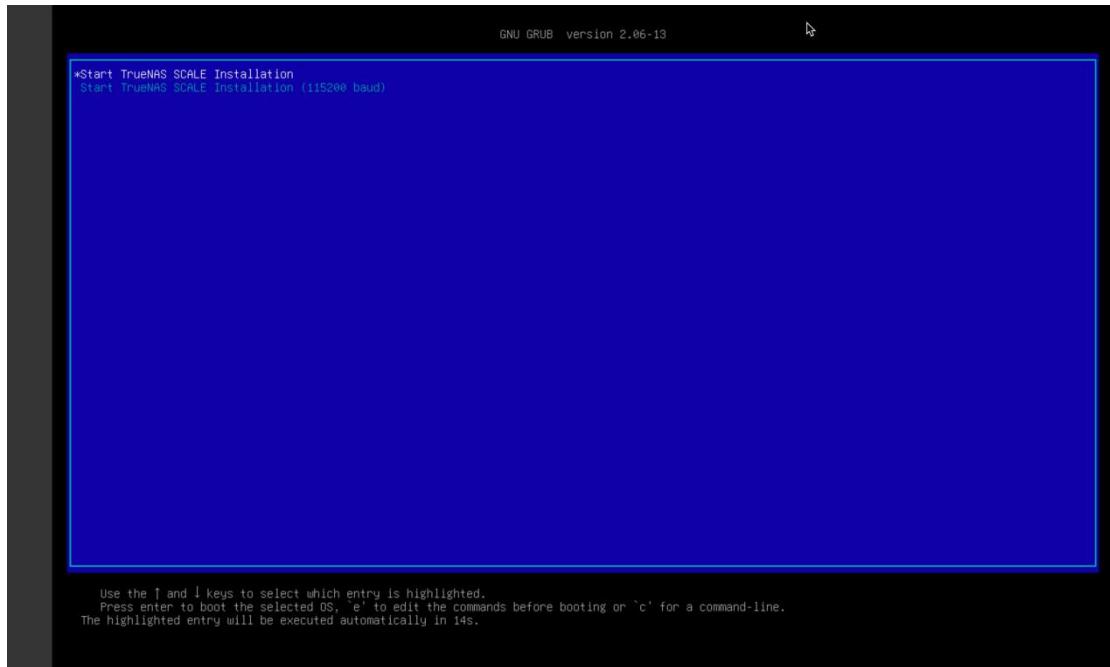
valid_lft forever preferred_lft forever

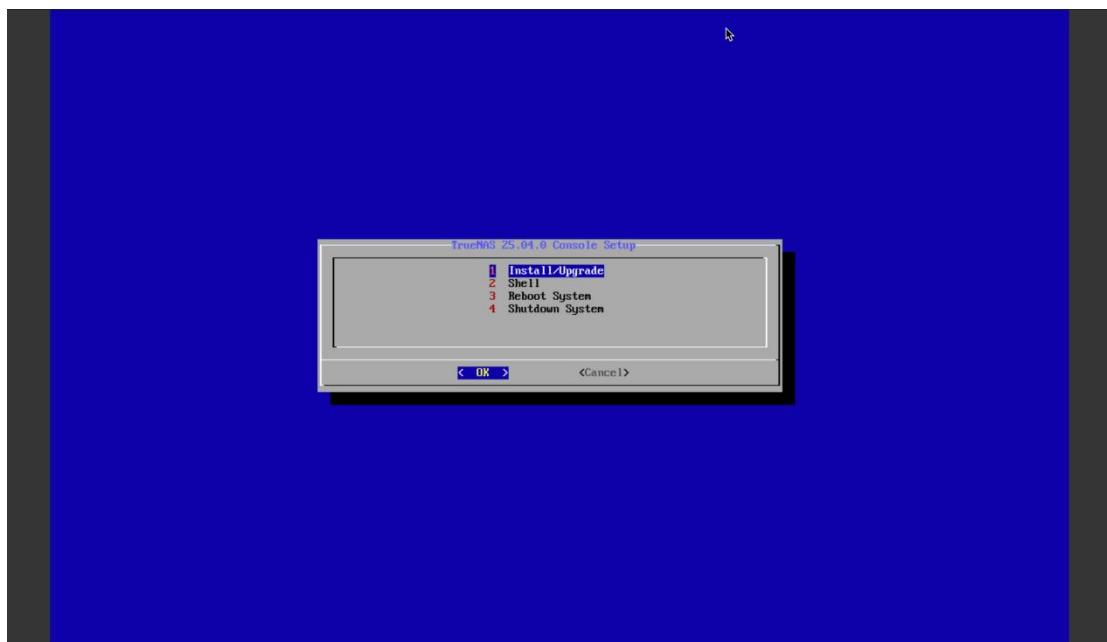
root@truenas[~]# ping zabbly.com
PING zabbly.com (45.45.148.7) 56(84) bytes of data.
64 bytes from rproxy.dcmtl.stgraber.org (45.45.148.7): icmp_seq=1 ttl=49 time=84.4 ms
64 bytes from rproxy.dcmtl.stgraber.org (45.45.148.7): icmp_seq=2 ttl=49 time=85.6 ms
64 bytes from rproxy.dcmtl.stgraber.org (45.45.148.7): icmp_seq=3 ttl=49 time=80.4 ms
^C
--- zabbly.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2004ms
rtt min/avg/max/mdev = 80.444/83.474/85.604/2.200 ms
root@truenas[~]# exit

```

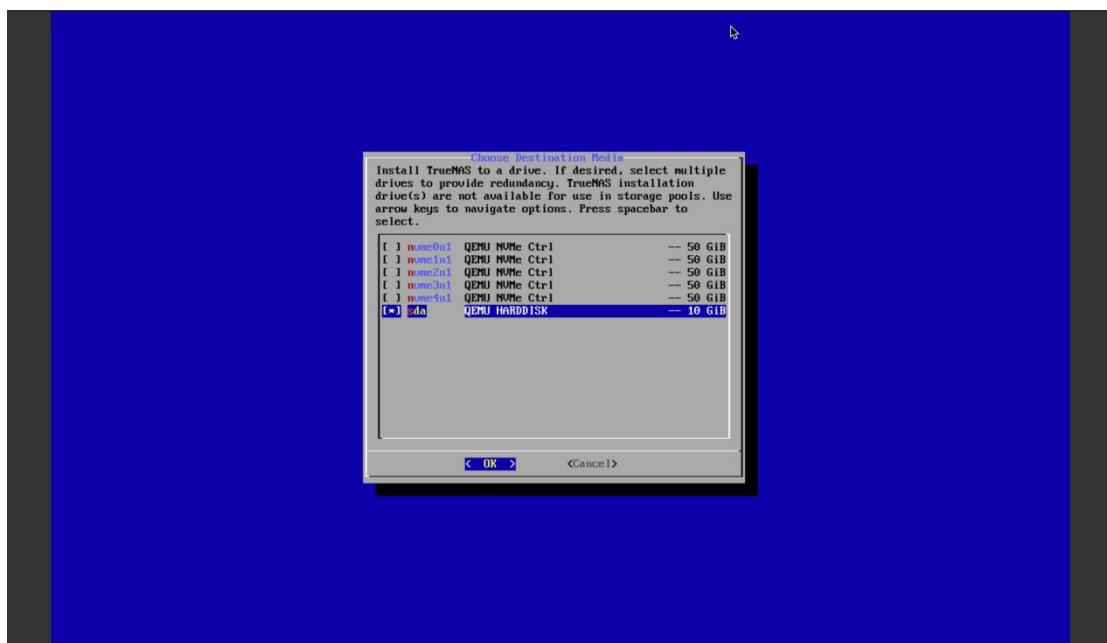
let us connect to the console of the virtual machine vm-truenas25 and continue with the installation of TrueNAS Community Edition.

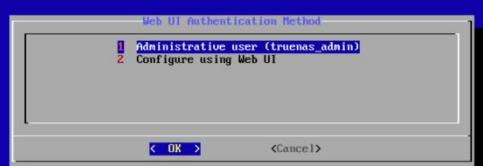
```
mac% incus start vm-truenas25 --console=vga
```

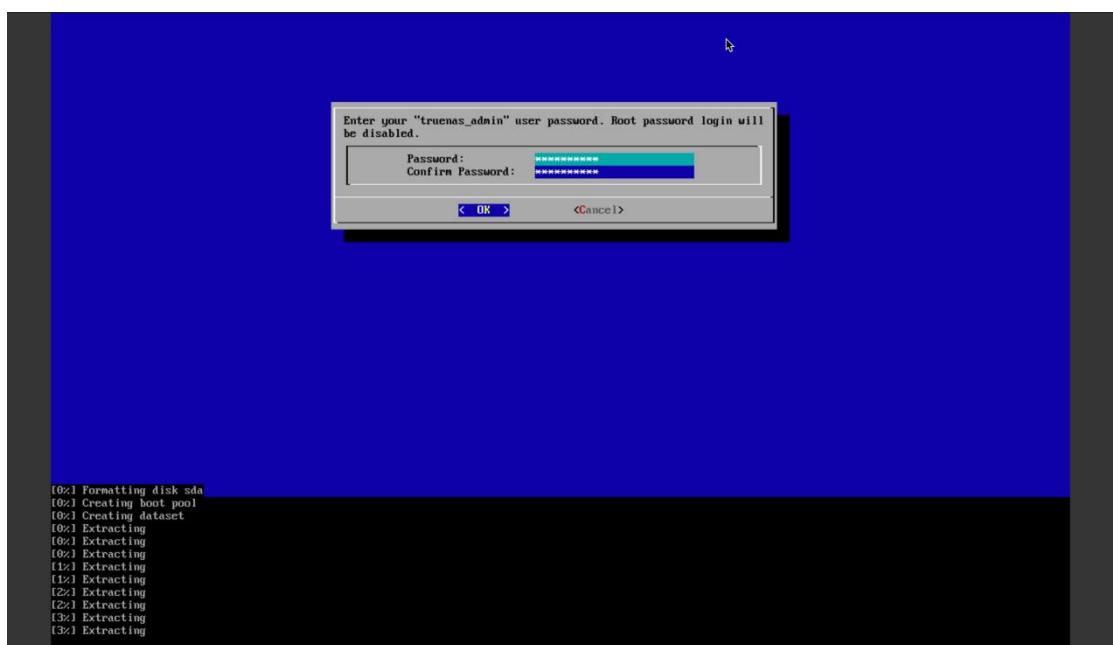
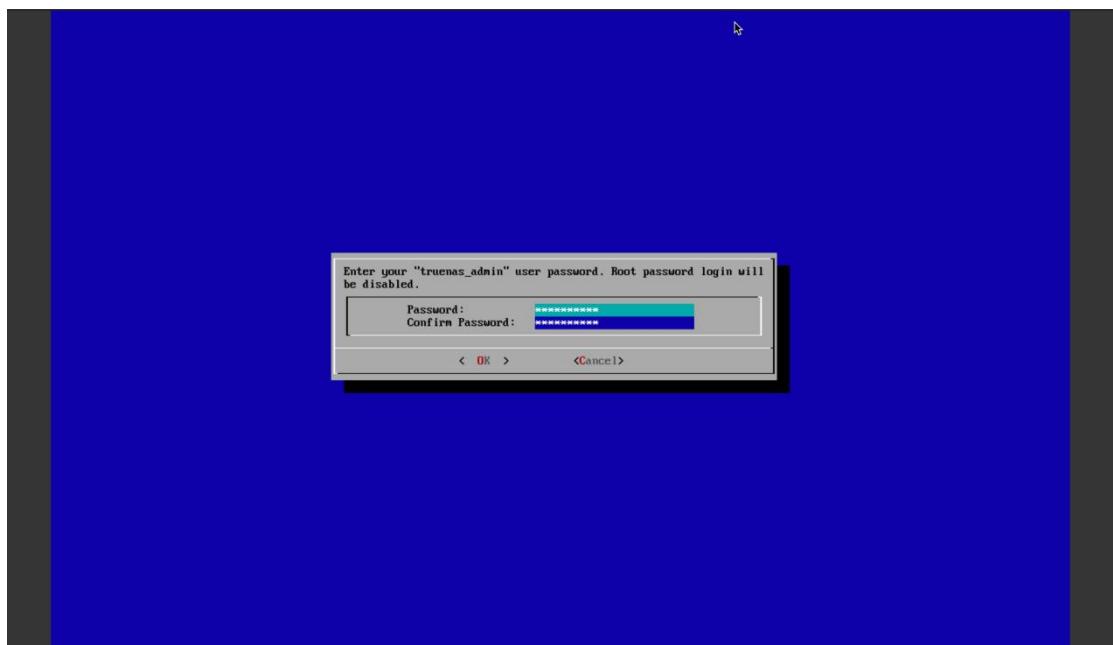


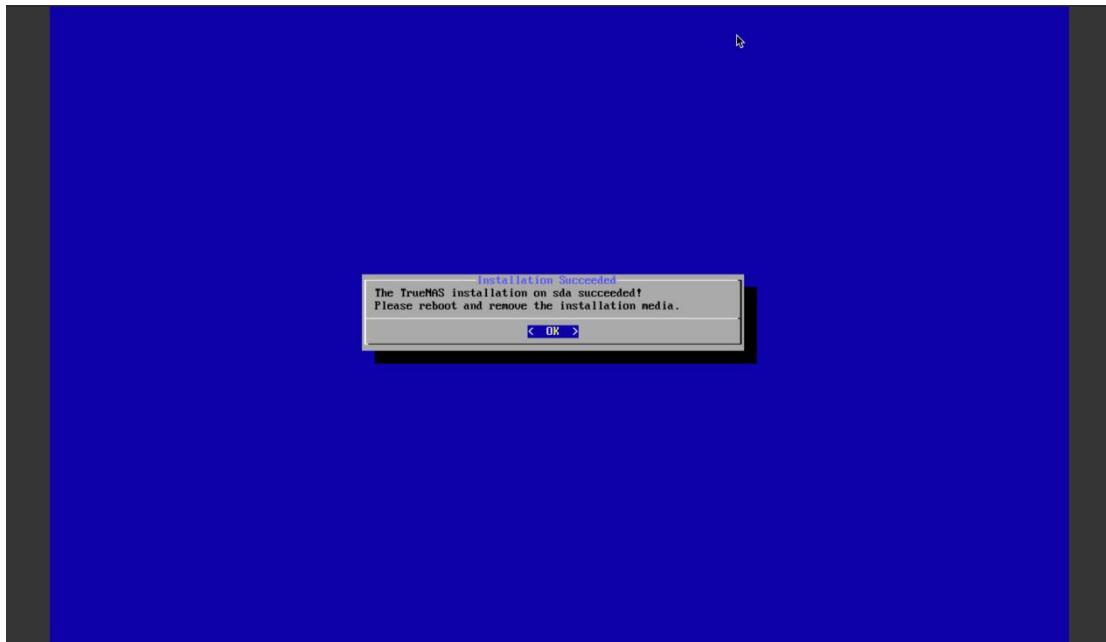
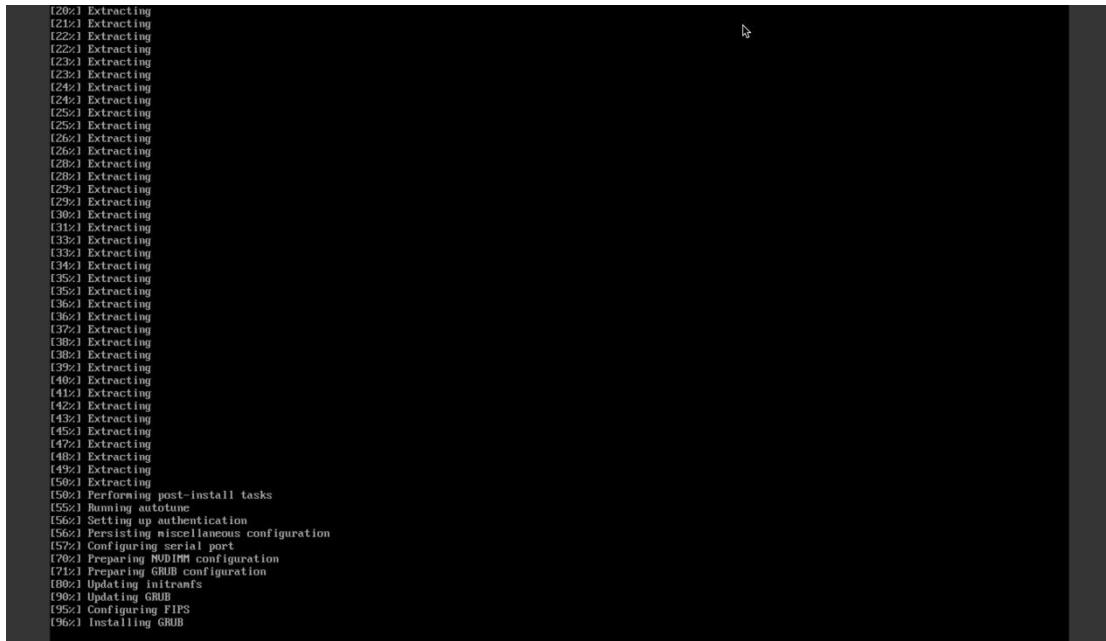


Let us select 'sda' as the root disk



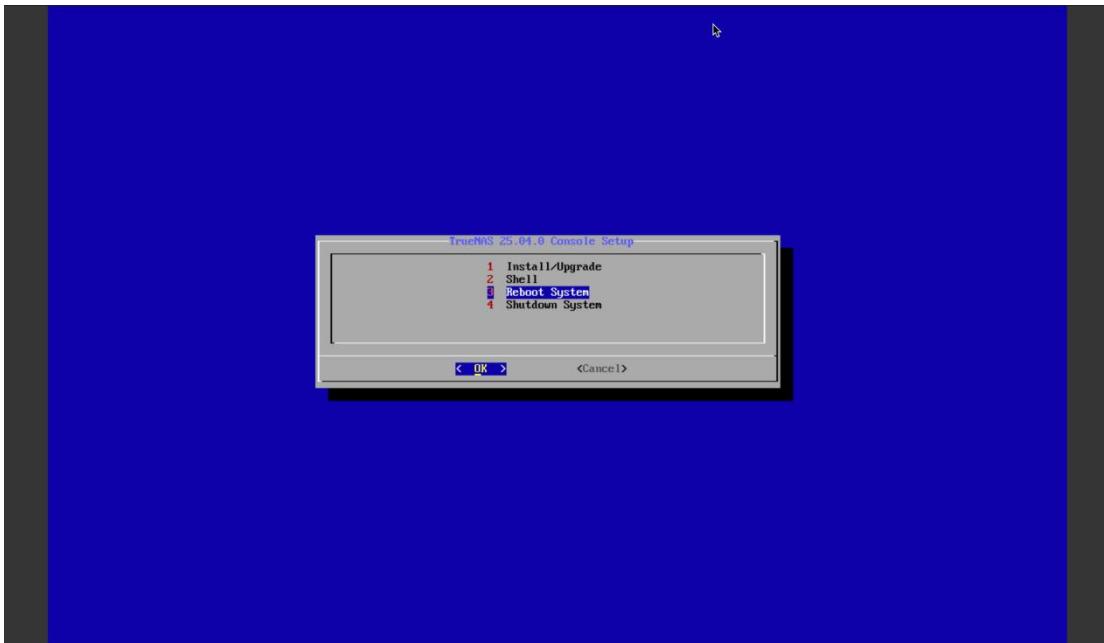






Now that the installation is completed, let us remove the attached iso file (volume) from virtual machine with the following command:

```
mac% incus config device remove vm-truenas25 install # 'install' being the install disk
```



Let us get connected back to console ones again and make sure the virtual machine boots up fine



```

[ 0.000000] Linux version 6.12.15-production+truenas (root@tnsbuilds01.tn.iks
ystems.net) (gcc (Debian 12.2.0-14) 12.2.0, GNU ld (GNU Binutils for Debian) 2.4
0) #1 SMP PREEMPT_DYNAMIC Tue Apr 15 20:07:13 UTC 2025
[ 0.069082] Kernel command line: BOOT_IMAGE=/ROOT/25.04.0/boot/vmlinuz-6.12.
15-production+truenas root=ZFS-boot-pool/ROOT/25.04.0 ro libata.allow_tpm=1 and_
iommu=on iommu=pt kvm_and.npt=1 kvm_and.avic=1 intel_iommu=on zfsforce=1 nome_co
re_multipath=M
[ 0.069158] AMD-UI: Unknown option - 'on'
[ 0.069267] Unknown kernel command line parameters "BOOT_IMAGE=/ROOT/25.04.0
/boot/vmlinuz-6.12.15-production+truenas zfsforce=1", will be passed to user spa
ce.
[ 0.069281] random: crng init done
[ 0.260154] audit: type=2000 audit(1746921619.736:1): state=initialized audit
_enabled=1 res=1
[ 1.161211] UFS: Disk quotas dquot 6.6.0
[ 1.335690] Initialise system trusted keyrings
[ 1.335724] Key type blacklist registered
[ 1.337132] integrity: Platform Keypair initialized
[ 1.337142] integrity: Machine keypair initialized
[ 1.358845] Key type asymmetric registered
[ 1.388045] Key type certificate registered
[ 2.595079] Loading compiled in X.509 certificates
[ 2.637721] Loaded X.509 cert 'Debian Secure Boot CA: 6ccece7e4cfc0d1f6149f3dd274fcc5ccb419ea1'
[ 2.637863] Loaded X.509 cert 'Debian Secure Boot Signer 2020: 00b55eb3b9'
[ 2.643739] Key type .fscript registered
[ 2.643787] Key type fscript-provisioning registered
[ 2.679277] Key type encrypted registered
Loading, please wait...
Starting systemd-udevd version 252.33-1~deb12u1
[ 3.153190] nome_core: module verification failed: signature and/or required key missing - tainting kernel
[ 3.174201] ipc: 0x0000:00:1f:0 space for GPIO uninitialized
[ 3.174201] SCSI subsystem initialized
[ 3.293169] ide4 0:0:0:1: Direct-Access QEMU QEMU HARDDISK 2.5+ PQ: 0 ANSI: 5
[ 3.320823] sd 0:0:0:1: Power-on or device reset occurred
[ 3.329163] sd 0:0:0:1: /sda1 20971520 512-byte logical blocks: (10.7 GB/10.0 GiB)
[ 3.329292] sd 0:0:0:1: sdal Write Protect is off
[ 3.329486] sd 0:0:0:1: /sda1 Write cache: enabled, read cache: enabled, doesn't support DPO or FUA
[ 3.355671] sd 0:0:0:1: /sda1 Attached SCSI disk
Begin Loading essential drivers ... done.
Begin Running /scripts/init-premount ... done.
Begin Mounting root file system ... Begin: Running /scripts/wfs-top ... done.
Begin Running /scripts/zfs-local-premount ... done.
Begin: Running /scripts/local-premount ... done.
[ 3.944560] spl: loading out-of-tree module taints kernel.
[ 4.123011] zfs: module license 'ODDL' taints kernel
[ 4.123049] Disabling lock debugging due to kernel taint
[ 4.123913] zfs: module license taints kernel.

```

```

[ OK ] Finished modprobe@drm.service - Load Kernel Module drm.
[ OK ] Finished modprobe@efi_pstore.service - Load Kernel Module efi_pstore.
[ OK ] Finished modprobe@loop.service - Load Kernel Module loop.
[ OK ] Mounted sys-fs-fuse-connections.mount - FUSE Control File System.
[ OK ] Mounted sys-fs-verity@.service - Kernel Command Line File System.
[ OK ] Finished zfs-luks-volume-wait.service - Wait for ZFS Volume (zvol) links in
[ OK ] Reached target zfs-volumes.target - ZFS volumes are ready.
[ OK ] Finished systemd-susyusers.service - Create System Users.
[ OK ] Starting systemd-tmpfiles-setup-dev.service - Create Static Device Node
[ OK ] Finished systemd-random-seed.service - Load/Save Random Seed.
[ OK ] Finished systemd-tmpfiles-setup-dev.service - Create Static Device Node
[ OK ] Reached target local-fs-pre.target - Preparation for Local File Systems
[ OK ] Reached target machines.target - Containers.
[ OK ] Starting systemd-udevd.service - Rule-based Manager for Device Events a
[ OK ] Started systemd-journald.service - Journal Service.
[ OK ] Starting Systemd Journal Input Service - Read Journal to Persistent St
[ OK ] Job systemd-udevd-trigger.service triggered. Collected alludev Devices.
[ OK ] Starting ifupdown-pre.service - Helper to synchronize boot up for ifupd
[ OK ] Finished ifupdown-pre.service - Helper to synchronize boot up for ifupd
[ OK ] Finished systemd-journal-flush.service - Flush Journal to Persistent St
[ OK ] Started systemd-udevd.service - Rule-based Manager for Device Events an
[ 26.947276] sd 0:0:0:1: Attached scsi generic sg0 type 0
[ 27.303265] netfs: FS-Cache loaded
[ 27.305696] Error: Driver 'pcspkr' is already registered, aborting...
[ OK ] Starting incus-agent.service - Incus - agent...
[ OK ] Started incus-agent.service - Incus - agent...
[ OK ] Finished incus-ndisks.service - Help on Disk Enumeration.
[ OK ] Starting middlewared.service - TrueNAS Middleware...
[ *** ] Job middlewared.service/start running (imin 27s / 6min 27s)
[ OK ] Started middlewared.service - TrueNAS Middleware.
[ OK ] Starting ix-syndisks.service - Sync Disk Cache Table...
[ OK ] Finished ix-syndisks.service - Sync Disk Cache Table.
[ OK ] Starting ix-etc.service - Generate TrueNAS /etc files...
[ OK ] Starting ix-zfs.service - Import ZFS pools...
[ OK ] Finished ix-zfs.service - Import ZFS pools.
[ OK ] Starting ix-cgroups.service - Enable required cgroups for TrueNAS...
[ OK ] Starting ix-preinit.service - Execute TrueNAS custom pre-init tasks...
[ OK ] Finished ix-cgroups.service - Enable required cgroups for TrueNAS.
[ OK ] Finished ix-preinit.service - Execute TrueNAS custom pre-init tasks.
[ *** ] Job ix-etc.service/start running (imin 43s / 6min 28s)

```

```
[ OK ] Reached target nss-user-lookup.target - User and Group Name Lookups.
[DEPEND] Dependency failed for sssd-nss.socket - SSSD NSS Service responder socket.
[DEPEND] Dependency failed for sssd-autofs.socket - SSSD AutoFS Service responder socket.
[DEPEND] Dependency failed for sssd-pac.socket - SSSD PAC Service responder socket.
[DEPEND] Dependency failed for sssd-pam-priv.socket - SSSD PAM Service responder private socket.
[DEPEND] Dependency failed for sssd-pam.socket - SSSD PAM Service responder socket.
[DEPEND] Dependency failed for sssd-ssh.socket - SSSD SSH Service responder socket.
Starting syslog-ng.service - System Logger Daemon...
Starting susstat.service - Resets System Activity Logs...
Starting systemd-logind.service - User Login Management...
[ OK ] Started zfs-zed.device - ZFS Event Daemon (zed).
[ OK ] Reached target zfs.target - ZFS Start...
[ OK ] Finished e2scrub_reap.service - Remove Stale Online ext4 Metadata Check Snapshots.
[ OK ] Started gssproxy.service - GSSAPI Proxy Daemon.
[ OK ] Finished ix-reboot.service - Exec TrueNAS reboot tasks.
[ OK ] Finished ix-shutdown.service - Exec TrueNAS shutdown tasks.
[ OK ] Started ced.sengere - Name Service Cache Daemon. g an
Starting chrony.service - chrony, an NTP client/server...
Starting containerd.service - containerd container runtime...
Starting e2scrub_reap.service - Remove Stale Online ext4 Metadata Check Snapshots...
Starting etc-setserial.service - controls configuration of serial ports...
Starting gssproxy.service - GSSAPI Proxy Daemon...
Starting ipvsadm.service - LSR / ipvsadm daemon...
Starting ix-reboot.service - Exec TrueNAS reboot tasks...
Starting ix-shutdown.service - Exec TrueNAS shutdown tasks...
Starting ix-ssh-keys.service - Save SSH keys...
Starting lm-sensors.service - Initialize hardware monitoring sensors...
[ OK ] Started lxcfs.service - FUSE filesystem for LXC.
Starting nscd.service - Name Service Cache Daemon...
Starting nvme-autoconnect.service - Connect NVMe-of subsystems automatically during boot...
Starting setserial.service - controls configuration of serial ports...
Starting smartmontools.service - Smart Monitoring and Reporting Technology (SMART) Daemon...
[ OK ] Reached target nss-user-lookup.target - User and Group Name Lookups.
[DEPEND] Dependency failed for sssd-nss.socket - SSSD NSS Service responder socket.
[DEPEND] Dependency failed for sssd-autofs.socket - SSSD AutoFS Service responder socket.
[DEPEND] Dependency failed for sssd-pac.socket - SSSD PAC Service responder socket.
[DEPEND] Dependency failed for sssd-pam-priv.socket - SSSD PAM Service responder private socket.
[DEPEND] Dependency failed for sssd-pam.socket - SSSD PAM Service responder socket.
[DEPEND] Dependency failed for sssd-ssh.socket - SSSD SSH Service responder socket.
Starting syslog-ng.service - System Logger Daemon...
Starting susstat.service - Resets System Activity Logs...
Starting systemd-logind.service - User Login Management...
[ OK ] Started zfs-zed.device - ZFS Event Daemon (zed).
[ OK ] Reached target zfs.target - ZFS Start...
[ OK ] Finished e2scrub_reap.service - Remove Stale Online ext4 Metadata Check Snapshots.
[ OK ] Started gssproxy.service - GSSAPI Proxy Daemon.
[ OK ] Finished ix-reboot.service - Exec TrueNAS reboot tasks.
[ OK ] Finished ix-shutdown.service - Exec TrueNAS shutdown tasks.
```

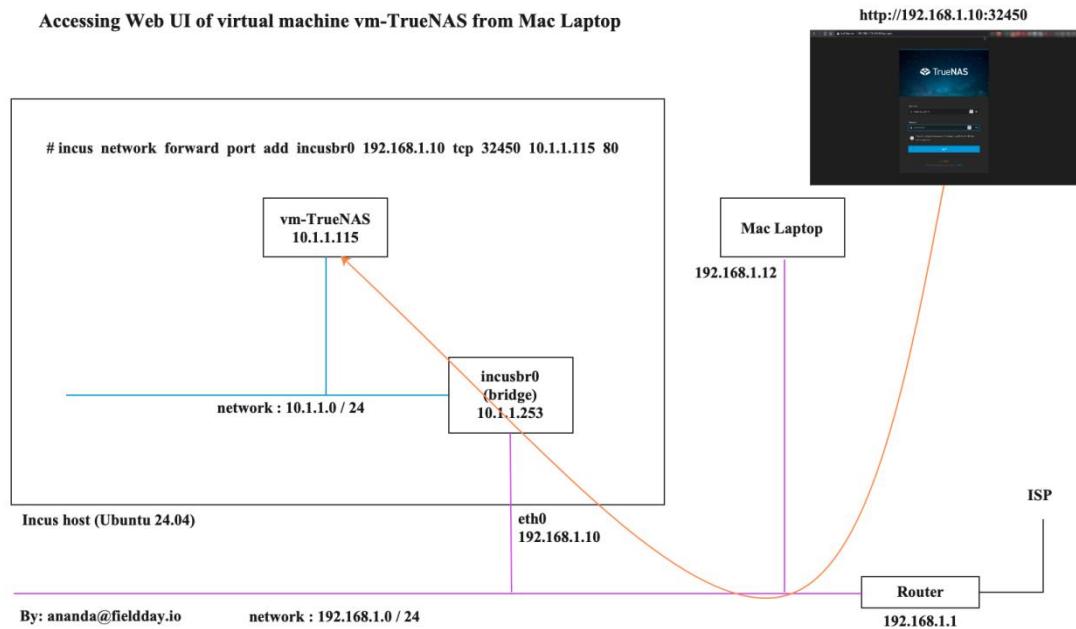
```
Console setup
-----
The web user interface is at:
http://10.1.1.115
https://10.1.1.115

1) Configure network interfaces
2) Configure network settings
3) Configure static routes
4) Change local administrator password
5) Create one-time password for "root"
6) Reset configuration to defaults
7) Open TrueNAS CLI Shell
8) Open Linux Shell
9) Reboot
10) Shutdown

Enter an option from 1-10:
```

We have reached a milestone where the installation is completed. Let us now configure TrueNAS from our Mac Laptop using the browser. But to accomplish that, we need to understand what it takes to reach out the administrative web UI running on vm-truenas25, which is on a different network segment (10.1.1.115), from the host os (192.168.1.10). which is the same network as our mac laptop.

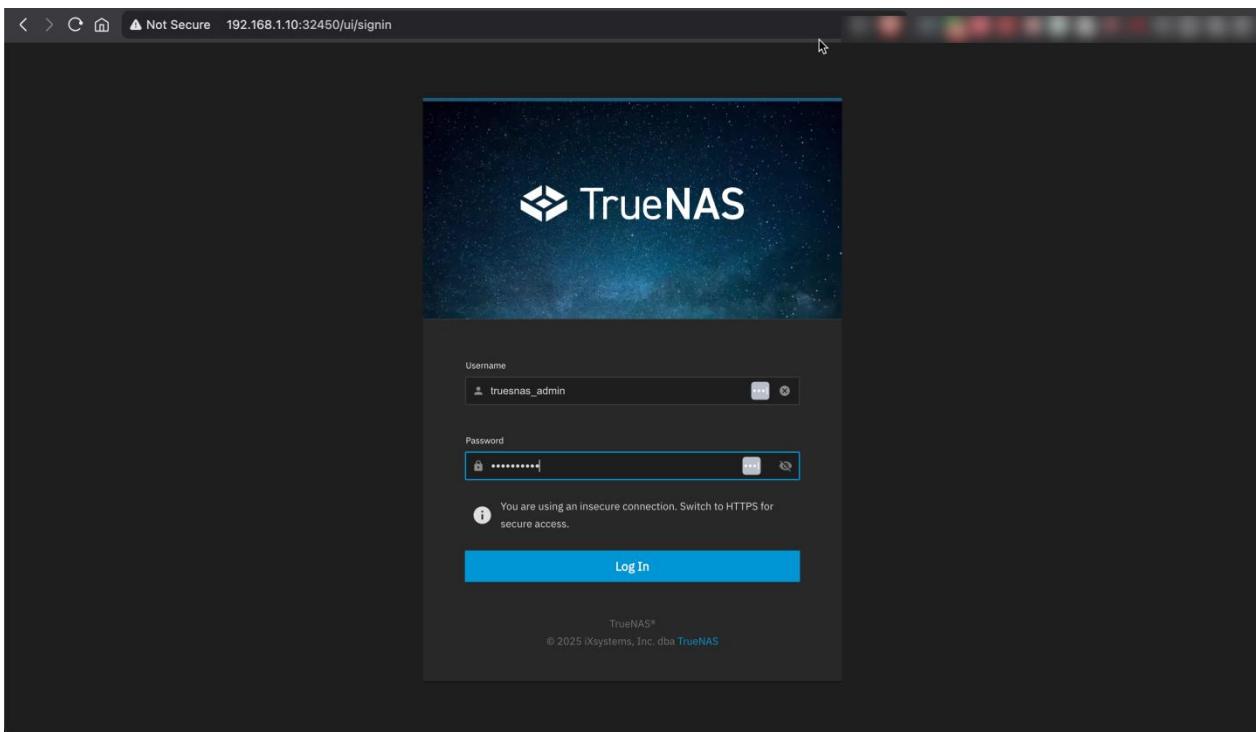
Let us understand our topology clearly so that we know what it takes to reach the incus admin web UI from mac laptop.



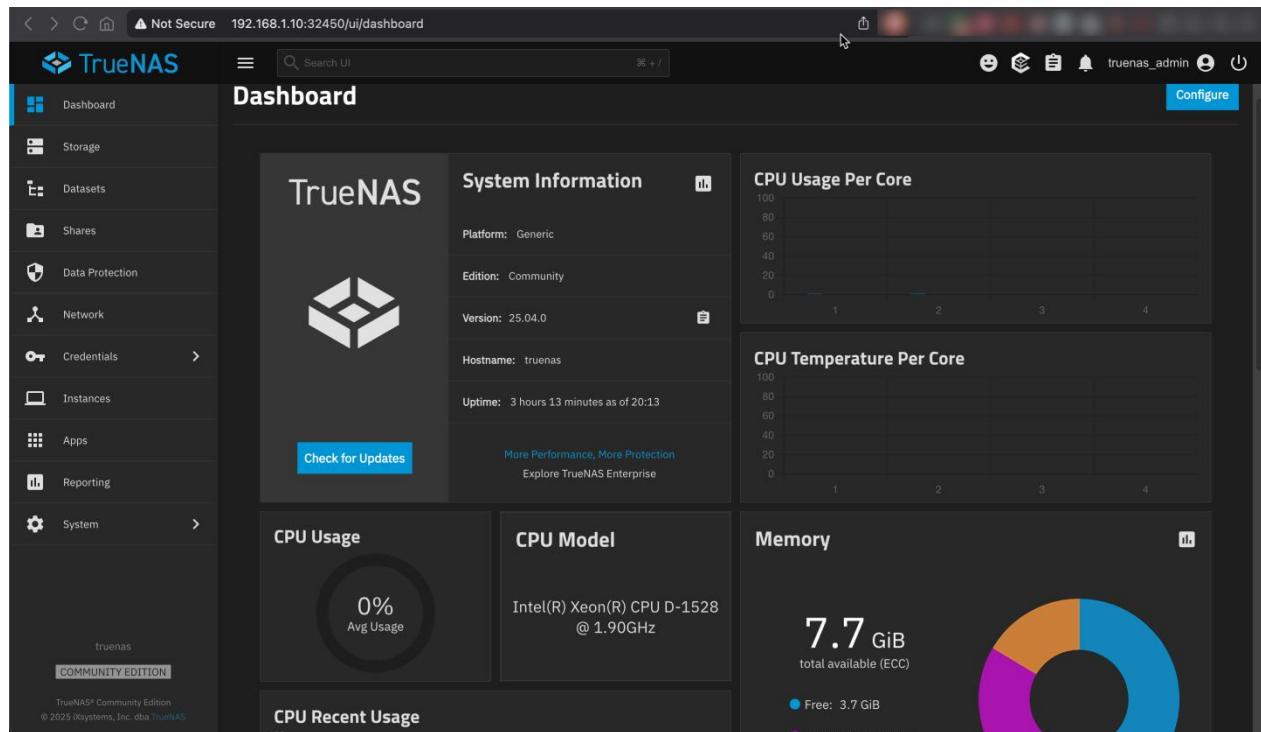
Issuing the network forward rule show below on the incus server at ip address 192.168.1.10, port 32450 will be forwarded to ip address 10.1.1.115, port number 80, via incusbr0

```
# incus network forward port add incusbr0 192.168.1.10 tcp 32450 10.1.1.115 80
```

Now pointing the browser on mac laptop to <http://192.168.1.10:32450> will get us the admin portal of TrueNAS running at 10.1.1.115



The Main Dashboard



All 5 disks will be listed

The screenshot shows the TrueNAS UI at the URL 192.168.1.10:32450/ui/storage/disks. The left sidebar has a 'Storage' section highlighted. The main area is titled 'Disks' and lists five disks:

	Name	Serial	Disk Size	Pool
<input type="checkbox"/>	sda	incus_root	10 GiB	boot-pool
<input type="checkbox"/>	nvme0n1	incus_disk1	50 GiB	N/A
<input type="checkbox"/>	nvme1n1	incus_disk2	50 GiB	N/A
<input type="checkbox"/>	nvme4n1	incus_disk3	50 GiB	N/A
<input type="checkbox"/>	nvme2n1	incus_disk4	50 GiB	N/A
<input type="checkbox"/>	nvme3n1	incus_disk5	50 GiB	N/A

At the bottom right of the table, there are buttons for 'Items per page:' (set to 50), '1 – 6 of 6', and navigation arrows.

Create a new storage pool

The screenshot shows the TrueNAS UI at the URL 192.168.1.10:32450/ui/storage. The left sidebar has a 'Storage' section highlighted with a red circle containing the number '1'. The main area is titled 'Storage Dashboard' and shows a message: 'No Pools'. It says, 'It seems you haven't configured pools yet. Please click the button below to create a pool.' Below the message is a large blue button labeled 'Create pool' with a red circle containing the number '2'.

3 Log (Optional)

4 Spare (Optional)

Select all 5 disks to create a RAIDZ1

5 Add

6

7

This screenshot shows the TrueNAS UI during the VDEV configuration process. The left sidebar includes options like Dashboard, Storage, Datasets, Shares, Data Protection, Network, Credentials, Instances, Apps, Reporting, and System. The main area displays a "Manually Configured VDEVs" section with a brief description and a "Edit Manual Disk Selection" button. To the right is a "Configuration Preview" panel showing settings for a VDEV named "testpool". A red arrow points to the "Name" field. Below it is an "Unassigned Disks" section stating "No disks available". Step 8 is highlighted with a red circle.

This screenshot shows the "Manual Selection" screen for RAIDZ1. The left sidebar is identical to the previous screenshot. The main area has two tabs: "Manual Selection" (selected) and "VDEVs". The "VDEVs" tab shows a "RAIDZ1" configuration with a total capacity of 200 GiB. It lists five disks: nvme0n1, nvme1n1, nvme2n1, nvme3n1, and nvme4n1, each with 50 GiB capacity. A red circle highlights the "Save Selection" button at the bottom right. Step 9 is highlighted with a red circle.

The screenshot shows the 'Pool Creation Wizard' step 10, titled 'Data'. It displays a 'Manually Configured VDEVs' section with a note about VDEVs being created through manual disk selection. A 'Description' section indicates 1 VDEV with no warnings. On the right, a 'Configuration Preview' panel shows pool details: Name: testpool, Cache: None, Data: Manual layout | 1 VDEVs, Dedup: None, Log: None, Spare: None, Metadata: None, Encryption: None, and Total Raw Capacity: 200 GiB. Below this is an 'Unassigned Disks' section stating 'No disks available.' At the bottom, navigation buttons include Back, Next (highlighted), Reset Step, Save And Go To Review, and a red '10' badge.

The screenshot shows the 'Pool Creation Wizard' step 11, titled 'Review'. It lists optional configurations: Log (Optional), Spare (Optional), Cache (Optional), Metadata (Optional), and Dedup (Optional). The 'Review' section shows the pool name is set to 'testpool'. The 'General Info' summary table includes:

Topology Summary	Details
Data	Manual layout 1 VDEVs
	Est. Usable Raw Capacity
	200 GiB

At the bottom, navigation buttons include Back, Inspect VDEVs, Start Over, Create Pool (highlighted), and a red '11' badge.

The screenshot shows the TrueNAS web interface during the storage pool creation process. The left sidebar includes options like Dashboard, Storage, Datasets, Shares, Data Protection, Network, Credentials, Instances, Apps, Reporting, and System. The main area is titled 'Storage' and shows steps 1 through 14. Step 8 is the 'Review' stage. A modal window titled 'Warning' contains the message: 'The contents of all added disks will be erased.' with two buttons: 'Confirm' (highlighted with a red circle) and 'Cancel'. Below the modal, the 'General Info' section shows a 'Pool Name' of 'testpool'. The 'Topology Summary' and 'Details' sections show 'Data' with a 'Manual layout | 1 VDEVs' and 'Est. Usable Raw Capacity' of '200 GiB'. At the bottom are buttons for 'Back', 'Inspect VDEVs', 'Start Over' (highlighted with a red circle), and 'Create Pool'.

This screenshot shows the same stage as the previous one, but the modal window has changed to a progress bar titled 'Create Pool' with the status 'pool.create 6.00% Formatting disks (1/5)'. Step 13 is highlighted with a red circle. The rest of the interface is identical to the first screenshot, showing the 'Review' step with its respective buttons and pool configuration details.

TrueNAS

☰ Search UI ☰ /

truenas_admin ⚙️ 📁 🗂️ 🔍

Import Pool Disks Create Pool

Storage Dashboard

testpool

Export/Disconnect Expand

Manage Devices

Manage Datasets

Topology ✓

Data VDEVs	1 x RAIDZ1 5 wide 50 GiB
Metadata VDEVs	VDEVs not assigned
Log VDEVs	VDEVs not assigned
Cache VDEVs	VDEVs not assigned
Spare VDEVs	VDEVs not assigned
Dedup VDEVs	VDEVs not assigned

Usage ✓

0%

Usable Capacity: 190.42 GiB

- Used: 11.28 MiB
- Available: 190.41 GiB

[View Disk Space Reports](#)

ZFS Health ✓

Scrub

Pool Status: Online

Total ZFS Errors: 0

Scheduled Scrub Task: Set

[View All Scrub Tasks](#)

Disk Health ✓

Manage Disks

Disks temperature related alerts: 0

Highest Temperature: No Data

Lowest Temperature: No Data

truenas

COMMUNITY EDITION

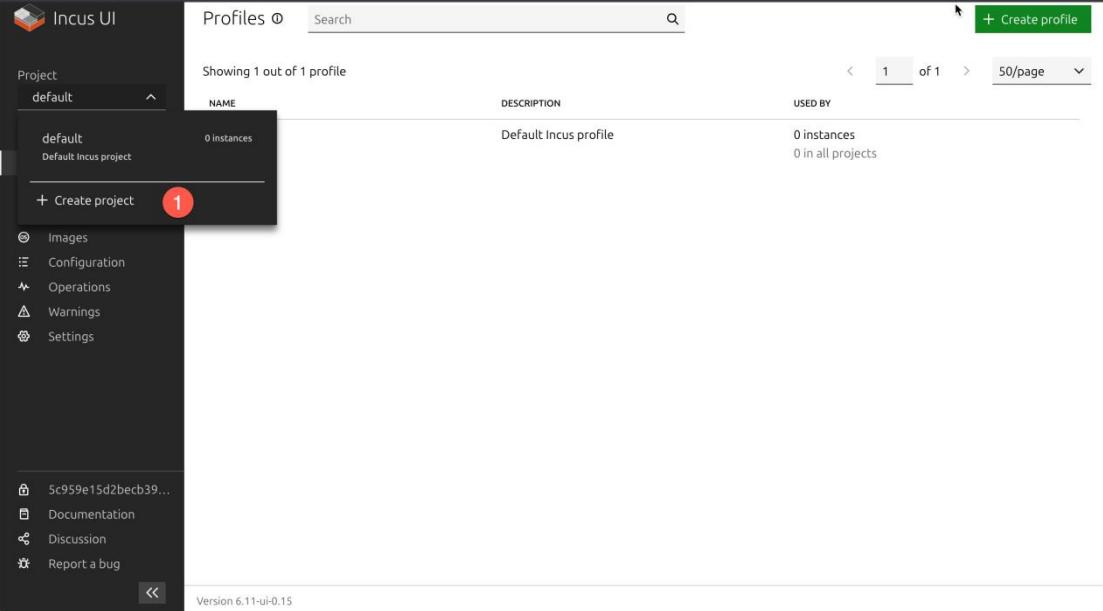
TrueNAS® Community Edition
© 2025 iXsystems, Inc. ota TrueNAS

Task #17: Creating Projects using web UI

In this task we will create projects using web UI. We will create 6 Projects in total. We will then assign each Project with its own storage, the ones we created in previous tasks.

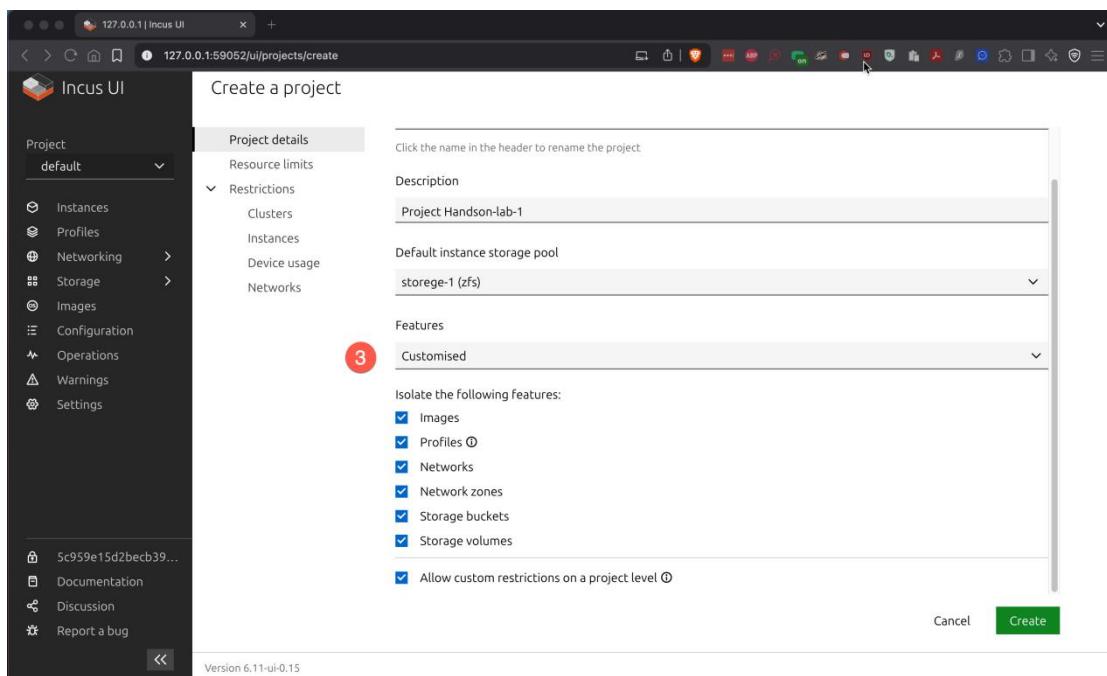
Running the command ‘incus webui’ on our Mac laptop will open up a browser session to our incus server (recollect that we have already established trust with the incus server).

Let us navigate the left panel and create a new project as shown above.



The screenshot shows the Incus UI web interface. On the left, there's a sidebar with a tree view containing 'Project' (selected), 'default' (expanded), 'Images', 'Configuration', 'Operations', 'Warnings', and 'Settings'. Under 'default', it lists '0 instances'. At the bottom of the sidebar is a red-highlighted button labeled '+ Create project' with a red circle containing the number '1'. The main area is titled 'Profiles' and shows a single profile named 'default' with '0 instances'. Below the table are links for 'Sc959e15d2becb39...', 'Documentation', 'Discussion', and 'Report a bug'. The footer indicates 'Version 6.11-ui-0.15'.

1. Let us give our first project a name ‘Handson-lab-1’.
2. Enter a relevant text under Description
3. From the dropdown menu we will select and assign ‘storage-1’ to this Project
4. Select ‘Customized’ under Features dropdown menu



- Under Restrictions, let us select ‘Instances’ and ‘Allow’ privileges to all instances

CONFIGURATION	INHERITED	OVERRIDE
Low level VM operations Whether to prevent using low-level VM options	Block From: incus	4 Allow Possible values are allow or block. When set to allow, low-level VM options like {config:option} instance-raw:raw.qemu, volatile.*_, etc. can be used.
Low level container operations Whether to prevent using low-level container options	Block From: incus	5 Allow Possible values are allow or block. When set to allow, low-level container options like {config:option} instance-raw:raw.lxc, {config:option} instance-raw:raw.idmap, volatile.*_, etc. can be used.
Container nesting Whether to prevent running nested Incus	Block From: incus	6 Allow Possible values are allow or block. When set to allow, {config:option} instance-security:security.nesting can be set to true for an instance.
Container privilege Which settings for privileged containers to prevent	Unprivileged From: incus	7 Allow Possible values are unprivileged, isolated, and allow.

- Under Restrictions, let us give privileges to disk devices and GPU devices
- Allow Snapshot creation features as well

The screenshot shows the Incus UI interface for creating a project. On the left, a sidebar lists various project settings like Instances, Profiles, Networking, Storage, Images, Configuration, Operations, Warnings, and Settings. The 'Instances' option is selected. The main area is titled 'Create a project' and contains a table for configuration settings:

CONFIGURATION	INHERITED	OVERRIDE
Container interception Whether to prevent using system call interception options	Block From: Incus	Allow
Snapshot creation Whether to prevent creating instance or volume snapshots	Block From: Incus	Allow
Idmap UID Which host UID ranges are allowed in <code>raw.idmap</code>	- From: Incus	Allow
Idmap GID Which host GID ranges are allowed in <code>raw.idmap</code>	- From: Incus	Allow

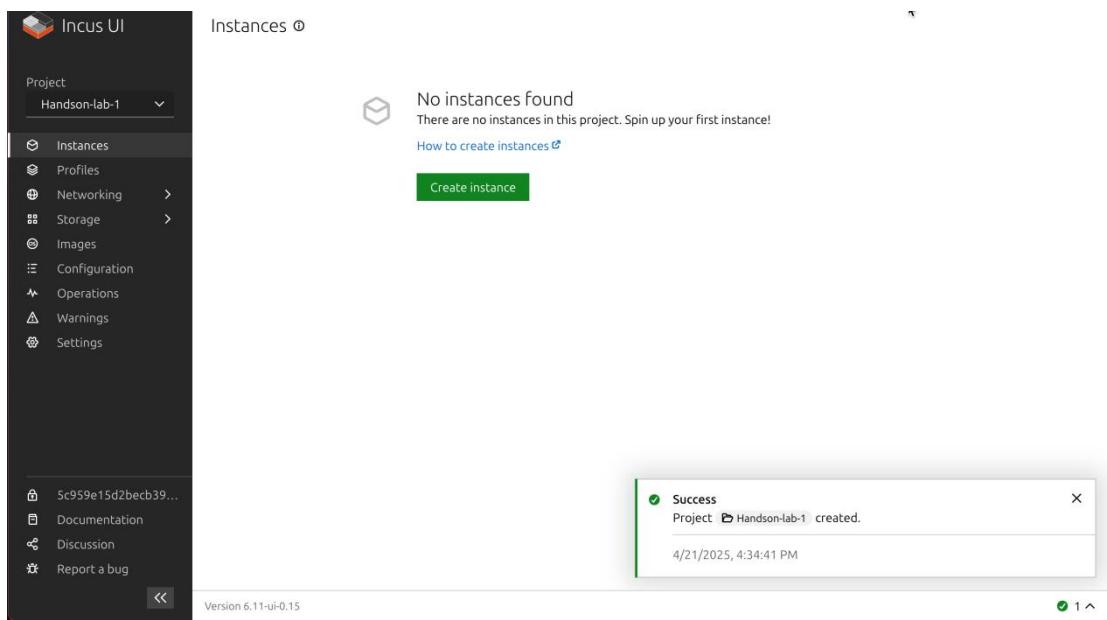
At the bottom right are 'Cancel' and 'Create' buttons. A note at the bottom says 'Version 6.11-ui-0.15'.

This screenshot shows the Incus UI interface for creating a project, specifically focusing on 'Device usage'. The sidebar and main configuration table are similar to the previous screenshot, but the 'Device usage' section is highlighted. The table includes the following rows:

CONFIGURATION	INHERITED	OVERRIDE
Disk devices (except the root one) Which disk devices can be used	Managed From: Incus	Managed
Disk devices path Which source can be used for disk devices	-	Allow
GPU devices Whether to prevent using devices of type gpu	Block From: Incus	Allow
Infiniband devices Whether to prevent using devices of type infiniband	Block From: Incus	Allow
Network devices	Managed	Allow

Red numbered circles (8, 9, 10) are placed over the 'Managed' dropdown in the first row, the 'Allow' dropdown in the fourth row, and the 'Allow' dropdown in the last row respectively. The 'Create' button is visible at the bottom right. A note at the bottom says 'Version 6.11-ui-0.15'.

8. We will get a confirmation on sucessful creation of the Project - Handson-lab-1
9. Note that on confirmation, the UI lists Handson-lab-1 under Project in the left panel
10. We need to keep in mind that any changes made in the UI will reflect on that particular project alone, until the user selects a different from project from the dropdown menu



Let us run the following command to see the current configuration of the project - handson-lab-1

```
mac$ incus project list -f compact
```

NAME	IMAGES	PROFILES	STORAGE VOLUMES	STORAGE BUCKETS	NETWORKS	NETWORK ZONES	DESCRIPTION	USED BY
Handson-lab-1	YES	YES	YES	YES	YES	YES	Project Handson-lab-1	1
default (current)	YES	YES	YES	YES	YES	YES	Default Incus project	2

```
mac$ incus project show Handson-lab-1
```

config:

```

features.images: "true"
features.networks: "false"
features.networks.zones: "false"
features.profiles: "true"
features.storage.buckets: "true"
features.storage.volumes: "true"
restricted: "true"
restricted.containers.lowlevel: allow
restricted.containers.nesting: allow
restricted.containers.privilege: allow
restricted.devices.disk: managed
restricted.devices.gpu: allow
restricted.snapshots: allow
restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-1
name: Handson-lab-1
used_by:
- /1.0/profiles/default?project=Handson-lab-1

```

From the storage side as well we can confirm that project Handson-lab-1 is indeed backed by storage-1

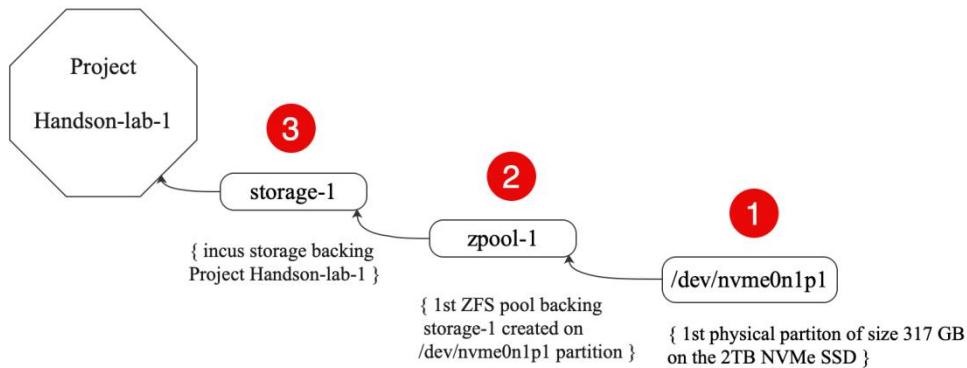
```
mac% incus storage show storage-1
config:
  source: zpool-1
  volatile.initial_source: zpool-1
  zfs.pool_name: zpool-1
description: Storage for Containers, VMs of Handson-lab-1
name: storage-1
driver: zfs
used_by:
- /1.0/profiles/default?project=Handson-lab-1
status: Created
locations:
- none
```

And we know storage-1 is indeed backed by the zfs pool `zpool-1`, which is shown above defined in the key `zfs.pool_name: zpool-1`

Task #18: Verifying storage for project: Handson-lab-1

In this task we will take a quick look at how the storage layout is being assigned to the project ‘Handson-lab-1’ (We are not planning on configuring anything in this task).

Storage layout for Project Handson-lab-1



A zfs pool named ‘zool-1’ was created on the first physical partition on the 2TB NVMe SSD `/dev/nvme0p1`. incus storage pool ‘storage-1’ was created on `zpool-1`, which is now managed by incus.

```
# lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    1 460.3G 0 disk
└─sda1     8:1    1   1M 0 part
└─sda2     8:2    1 460.3G 0 part /
zd0       230:0   0   5.8G 0 disk
nvme0n1   259:0   0   1.9T 0 disk
└─nvme0n1p1 259:1   0  317G 0 part ---> zpool-1
└─nvme0n1p2 259:2   0  317G 0 part
└─nvme0n1p3 259:3   0  317G 0 part
└─nvme0n1p4 259:4   0  317G 0 part
└─nvme0n1p5 259:5   0  317G 0 part
└─nvme0n1p6 259:6   0  317G 0 part
└─nvme0n1p7 259:7   0   5.7G 0 part
nvme1n1   259:8   0 931.5G 0 disk
└─nvme1n1p1 259:9   0  155G 0 part
└─nvme1n1p2 259:10  0  155G 0 part
└─nvme1n1p3 259:11  0  155G 0 part
└─nvme1n1p4 259:12  0  155G 0 part
```

```

└─nvme1n1p5 259:13 0 155G 0 part
└─nvme1n1p6 259:14 0 155G 0 part
└─nvme1n1p7 259:15 0 1.5G 0 part

```

```
# zpool list
```

NAME	SIZE	ALLOC	FREE	CKPOINT	EXPANDSZ	FRAG	CAP	DEDUP	HEALTH	ALTROOT
basepool	5.50G	7.13M	5.49G	-	-	22%	0%	1.00x	ONLINE	-
disks-pool	308G	2.92M	308G	-	-	0%	0%	1.00x	ONLINE	-
images-pool	462G	1.15M	462G	-	-	0%	0%	1.00x	ONLINE	-
iso-pool	154G	5.86G	148G	-	-	0%	3%	1.00x	ONLINE	-
zpool-1	316G	1002K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-2	316G	804K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-3	316G	741K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-4	316G	724K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-5	316G	747K	316G	-	-	0%	0%	1.00x	ONLINE	-
zpool-6	316G	747K	316G	-	-	0%	0%	1.00x	ONLINE	-

```
# incus storage list -f compact
```

NAME	DRIVER	DESCRIPTION	USED BY	STATE
basepool	zfs	passed to incus admin init	1	CREATED
shared-disks	zfs	shared disks accross Projects	0	CREATED
shared-images	zfs	shared images accross Projects	1	CREATED
shared-iso	zfs	shared ISOs accross Projects	1	CREATED
storage-1	zfs	Storage for Containers, VMs of Handson-lab-1	1	CREATED
storage-2	zfs	Storage for Containers, VMs of Handson-lab-2	0	CREATED
storage-3	zfs	Storage for Containers, VMs of Handson-lab-3	0	CREATED
storage-4	zfs	Storage for Containers, VMs of Handson-lab-4	0	CREATED
storage-5	zfs	Storage for Containers, VMs of Handson-lab-5	0	CREATED
storage-6	zfs	Storage for Containers, VMs of Handson-lab-6	0	CREATED

Task #19: default profile for project: Handon-lab-1

In this task we will look at the ‘default profile’ created for Project Handon-lab-1.

In Incus, a profile is a named set of configuration settings that can be applied to one or more instances (containers or VMs). Profiles allow you to define reusable configurations, making it easier to manage multiple instances with similar requirements.

We can also create profiles at various levels namely - storage, network and projects. As part of creating a new project, incus will create a default profile. Let us take a look at it via web UI and command line.

The screenshot shows the Incus UI interface. On the left, there is a sidebar with a tree view of project resources. The 'Profiles' node is expanded, and the 'default' profile is selected, indicated by a red circle with the number 3. A red circle with the number 2 is on the 'Profiles' node itself. A red circle with the number 1 is on the project dropdown menu. The main content area displays a table of profiles. The 'default' profile is listed with the following details:

NAME	DESCRIPTION	USED BY
default	Default Incus profile for project Handon-lab-1	0 instances

At the top right, there is a '+ Create profile' button. Below the table, there are navigation controls (back, forward, page number, page size) and a search bar. At the bottom, there is a footer with version information: Version 6.12-ui-0.15.

The screenshot shows the detailed view of the 'default' profile. The sidebar on the left is identical to the previous screenshot. The main content area has a header 'Profiles / default' with a 'Delete' button. Below the header, there are tabs for 'Overview' (selected) and 'Configuration'. The 'Overview' tab displays the following information:

General	Name	default	
	Description	Default Incus profile for project Handon-lab-1	
Networks	-		
Devices	NAME	TYPE	DETAILS
	root	disk (root)	Pool ● storage-1
Limits	CPU	-	
	Memory	-	

Below the table, there is a section titled 'Usage (0)'. At the bottom, there is a footer with version information: Version 6.12-ui-0.15.

Incus UI

Project: Handson-lab-1

- Instances
- Profiles**
- Networking >
- Storage >
- Pools
- Volumes
- Custom ISOs
- Images
- Configuration
- Operations
- Warnings
- Settings

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

YAML Configuration Version 6.12-ui-0.15

Profiles / default

5

Overview Configuration

Main configuration

Devices

Disk 6

Root storage

CONFIGURATION	INHERITED	OVERRIDE
Root storage		X
Pool		storage-1 ↲
Size		unlimited ↲

+ Attach disk device

Resource limits

Security policies

Snapshots

Migration

Boot

Cloud init

Delete

Incus UI

Project: Handson-lab-1

- Instances
- Profiles**
- Networking >
- Storage >
- Pools
- Volumes
- Custom ISOs
- Images
- Configuration
- Operations
- Warnings
- Settings

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

YAML Configuration Version 6.12-ui-0.15

Profiles / default

7

Overview Configuration

Main configuration

Devices

Disk

Exposed CPU limit

Which CPUs to expose to the instance

INHERITED From: Incus OVERRIDE ↲

Memory limit

Usage limit for the host's memory

INHERITED From: Incus OVERRIDE ↲

Memory swap (Containers only)

Control swap usage by the instance

INHERITED From: Incus OVERRIDE ↲

Disk priority

Priority of the instance's I/O requests

INHERITED From: Incus OVERRIDE ↲

Max number of processes (Containers only)

Maximum number of processes that can run in the instance

INHERITED From: Incus OVERRIDE ↲

Delete

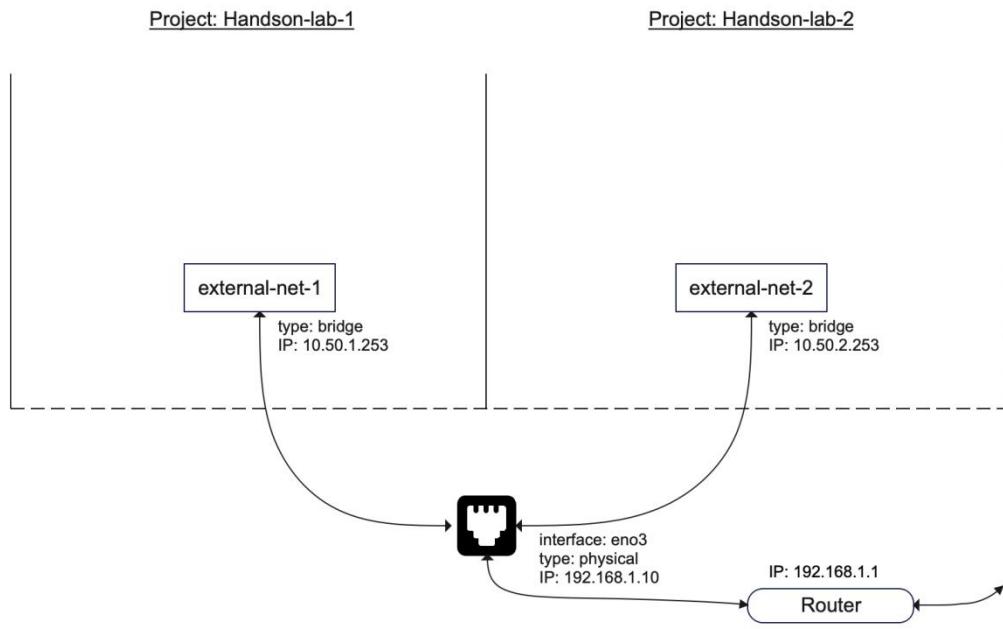
The same defaults can also be listed via command lines

```
mac% incus profile list --project Handson-lab-1
+-----+-----+
| NAME | DESCRIPTION | USED BY |
+-----+-----+
| default | Default Incus profile for project Handson-lab-1 | 0 |
+-----+-----+
```

```
mac% incus profile show default --project Handson-lab-1
config:
  security.nesting: "true"
  security.privileged: "true"
description: Default Incus profile for project Handson-lab-1
devices:
  root:
    path: /
    pool: storage-1
    type: disk
name: default
used_by: []
project: Handson-lab-1
```

Task #20: Creating external network for project: Handson-lab-1

In this task (and the ones following) we will focus primarily on the networking aspects of incus. To begin with, we will create an external network for project Handson-lab-1. Let us take a look at its network topology.



To get an overall idea of what we eventually wanted to build, the above diagram showcases the networking topology of two different projects. For now, we will focus on just the first project alone (Handson-lab-1).

Let us first switch to the project Handson-lab-1

```
mac% incus project list
+-----+-----+-----+-----+-----+-----+
|     NAME      | IMAGES | PROFILES | STORAGE VOLUMES | STORAGE BUCKETS | NETWORKS | NETWORK ZONES |     DESCRIPTION      | USED BY |
+-----+-----+-----+-----+-----+-----+
| Handson-lab-1 | YES   | YES   | YES   | YES   | YES   | YES   | Project Handson-lab-1 | 1           |
+-----+-----+-----+-----+-----+-----+
| default (current) | YES   | YES   | YES   | YES   | YES   | YES   | Default Incus project | 4           |
+-----+-----+-----+-----+-----+-----+
```

```
mac% incus project switch Handson-lab-1
```

```
mac% incus project show Handson-lab-1
```

```
config:
```

```
  features.images: "true"
```

```

features.networks: "true"
features.networks.zones: "true"
features.profiles: "true"
features.storage.buckets: "true"
features.storage.volumes: "true"
restricted: "true"
restricted.containers.lowlevel: allow
restricted.containers.nesting: allow
restricted.containers.privilege: allow
restricted.devices.disk: managed
restricted.devices.gpu: allow
restricted.snapshots: allow
restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-1
name: Handson-lab-1
used_by:
- /1.0/profiles/default?project=Handson-lab-1

```

Let us take a look at the current network setup

```

mac% incus network list
+-----+-----+-----+-----+-----+-----+
|   NAME   |   TYPE   | MANAGED |   IPV4    |   IPV6   | DESCRIPTION | USED BY | STATE  |
+-----+-----+-----+-----+-----+-----+
| br-int   | bridge   | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+
| docker0  | bridge   | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+
| eno1     | physical | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+
| eno2     | physical | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+
| eno3     | physical | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+
| eno4     | physical | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+
| incusbr0 | bridge   | YES     | 10.1.1.253/24 | none |           | 1        | CREATED |
+-----+-----+-----+-----+-----+-----+
| lo       | loopback | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+
| ovs-system | unknown | NO      |           |           |           | 0        |         |
+-----+-----+-----+-----+-----+-----+

```

Let us create an external network. Note that we are NOT using the incus network bridge (incusbr0) in our setup. It means any container or a virtual machine that connects to the external network, will not go through the incusbr0 bridge but its network traffic will directly go to the physical network interface ‘eno3’. We may want to have the network topology this way to get some clarity on incus networking features and concepts.

Let us also have some flexibility in the network access among projects. To have network isolation in place we will create network (access control lists) ACL rules, which we will work on a later task (#28). For now let us de-select networks and network zones for project Handson-lab-1 using web UI.

Project configuration ⚡ / Handson-lab-1

Project details

Description
Project Handson-lab-1

Default instance storage pool
storage-1 (zfs)

Features
Customised

Isolate the following features:

- Images
- Profiles ⓘ
- Networks
- Network zones
- Storage buckets
- Storage volumes
- Allow custom restrictions on a project level ⓘ

Cancel **Save 2 changes**

Project configuration ⚡ / Handson-lab-1

Project details

Description
Project Handson-lab-1

Default instance storage pool
storage-1 (zfs)

Features
Customised

Isolate the following features:

- Images
- Profiles ⓘ
- Networks
- Network zones
- Storage buckets
- Storage volumes
- Allow custom restrictions on a project level ⓘ

Success
Project Handson-lab-1 updated.
5/19/2025, 12:35:41 PM

In the following command we pass all the required parameters to create an external network (named ‘external-net-1’) - exclusively dedicated to Project Handson-lab-1.

```
mac% incus network create external-net-1 ipv4.address=10.50.1.253/24 ipv4.nat=true
ipv6.address=none ipv4.dhcp.ranges=10.50.1.101-10.50.1.200 --project=Handson-lab-1
```

```
mac% incus network list
```

NAME	TYPE	MANAGED	IPV4	IPV6	DESCRIPTION	USED BY	STATE
br-int	bridge	NO				0	
docker0	bridge	NO				0	
eno1	physical	NO				0	
eno2	physical	NO				0	
eno3	physical	NO				0	
eno4	physical	NO				0	
external-net-1	bridge	YES	10.50.1.253/24	none		0	CREATED
incusbr0	bridge	YES	10.1.1.253/24	none		1	CREATED
lo	loopback	NO				0	
ovs-system	unknown	NO				0	

```
mac$ incus network show external-net-1
```

```
config:
  ipv4.address: 10.50.1.253/24
  ipv4.dhcp.ranges: 10.50.1.101-10.50.1.200
  ipv4.nat: "true"
  ipv6.address: none
description: ""
name: external-net-1
type: bridge
used_by: []
managed: true
status: Created
locations:
```

```
- none  
project: default
```

Let us go ahead and edit the Description field of the newly created external-net-1 to make it verbatim

```
mac$ incus network edit external-net-1  
config:  
  ipv4.address: 10.50.1.253/24  
  ipv4.dhcp.ranges: 10.50.1.101-10.50.1.200  
  ipv4.nat: "true"  
  ipv6.address: none  
  description: "uplink for Handson-lab-1"  
name: external-net-1  
type: bridge  
used_by: []  
managed: true  
status: Created  
locations:  
- none  
project: default
```

```
mac$ incus network list
```

	NAME		TYPE		MANAGED		IPV4		IPV6		DESCRIPTION		USED BY		STATE	
	br-int		bridge		NO								0			
	docker0		bridge		NO								0			
	eno1		physical		NO								0			
	eno2		physical		NO								0			
	eno3		physical		NO								0			
	eno4		physical		NO								0			
	external-net-1		bridge		YES		10.50.1.253/24		none		uplink for Handson-lab-1		0		CREATED	
	incusbr0		bridge		YES		10.1.1.253/24		none				1		CREATED	
	lo		loopback		NO								0			

```

| ovs-system | unknown | NO | | | | 0 | | |
+-----+-----+-----+-----+-----+-----+

```

Now that we have created the external-net-1, let us create a network profile for it. By passing this network profile while launching a container or a virtual machine, will make that instance connect to the external-net-1. This concept will be clear once we see an example (which we will carry it out as a separate task).

```

mac$ incus profile create external-net
mac$ incus profile edit external-net
config: {}
description: external-net profile for project Handson-lab-1
devices:
eth0:
  name: eth0
  network: external-net-1
  type: nic
name: external-net
used_by: []
project: Handson-lab-1

```

The network profile can also be viewed via web UI

NAME	DESCRIPTION	USED BY
default	Default Incus profile for project Handson-lab-1	0 instances
external-net	external-net profile for project Handson-lab-1	0 instances

Similar to creating a network profile, let us create a storage profile where we define the storage pool dedicated to project Handson-lab-1 alone (not sharing with any other projects, delivers storage isolation and multi-tenancy)

```
mac$ incus profile create storage
mac$ incus profile edit storage
config: {}
description: storage profile for project Handson-lab-1
devices:
root:
path: /
pool: storage-1
```

```

type: disk
name: storage
used_by: []
project: Handson-lab-1

```

Let us again verify this via web UI

The screenshot shows the Incus UI interface. The left sidebar is titled "Incus UI" and includes a "Project" dropdown set to "Handson-lab-1". The sidebar also lists "Instances", "Profiles" (which is highlighted with a red circle labeled 2), "Networking", "Networks", "ACLs", "Storage" (with a red circle labeled 3), "Images", "Configuration", "Operations", "Warnings", and "Settings". Below the sidebar are links for "c88b6a7dd7e8fa5f...", "Documentation", "Discussion", and "Report a bug". The main content area is titled "Profiles" and shows a table with three rows:

NAME	DESCRIPTION	USED BY
default	Default Incus profile for project Handson-lab-1	0 instances
external-net	external-net profile for project Handson-lab-1	0 instances
storage	storage profile for project Handson-lab-1	0 instances

At the top right of the main content area is a green button labeled "+ Create profile". At the bottom right is a green checkmark icon with the number "1" and an upward arrow.

The screenshot shows the "Profiles / storage" configuration page for the "storage" profile. The left sidebar is identical to the previous screenshot. The main content area has tabs for "Overview" (which is selected) and "Configuration". The "Overview" tab displays the following details for the "storage" profile:

- General**: Name is "storage", Description is "storage profile for project Handson-lab-1".
- Networks**: None listed.
- Devices**: A table with one row:

NAME	TYPE	DETAILS
root	disk (root)	Pool: storage-1
- Limits**: CPU and Memory limits are listed.
- Usage (0)**: None listed.

At the top right is a "Delete" button. At the bottom right is a green checkmark icon with the number "1" and an upward arrow.

Incus UI

Profiles / storage

Project: Handson-lab-1

Overview Configuration 4

Devices 5

	CONFIGURATION	INHERITED	OVERRIDE
Root storage	x		storage-1 ↗ unlimited ↗
Network			
GPU			
Proxy			
Other			
Resource limits			
Security policies			
Snapshots			
Migration			
Boot			
Cloud init			

+ Attach disk device

YAML Configuration

c88bb6a7dd7ebfa5f...
Documentation
Discussion
Report a bug

Version 6.12-ui-0.15 1^

Task #21: Sanity testing external network of project: Handson-lab-1

Now that we have completed creating an external network for project Handson-lab-1, let us launch a container instance (c1), passing the profiles that we have created and see how the container consume resources.

```
mac% incus launch images:ubuntu/24.04 c1 --profile=external-net --profile=storage -d root,size=5GiB
```

```
# incus list
```

NAME	STATE	IPV4	IPV6	TYPE	SNAPSHOTS
c1	RUNNING	10.50.1.181 (eth0)		CONTAINER	0

The first thing to notice is the IP address assigned to the container c1 (10.50.1.181). You may recollect that when we created the external network (external-net-1) in the previous task, we passed a parameter ‘ipv4.dhcp.ranges=10.50.1.101-10.50.1.200’ to the ‘incus network create’ command. We can now confirm that the ip address assigned to the container (c1) is provided from that range.

```
mac% incus info c1
Name: c1
Description:
Status: RUNNING
Type: container
Architecture: x86_64
PID: 84472
Created: 2025/05/19 21:46 UTC
Last Used: 2025/05/19 21:46 UTC
Started: 2025/05/19 21:46 UTC
```

Resources:

```
Processes: 12
Disk usage:
    root: 1.06MiB
CPU usage:
    CPU usage (in seconds): 2
Memory usage:
    Memory (current): 44.16MiB
Network usage:
    eth0:
```

```

Type: broadcast
State: UP
Host interface: veth4639fcf4
MAC address: 10:66:6a:dd:75:84
MTU: 1500
Bytes received: 389B
Bytes sent: 1.44kB
Packets received: 2
Packets sent: 16
IP addresses:
  inet: 10.50.1.181/24 (global)
    inet6: fe80::1266:6aff:fedd:7584/64 (link)

lo:
Type: loopback
State: UP
MTU: 65536
Bytes received: 0B
Bytes sent: 0B
Packets received: 0
Packets sent: 0
IP addresses:
  inet: 127.0.0.1/8 (local)
    inet6: ::1/128 (local)

```

Let us get inside the container and run a few commands.

```

mac% incus shell c1
root@c1:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
  link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
  inet 127.0.0.1/8 scope host lo
    valid_lft forever preferred_lft forever
  inet6 ::1/128 scope host
    valid_lft forever preferred_lft forever
21: eth0@if22: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
  link/ether 10:66:6a:dd:75:84 brd ff:ff:ff:ff:ff:ff link-netnsid 0
  inet 10.50.1.181/24 metric 100 brd 10.50.1.255 scope global dynamic eth0
    valid_lft 2262sec preferred_lft 2262sec
  inet6 fe80::1266:6aff:fedd:7584/64 scope link
    valid_lft forever preferred_lft forever

root@c1:~# apt update -y
root@c1:~# apt upgrade -y

```

```
root@c1:~# apt install -y traceroute
```

Let us run the traceroute command reaching out to website zabbly.com (maintainers of project incus).

```
root@c1:~# traceroute zabbly.com
traceroute to zabbly.com (45.45.148.7), 30 hops max, 60 byte packets
 1 _gateway.incus (10.50.1.253)  0.066 ms  0.025 ms  0.020 ms
 2 _gateway (192.168.1.1)  0.333 ms  0.874 ms  0.828 ms
 3 10.61.209.3 (10.61.209.3)  19.560 ms  19.519 ms  19.478 ms
 4 po-323-339-rur201.sanjose.ca.sfba.comcast.net (96.216.8.237)  19.359 ms * *
[ removed the rest of the output ]
```

To reach zabbly.com (45.45.148.7) the network traffic goes to external-net-1 first (10.50.1.253). From there it goes to 192.168.1.1 which is the ip address of the router on the physical network.

It is important to note that the network traffic did NOT go through incus network bridge (incusbr0), which was our intention when we designed the network topology in the previous task.

Let us take a look at how the storage is spread cross for the container c1.

```
mac% incus storage show storage-1
config:
  source: zpool-1
  volatile.initial_source: zpool-1
  zfs.pool_name: zpool-1
description: Storage for Containers, VMs of Handson-lab-1
name: storage-1
driver: zfs
used_by:
-
/1.0/images/13fbcccd9a71b3110757f537ed43f3bd23d56f655c9915a654baf25cb6f1da80e?project=Ha
ndson-lab-1
- /1.0/instances/c1?project=Handson-lab-1
- /1.0/profiles/default?project=Handson-lab-1
- /1.0/profiles/storage?project=Handson-lab-1
status: Created
locations:
- none
```

And finally let us look at the same container via web UI

Incus UI

Project: Handson-lab-1

Instances ①

Showing 1 out of 1 instance

NAME	TYPE	DESCRIPTION	IPV4	STATUS
c1 ③	Container		10.50.1.181	Running

+ Create instance

1 Instances ②

2 Profiles

3 Networking Networks ACLs Storage Images Configuration Operations Warnings Settings

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

Version 6.12-ui-0.15

This screenshot shows the 'Instances' page of the Incus UI. The sidebar on the left is titled 'Handson-lab-1' and contains links for Instances (highlighted with a red circle), Profiles, Networking, Storage, Images, Configuration, Operations, Warnings, and Settings. The main area shows a table with one instance: 'c1' (highlighted with a red circle), which is a Container type with an IPv4 address of 10.50.1.181 and a status of 'Running'. There are buttons for 'Create instance', 'Duplicate', 'Export', and 'Delete' at the top right. Below the table, there are links for documentation, discussion, and reporting bugs.

Incus UI

Project: Handson-lab-1

Instances / c1

Running

+ Create Image | Duplicate | Export | Delete

Overview Configuration Snapshots Terminal Console Logs

Usage ④

	CPU Time(s)	
Memory	12.44	116.7 MiB of 122.7 GiB memory used
Disk	76.3 MiB of 5.0 GiB disk used	

Networks

NAME	INTERFACE	TYPE	MANAGED
external-net-1	eth0	bridge	Yes

Devices

NAME	TYPE	DETAILS
root	disk (root)	Pool storage-1

Profiles

NAME	DESCRIPTION
external-net	external-net profile for project Handson-lab-1
storage	storage profile for project Handson-lab-1

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

Version 6.12-ui-0.15

This screenshot shows the 'Instances / c1' page of the Incus UI. The sidebar on the left is identical to the previous screenshot. The main area has tabs for Overview, Configuration, Snapshots, Terminal, Console, and Logs, with 'Overview' selected (highlighted with a red circle). It displays resource usage statistics: CPU time (12.44), memory usage (116.7 MiB of 122.7 GiB), and disk usage (76.3 MiB of 5.0 GiB). Below this are sections for Networks (listing 'external-net-1' with interface eth0, type bridge, and managed status) and Devices (listing 'root' as disk (root) in pool 'storage-1'). A 'Profiles' section lists 'external-net' and 'storage' profiles. The bottom of the page includes links for documentation, discussion, and reporting bugs.

Incus UI

Instances / c1

Running | 5 | 1 | 0

+ Create Image | Duplicate | Export | Delete

Project: Handson-lab-1

Overview Configuration Snapshots Terminal Console Logs

Main configuration

Instances Profiles Networking Networks ACLs Storage Images Configuration Operations Warnings Settings

Disk 5

Name: c1
Click the instance name in the header to rename the instance

Description: Enter description

Profiles: external-net (Remove) storage (Remove)

Resource limits Security policies Snapshots Migration Boot Cloud init

Add profile

YAML Configuration

Version 6.12-ui-0.15

1

Incus UI

Instances / c1

Running | 5 | 1 | 0

+ Create Image | Duplicate | Export | Delete

Project: Handson-lab-1

Overview Configuration Snapshots Terminal Console Logs

Main configuration

Instances Profiles Networking Networks ACLs Storage Images Configuration Operations Warnings Settings

Disk

Root storage

CONFIGURATION	INHERITED	OVERRIDE
Root storage		X
Pool	storage-1 From: storage-profile	storage-1
Size	unlimited From: storage-profile	5GiB

+ Attach disk device

YAML Configuration

Version 6.12-ui-0.15

1

Incus UI

Instances / c1

Running | ⏪ | ⏴ | II | ☐

Configuration

CONFIGURATION	INHERITED	OVERRIDE
Root storage		X
Pool	storage-1 From:storage-profile	storage-1
Size	unlimited From:storage-profile	5GiB

+ Attach disk device

Main configuration

- Devices
 - Disk **6**
 - Network
 - GPU
 - Proxy
 - Other
- Resource limits
- Security policies
- Snapshots
- Migration
- Boot
- Cloud init

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

YAML Configuration

Version 6.12-ui-0.15

Incus UI

Instances / c1

Running | ⏪ | ⏴ | II | ☐

Configuration

CONFIGURATION	INHERITED	OVERRIDE
eth0	external-net-1 From: external-net profile	ⓘ

+ Attach network

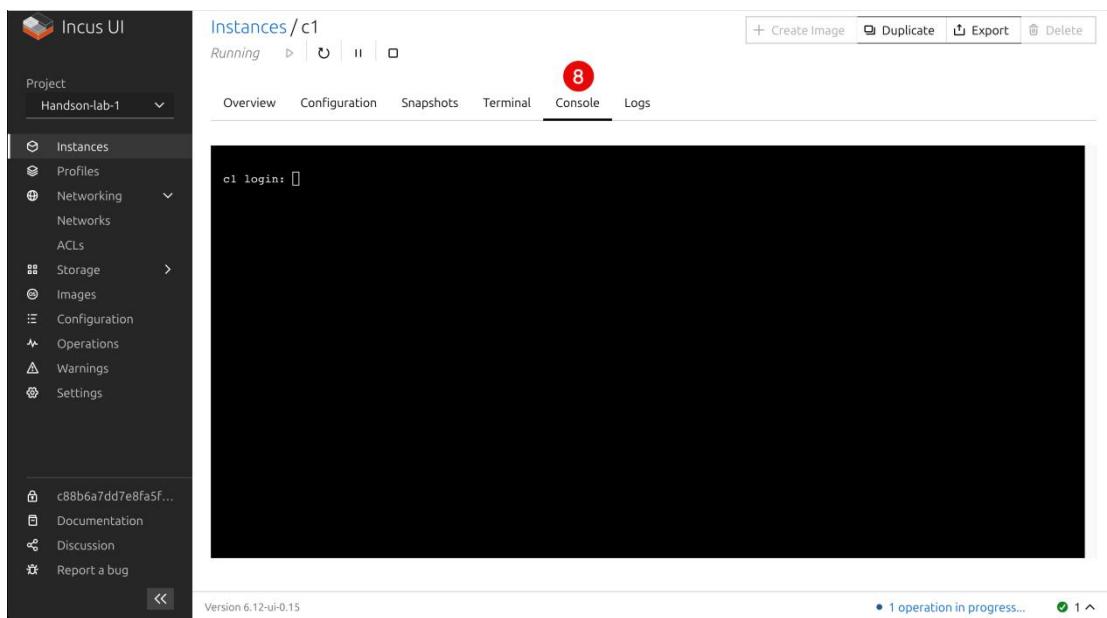
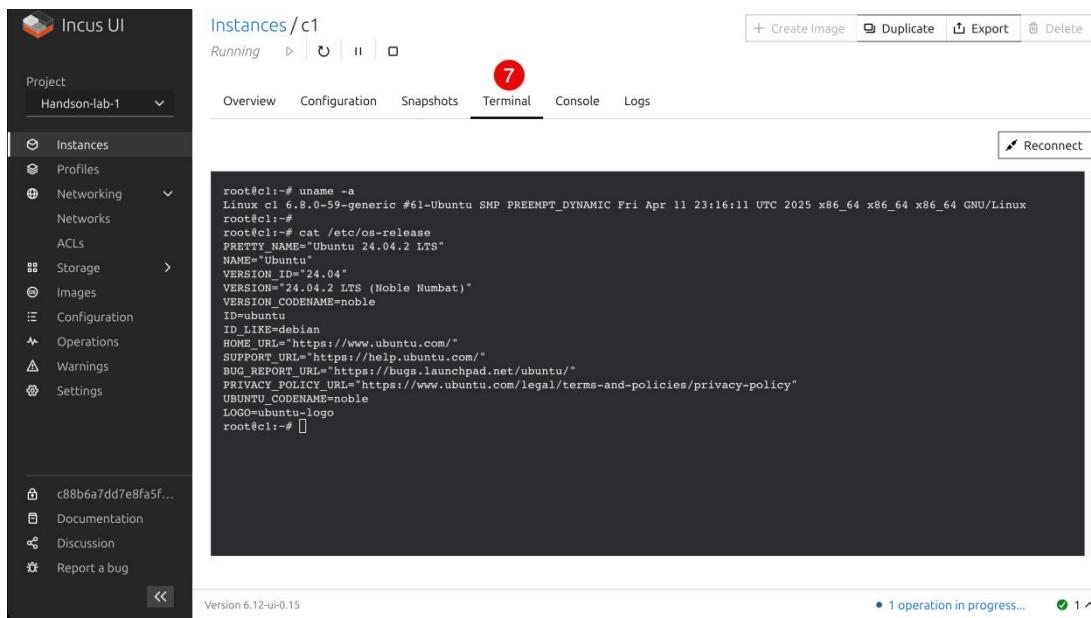
Main configuration

- Devices
 - Disk
 - Network **1**
 - GPU
 - Proxy
 - Other
- Resource limits
- Security policies
- Snapshots
- Migration
- Boot
- Cloud init

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

YAML Configuration

Version 6.12-ui-0.15

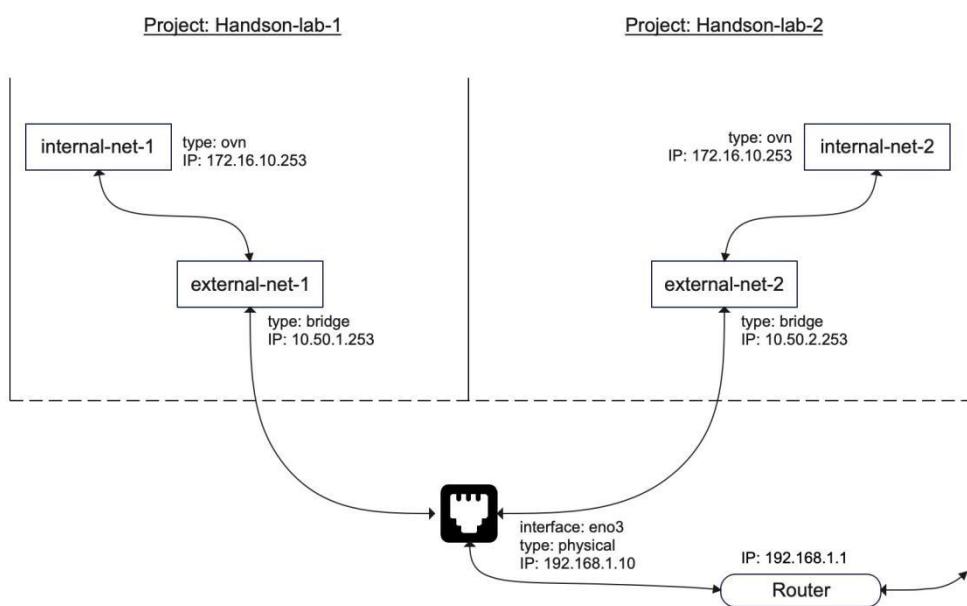


Task #22: Creating an Internal network for project: Handson-lab-1

Let us now move onto creating a second network dedicated to the project Handson-lab-1. Let us call it ‘internal-net-1’. The reason we want to name it ‘internal network’ is because we want to limit the scope of this network, exclusively for all the containers and virtual machines that run within project Handson-lab-1.

The second important concept we want to understand in this task is about OVN networks (Open Virtual Networking). By using the OVN network type for internal-net-1, we can repurpose the same IP address range across Projects. The importance of this feature may not stand out right away immediately. But as we make progress, we will see its power and advantage down the line when we build more usecases/tasks around it.

Let us take a look at its network topology where we are introducing internal network and its dependency on the external network that we already have in place.



Let us set an ipv4 ovn range on the external network

```
mac% incus network set external-net-1 ipv4.ovn.ranges=10.50.1.1-10.50.1.25
```

Let us enable ipv4.routing on the external network

```
mac% incus network set external-net-1 ipv4.routing=true
```

Let us go ahead and create the internal network of type OVN

```
mac% incus network create internal-net-1 --type=ovn ipv4.address=172.16.10.253/24
```

```
ipv4.nat=true ipv6.address=none network=external-net-1
```

We will set a dhcp range as well for the internal network

Let us edit the description of internal network to make it verbatim

```
mac% incus network edit internal-net-1
config:
  bridge.mtu: "1500"
  ipv4.address: 172.16.10.253/24
  ipv4.dhcp.ranges: 172.16.10.101-172.16.10.200
  ipv4.nat: "true"
  ipv6.address: none
  network: external-net-1
  volatile.network.ipv4.address: 10.50.1.1
description: "routes via external-net-1"
name: internal-net-1
type: ovn
used_by: []
managed: true
status: Created
locations:
- none
project: default
```

```
mac% incus network list
+-----+-----+-----+-----+-----+-----+-----+
| NAME | TYPE | MANAGED | IPV4 | IPV6 | DESCRIPTION | USED BY | STATE |
+-----+-----+-----+-----+-----+-----+-----+
| external-net-1 | bridge | YES | 10.50.1.253/24 | none | uplink for Handson-lab-1 | 3 | CREATED |
+-----+-----+-----+-----+-----+-----+-----+
| incusbr0 | bridge | YES | 10.1.1.253/24 | none | | 1 | CREATED |
+-----+-----+-----+-----+-----+-----+-----+
| internal-net-1 | ovn | YES | 172.16.10.253/24 | none | routes via external-net-1 | 0 | CREATED |
+-----+-----+-----+-----+-----+-----+-----+
```

While we are at it, we will go ahead and make changes to the description of incus bridge as well (incusbr0)

```
mac% incus network edit incusbr0
config:
  ipv4.address: 10.1.1.253/24
  ipv4.nat: "true"
  ipv6.address: none
```

```

description: "routes traffic via eno3"
name: incusbr0
type: bridge
used_by:
- /1.0/profiles/default
managed: true
status: Created
locations:
- none
project: default

```

```

mac% incus network list
+-----+-----+-----+-----+-----+-----+
|     NAME      | TYPE   | MANAGED | IPV4          | IPV6          | DESCRIPTION           | USED BY | STATE  |
+-----+-----+-----+-----+-----+-----+
| external-net-1 | bridge | YES    | 10.50.1.253/24 | none          | uplink for Handson-lab-1 | 3        | CREATED |
+-----+-----+-----+-----+-----+-----+
| incusbr0       | bridge | YES    | 10.1.1.253/24  | none          | routes traffic via eno3  | 1        | CREATED |
+-----+-----+-----+-----+-----+-----+
| internal-net-1 | ovn    | YES    | 172.16.10.253/24 | none          | routes via external-net-1 | 0        | CREATED |
+-----+-----+-----+-----+-----+-----+

```

Similar to external network, let us create a profile for the internal network as well.

```
mac% incus profile create internal-net
```

We will make the description field verbatim

```

mac% incus profile edit internal-net
config:
description: "internal network profile for project Handson-lab-1"
devices:
eth0:
  name: eth0
  network: internal-net-1
  type: nic
name: internal-net
used_by: []
project: Handson-lab-1

```

Task #23: Sanity testing internal network of project: Handson-lab-1

Now that we have completed creating an internal network for project Handson-lab-1, let's launch a container instance, pass the profiles that we have created for internal network and see how the container consumes resources.

```
mac% incus project switch Handson-lab-1
```

```
mac% incus launch images:ubuntu/24.04 c2 --profile=internal-net --profile=storage -d root,size=5GiB
```

```
mac% incus list
```

NAME	STATE	IPV4	IPV6	TYPE	SNAPSHOTS
c1	RUNNING	10.50.1.181 (eth0)		CONTAINER	0
c2	RUNNING	172.16.10.2 (eth0)		CONTAINER	0

Note the IP address assigned to the container c2 (172.16.10.2). You may recollect that when we created the internal network (internal-net-1) in the previous task, we passed a parameter ‘ipv4.dhcp.ranges=172.16.10.100-172.16.10.200’ to the ‘incus network create’ command. We can now confirm that the ip address assigned to the container (c2) is provided from that range.

```
mac% incus info c2
```

Name: c2

Description:

Status: RUNNING

Type: container

Architecture: x86_64

PID: 88988

Created: 2025/05/20 15:58 UTC

Last Used: 2025/05/20 15:58 UTC

Started: 2025/05/20 15:58 UTC

Resources:

Processes: 12

Disk usage:

root: 1.05MiB

CPU usage:

CPU usage (in seconds): 2

Memory usage:

```

Memory (current): 43.95MiB
Network usage:
eth0:
    Type: broadcast
    State: UP
    Host interface: veth4c60ac20
    MAC address: 10:66:6a:ca:9e:f1
    MTU: 1500
    Bytes received: 688B
    Bytes sent: 1.58kB
    Packets received: 2
    Packets sent: 14
    IP addresses:
        inet: 172.16.10.2/24 (global)
        inet6: fe80::1266:6aff:fea:9ef1/64 (link)
lo:
    Type: loopback
    State: UP
    MTU: 65536
    Bytes received: 0B
    Bytes sent: 0B
    Packets received: 0
    Packets sent: 0
    IP addresses:
        inet: 127.0.0.1/8 (local)
        inet6: ::1/128 (local)

```

Let us get inside the container and run a few commands.

```

mac% incus shell c2
root@c2:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
26: eth0@if27: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group
default qlen 1000
    link/ether 10:66:6a:ca:9e:f1 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.16.10.2/24 metric 100 brd 172.16.10.255 scope global dynamic eth0
        valid_lft 3243sec preferred_lft 3243sec
    inet6 fe80::1266:6aff:fea:9ef1/64 scope link
        valid_lft forever preferred_lft forever

```

```

root@c2:~# apt update -y
root@c2:~# apt upgrade -y
root@c2:~# apt install -y traceroute

root@c2:~# traceroute zabbly.com
traceroute to zabbly.com (45.45.148.7), 30 hops max, 60 byte packets
 1 _gateway (172.16.10.253) 4.353 ms 4.587 ms 4.812 ms
 2 _gateway.incus (10.50.1.253) 6.278 ms 6.290 ms 6.267 ms
 3 _gateway (192.168.1.1) 6.150 ms 5.777 ms 6.046 ms
 4 10.61.209.2 (10.61.209.2) 17.201 ms 10.61.209.3 (10.61.209.3) 17.086 ms 10.61.209.2
(10.61.209.2) 23.465 ms
..
[ removed the rest of the output ]

```

To reach zabbly.com (45.45.148.7) the network traffic goes to internal-net-1 first (172.16.10.253). The next router it reaches is the external-net-1 (10.50.1.253). Then finally from there it goes to 192.168.1.1 which is the ip address of the router on the physical network.

It is important to note that the network traffic did NOT go through incus network bridge (incusbr0), which was our intention when we designed the network topology in the previous task.

We can see from the web UI that currently we have 2 containers c1 and c2 running inside the project Handson-lab-1. container c1 is connected to the external-net-1 and c2 is connected to internal-net-1

The screenshot shows the Incus UI interface. On the left, a sidebar menu lists 'Project' (Handson-lab-1), 'Instances' (2), 'Profiles', 'Networking', 'Storage', 'Images', 'Configuration', 'Operations', 'Warnings', and 'Settings'. A red circle with the number '1' is on the 'Project' dropdown. A red circle with the number '2' is on the 'Instances' link. At the bottom of the sidebar, there are links for 'c88b6a7dd7e8fa5f...', 'Documentation', 'Discussion', and 'Report a bug'. The main area is titled 'Instances' with a search bar and a 'Create instance' button. It displays a table with the following data:

	NAME	TYPE	DESCRIPTION	IPV4	STATUS
<input type="checkbox"/>	c1	Container		10.50.1.181	Running
<input type="checkbox"/>	c2	Container		172.16.10.2	Running

Below the table, it says 'Showing all 2 instances'. There are navigation arrows, a page size selector ('50/page'), and a refresh icon. At the bottom right, it says '1 operation in progress...'.

As we can see, containers c1 and c2 are on different networks. Now let us test if we can ping from c2 to c1.

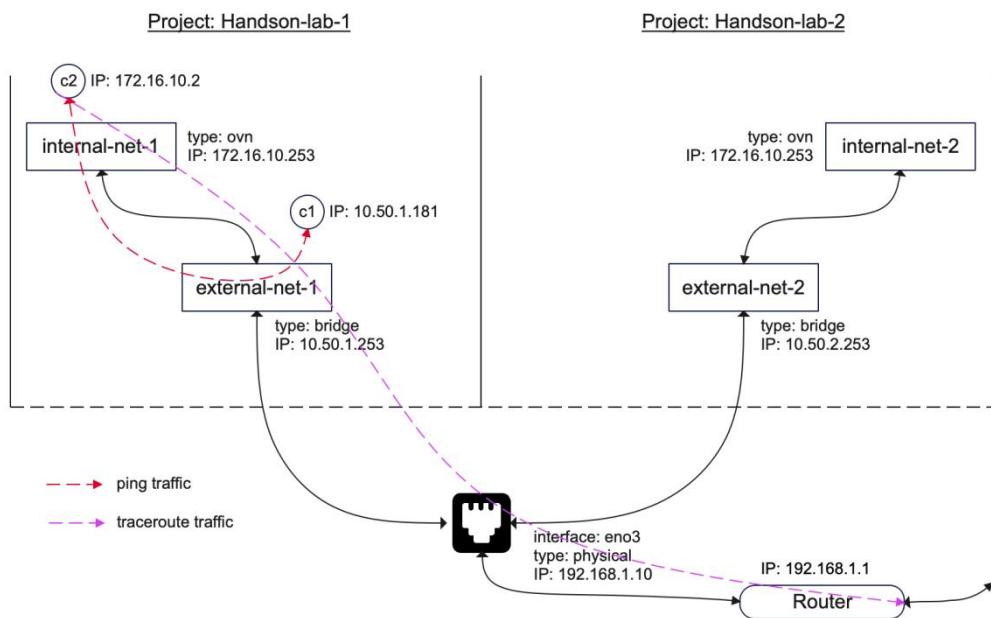
```

root@c2:~# ping c1
PING c1.incus (10.50.1.181) 56(84) bytes of data.
64 bytes from c1.incus (10.50.1.181): icmp_seq=1 ttl=63 time=6.48 ms
64 bytes from c1.incus (10.50.1.181): icmp_seq=2 ttl=63 time=0.634 ms
64 bytes from c1.incus (10.50.1.181): icmp_seq=3 ttl=63 time=0.113 ms
^C
--- c1.incus ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 0.113/2.407/6.476/2.884 ms
root@c2:~#

```

This proves network traffic works fine as expected between the internal-net-1 and external-net-1.

And here is a diagrammatic representation of the network traffic flow of both the tests that we have conducted so far - ping and traceroute test.



The red colored arrow shows the network traffic flow of the ping test. And the purple colored arrow shows the network traffic flow of the traceroute test.

Task #24: Creating of a profile for a desktop virtual machine

In this task, let us create or profile for a desktop virtual machine defining its configurations. In the next task we will look at how to use this profile. Run the following command to create a new profile named ‘desktop’.

```
mac% incus project switch Handson-lab-1
mac% incus profile create desktop
mac% incus profile edit desktop
config:
  limits.memory: 8GiB
  limits.cpu: 2
  security.nesting: "true"
description: desktop profile or project Handson-lab-1
devices:
  eth0:
    name: eth0
    network: external-net-1
    type: nic
  eth1:
    name: eth1
    network: internal-net-1
    type: nic
name: desktop
used_by: []
project: Handson-lab-1
```

The profile is self-explanatory. It defines the memory (8GiB), CPU count=2, supports virtualization nesting, defines 2 network interfaces. The first network interface (eth0) will be connected to the external network (external-net-1) and the second network interface (eth1) will be connected to internal network (internal-net-1).

Note that we did not define anything with regards to the storage. As you may recall we have already created a profile for storage before (aptly named ‘storage’) which defines which storage pool needs to be used. Here it is for a quick reference:

```
mac$ incus profile show storage
config: {}
description: storage profile for project Handson-lab-1
devices:
  root:
    path: /
    pool: storage-1
```

```
type: disk
name: storage
used_by:
- /1.0/instances/c1?project=Handson-lab-1
- /1.0/instances/c2?project=Handson-lab-1
project: Handson-lab-1
```

Also note that ‘storage’ profile only defines the pool (storage-1) where the root disk will be stored. But it does not define the size of the root disk the virtual machine that will eventually be allocated to the virtual machine. We will see in the next task as where and how the size of the root disk will be defined.

Task #25: Creation of desktop virtual machine

In this task, let us get into the action of launching an ubuntu desktop leveraging the profiles that we have created before - namely ‘storage’ and ‘desktop’ profiles.

Let us launch an ubuntu desktop passing the profiles as parameters to the incus launch command line..

```
mac% incus project switch Handson-lab-1

mac% incus launch images:ubuntu/plucky/desktop ubuntu-desktop-vm --profile=storage
--profile=desktop -d root,size=20GiB --vm

mac% incus list
```

NAME	STATE	IPV4	IPV6	TYPE	SNAPSHOTS
c1	RUNNING	10.50.1.181 (eth0)		CONTAINER	0
c2	RUNNING	172.16.10.2 (eth0)		CONTAINER	0
ubuntu-desktop-vm	RUNNING	172.16.10.3 (enp6s0)	10.50.1.110 (enp5s0)	VIRTUAL-MACHINE	0

The first thing to notice are the ip addresses of the two network interfaces and the networks to which they are connected to.

- enp5s0 is the first network interface connected to the external-net-1
- enp6s0 is the second network interface connected to the internal-net-1

The reason being , the profile ‘desktop’ which is passed as the parameter, defines to which networks the two network interfaces should be connected to.

Similarly, the profile ‘storage’ defines to which storage pool (storage-1) the root disk of the desktop virtual machine should be stored.

The same configuration can be seen using web UI

Incus UI

Project: Handson-lab-1

Instances ①

Showing all 3 instances

NAME	TYPE	DESCRIPTION	IPV4	STATUS
c1	Container		10.50.1.181	Running
c2	Container		172.16.10.2	Running
ubuntu-desktop-vm	VM		2 addresses	Running

+ Create instance

1 Instances ②

2 Profiles

Networking >

Storage >

Images

Configuration

Operations

Warnings

Settings

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

<< Version 6.12-ui-0.15 2 ^

This screenshot shows the Incus UI interface. On the left is a sidebar with project settings and navigation links. The main area displays a table of instances, showing their names, types, descriptions, IPv4 addresses, and statuses. There are buttons for creating a new instance and managing existing ones.

Incus UI

Project: Handson-lab-1

Instances / ubuntu-desktop-vm

+ Create Image | Duplicate | Export | Delete

Overview Configuration Snapshots Terminal Console Logs

Memory: 1.7 GiB of 7.2 GiB memory used

Disk: 45.7 MiB of 20.0 GiB disk used

Networks

NAME	INTERFACE	TYPE	MANAGED
external-net-1	eth0	bridge	Yes
internal-net-1	eth1	ovn	Yes

Devices

NAME	TYPE	DETAILS
root	disk (root)	Pool storage-1

Profiles

NAME	DESCRIPTION
storage	storage profile for project Handson-lab-1
desktop	desktop profile for project Handson-lab-1

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

<< Version 6.12-ui-0.15 2 ^

This screenshot shows the detailed configuration for the 'ubuntu-desktop-vm' instance. It includes sections for memory and disk usage, a table of network interfaces, a table of devices, and a table of profiles. The 'storage' and 'desktop' profiles are highlighted with red boxes.

Incus UI

Instances / ubuntu-desktop-vm

Running | ↻ | II | □

+ Create Image | Duplicate | Export | Delete

Project: Handson-lab-1

Overview Configuration Snapshots Terminal Console Logs

Main configuration

Devices

CONFIGURATION	INHERITED	OVERRIDE
Root storage		X
Pool	storage-1 From: storage-profile	storage-1
Size	unlimited From: storage-profile	20GiB

Root storage

+ Attach disk device

YAML Configuration

c88b6a7dd7e8fa5f...

Documentation Discussion Report a bug

Version 6.12-ui-0.15

2 ^

Incus UI

Instances / ubuntu-desktop-vm

Running | ↻ | II | □

+ Create Image | Duplicate | Export | Delete

Project: Handson-lab-1

Overview Configuration Snapshots Terminal Console Logs

Main configuration

Devices

CONFIGURATION	INHERITED	OVERRIDE
eth0	external-net-1 From: desktop profile	O
eth1	internal-net-1 From: desktop profile	O

+ Attach network

YAML Configuration

c88b6a7dd7e8fa5f...

Documentation Discussion Report a bug

Version 6.12-ui-0.15

2 ^

Incus UI

Instances / ubuntu-desktop-vm

Running | ⏪ | ⏹ | ⏸ | ⏷

Configuration

Main configuration	CONFIGURATION	INHERITED	OVERRIDE
Exposed CPU limit	2 From: desktop profile	Inherited	<input type="button" value="Override"/>
Memory limit	8GB From: desktop profile	Inherited	<input type="button" value="Override"/>
Memory swap (Containers only)	Allow From: Incus	Inherited	<input type="button" value="Override"/>
Disk priority	5 From: Incus	Inherited	<input type="button" value="Override"/>
Max number of processes (Containers only)	- From: Incus	Inherited	<input type="button" value="Override"/>

YAML Configuration

Version 6.12-ui-0.15

Incus UI

Instances / ubuntu-desktop-vm

Running | ⏪ | ⏹ | ⏸ | ⏷

Configuration

Main configuration	CONFIGURATION	INHERITED	OVERRIDE
Protect deletion	No From: Incus	Inherited	<input type="button" value="Override"/>
Privileged (Containers only)	Deny From: Incus	Inherited	<input type="button" value="Override"/>
Nesting (Containers only)	Allow From: desktop profile	Inherited	<input type="button" value="Override"/>
Protect UID/GID shift (Containers only)	No From: Incus	Inherited	<input type="button" value="Override"/>
Base host id (Containers only)	- From: Incus	Inherited	<input type="button" value="Override"/>

YAML Configuration

Version 6.12-ui-0.15

Let us run the same sanity tests that we have done before.

```
mac% incus shell ubuntu-desktop-vm
root@ubuntu-desktop-vm:~# apt update -y
root@ubuntu-desktop-vm:~# apt upgrade -y
root@ubuntu-desktop-vm:~# apt install traceroute -y
```

```
root@ubuntu-desktop-vm:~# ping c1
PING c1.incus (10.50.1.181) 56(84) bytes of data.
64 bytes from c1.incus (10.50.1.181): icmp_seq=1 ttl=64 time=0.403 ms
64 bytes from c1.incus (10.50.1.181): icmp_seq=2 ttl=64 time=0.314 ms
64 bytes from c1.incus (10.50.1.181): icmp_seq=3 ttl=64 time=0.319 ms
```

```

64 bytes from c1.incus (10.50.1.181): icmp_seq=4 ttl=64 time=0.315 ms
^C
--- c1.incus ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3062ms
rtt min/avg/max/mdev = 0.314/0.337/0.403/0.037 ms

```

```

root@ubuntu-desktop-vm:~# ping c2
PING c2.incus (172.16.10.2) 56(84) bytes of data.
64 bytes from 172.16.10.2: icmp_seq=1 ttl=64 time=1.95 ms
64 bytes from 172.16.10.2: icmp_seq=2 ttl=64 time=0.247 ms
64 bytes from 172.16.10.2: icmp_seq=3 ttl=64 time=0.227 ms
64 bytes from 172.16.10.2: icmp_seq=4 ttl=64 time=0.281 ms
^C
--- c2.incus ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3037ms
rtt min/avg/max/mdev = 0.227/0.676/1.952/0.736 ms

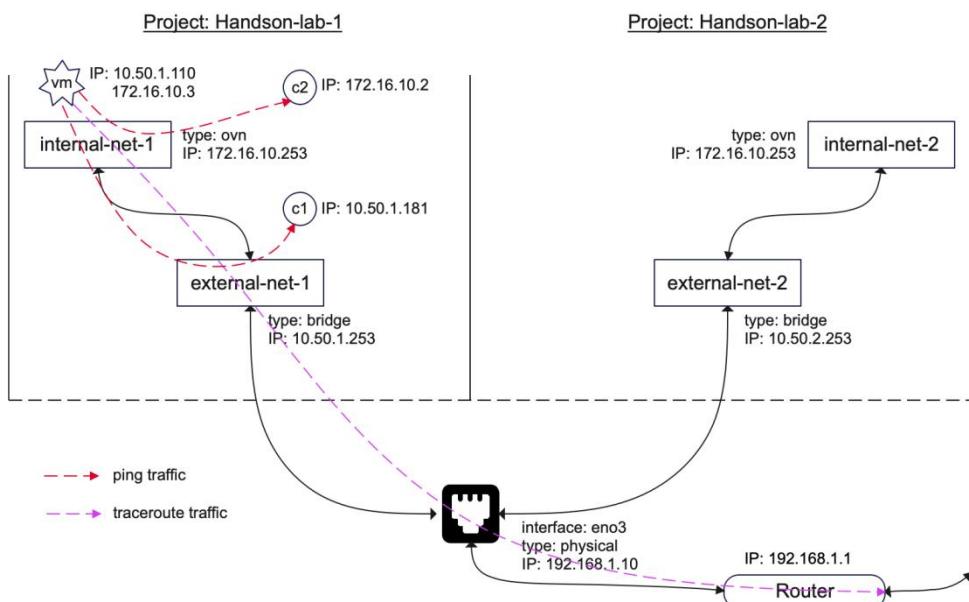
```

```

root@ubuntu-desktop-vm:~# traceroute zabbly.com
traceroute to zabbly.com (45.45.148.7), 30 hops max, 60 byte packets
 1 _gateway.incus (10.50.1.253) 1.139 ms 1.038 ms 1.008 ms
 2 _gateway (192.168.1.1) 0.973 ms 32.912 ms 32.905 ms
 3 10.61.209.2 (10.61.209.2) 12.192 ms 21.238 ms 21.104 ms
...
[ removed the rest of the output ]

```

Here is diagrammatic representation of the ping tests



There are two ways we can access the console of this virtual machine

- Method #1: web UI
- Method #2: From Mac Laptop using command line

Method #1: web UI

The screenshot shows the Incus UI interface. On the left, a sidebar menu is open under the 'Project' section for 'Handson-lab-1'. The 'Instances' option is highlighted with a red circle labeled '1'. Below it, the 'ubuntu-desktop-vm' instance is also highlighted with a red circle labeled '3'. The main content area displays a table of instances:

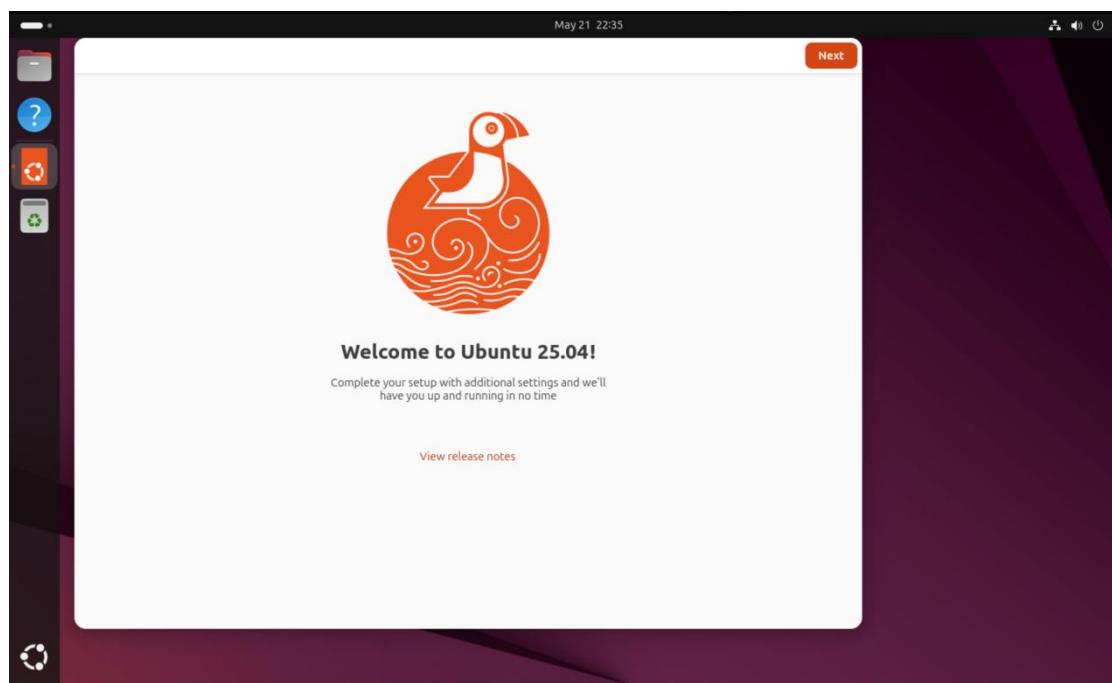
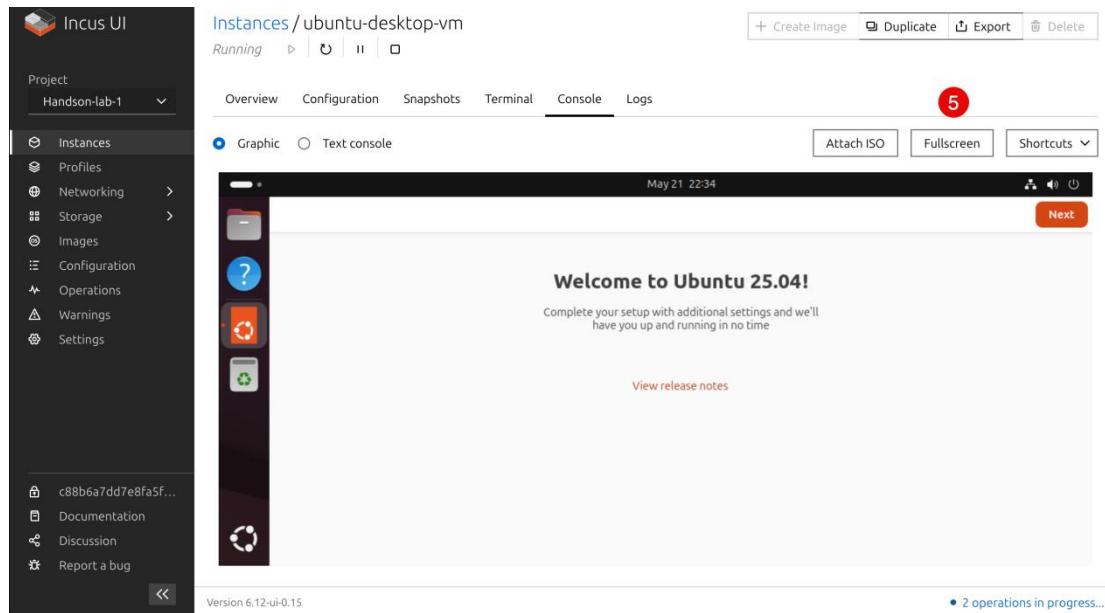
NAME	TYPE	DESCRIPTION	IPV4	STATUS
c1	Container		10.50.1.181	• Running
c2	Container		172.16.10.2	• Running
ubuntu-desktop-vm	VM		2 addresses	• Running

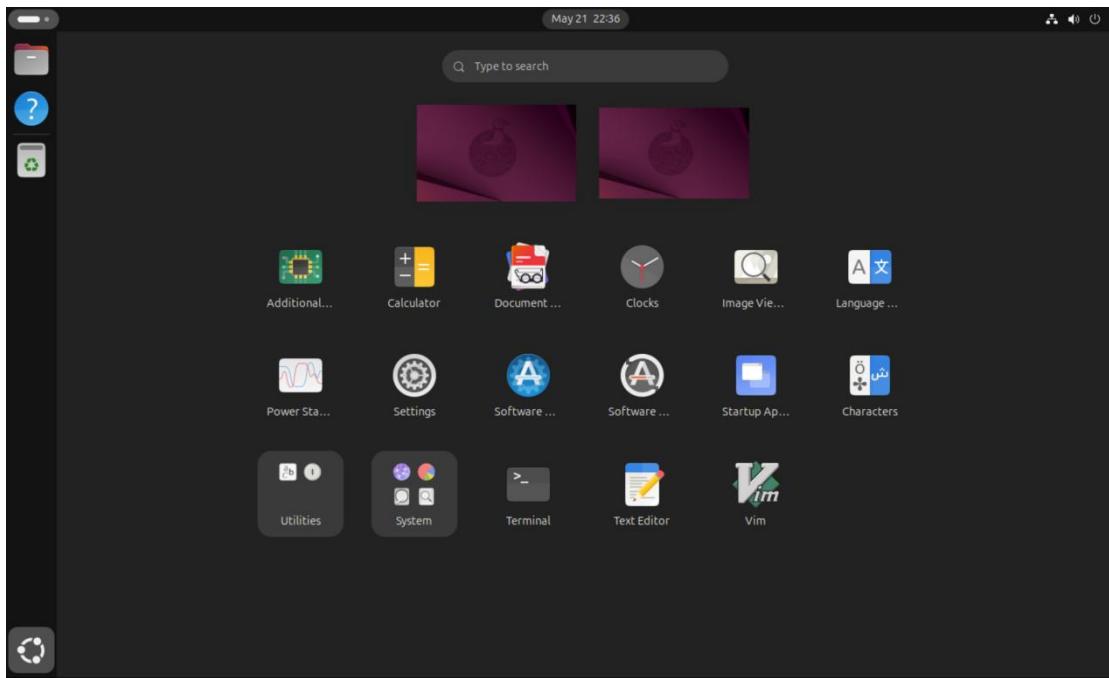
At the bottom right of the main area, there is a message: '• 1 operation in progress...'.

The screenshot shows the detailed view for the 'ubuntu-desktop-vm' instance. The top navigation bar includes options for 'Create Image', 'Duplicate', 'Export', and 'Delete'. The main content is organized into several sections:

- Usage:** CPU Time(s) 498.98, Memory 2.0 GiB of 7.2 GiB memory used, Disk 314.3 MiB of 20.0 GiB disk used.
- Networks:** Lists external-net-1 (eth0, bridge, Yes) and internal-net-1 (eth1, ovn, Yes).
- Devices:** Lists root (disk (root), Pool, storage-1).
- Profiles:** Lists storage (storage profile for project Handson-lab-1).

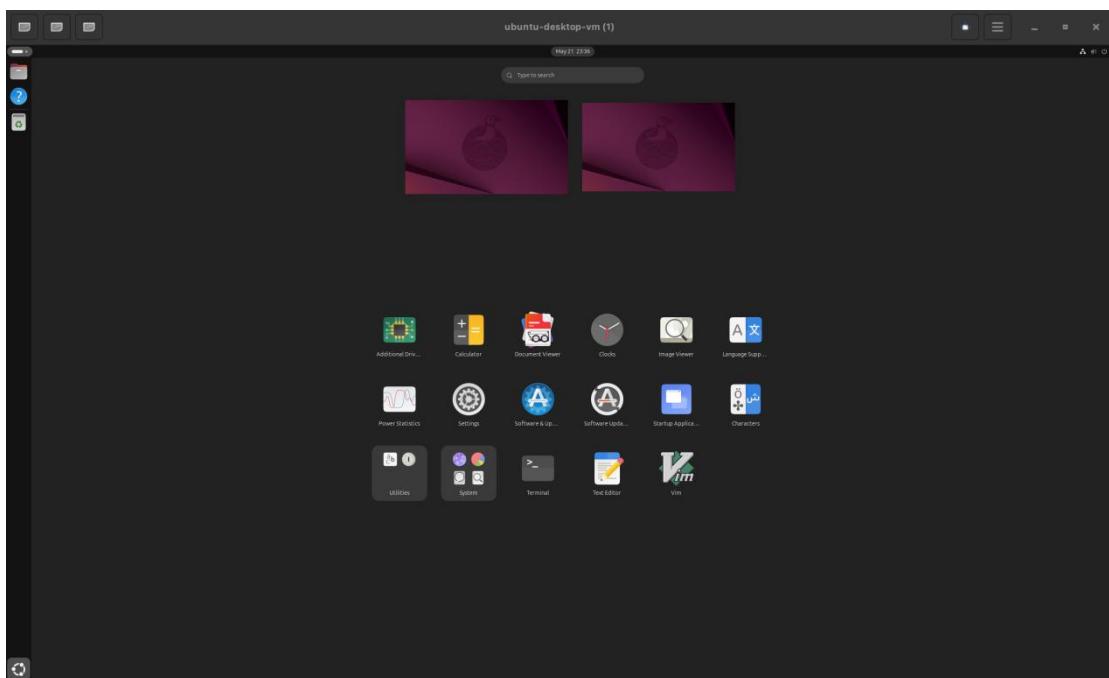
At the bottom right of the main area, there is a message: '• 1 operation in progress...'.





Method #2: From Mac Laptop using command line

```
mac% incus project switch Handson-lab-1  
mac% incus console ubuntu-desktop-vm --type=vga
```



Task #26: Creation of project: Handson-lab-2

Let us take a look at how project Handson-lab-1 has evolved so far with regards to all the networks, profile, storage pool and all the virtual machine and container instances that we have launched so far.

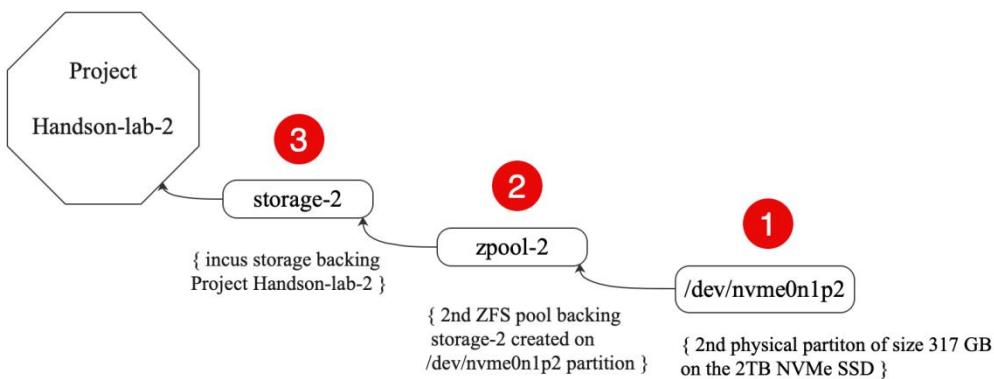
We are going to follow the same procedures to create a second project Handson-lab-2. Once we have it in place, we will understand how to segregate the networks among projects, which is a crucial concept to understand to attain multi-tenancy.

```
mac% incus project show Handson-lab-1
config:
  features.images: "true"
  features.networks: "false"
  features.networks.zones: "false"
  features.profiles: "true"
  features.storage.buckets: "true"
  features.storage.volumes: "true"
  restricted: "true"
  restricted.containers.lowlevel: allow
  restricted.containers.nesting: allow
  restricted.containers.privilege: allow
  restricted.devices.disk: managed
  restricted.devices.gpu: allow
  restricted.snapshots: allow
  restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-1
name: Handson-lab-1
used_by:
- /1.0/instances/c1?project=Handson-lab-1
- /1.0/instances/c2?project=Handson-lab-1
- /1.0/instances/ubuntu-desktop-vm?project=Handson-lab-1
-
/1.0/images/6b4ffc77179ffc19b47d4fe9044d3e6b81408ddb8ec849fdc2c2b256c5a55a0d?project=Ha
ndson-lab-1
-
/1.0/images/c9679e4c68822195ef69182e8b97261ef534a18f48e72d93ba15a0ffb3a0f835?project=Ha
ndson-lab-1
- /1.0/profiles/default?project=Handson-lab-1
- /1.0/profiles/desktop?project=Handson-lab-1
- /1.0/profiles/external-net?project=Handson-lab-1
- /1.0/profiles/internal-net?project=Handson-lab-1
- /1.0/profiles/storage?project=Handson-lab-1
```

To create a new project named Handson-lab-2, we need to repeat from task #17 to task #23. (I am skipping to showcase the same document contents once again as it will be redundant).

Here is the overall design for project Handson-lab-2, which is almost similar to project Handson-lab-1

Storage layout for Project Handson-lab-2



Here are the details after creation of project Handson-lab-2

```
mac% incus project switch Handson-lab-2
```

Project Handson-lab-2

```
mac% incus project show Handson-lab-2
config:
```

```
features.images: "true"
features.networks: "false"
features.networks.zones: "false"
features.profiles: "true"
features.storage.buckets: "true"
features.storage.volumes: "true"
restricted: "true"
restricted.containers.lowlevel: allow
restricted.containers.nesting: allow
restricted.containers.privilege: allow
restricted.devices.disk: managed
restricted.devices.gpu: allow
restricted.snapshots: allow
```

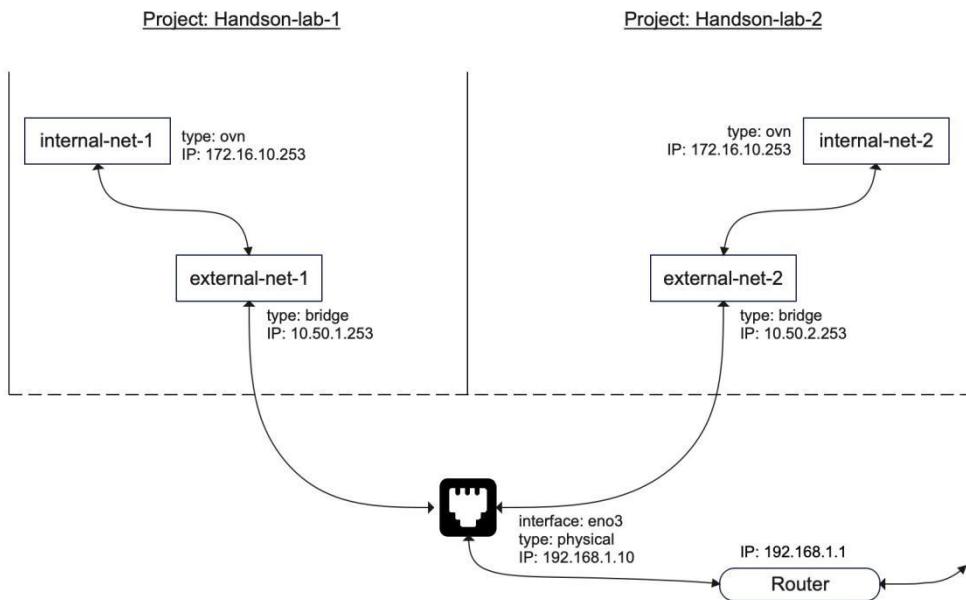
```
restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-2
name: Handson-lab-2
used_by:
- /1.0/profiles/default?project=Handson-lab-2
```

default profile

```
mac% incus profile show default
config:
  security.nesting: "true"
  security.privileged: "true"
description: Default Incus profile for project Handson-lab-2
devices:
  root:
    path: /
    pool: storage-2
    type: disk
name: default
used_by: []
project: Handson-lab-2
Storage profile
```

```
mac% incus storage show storage
config: {}
description: storage profile for project Handson-lab-2
devices:
  root:
    path: /
    pool: storage-2
    type: disk
name: storage
used_by: []
project: Handson-lab-2
```

Network Architecture for project Handson-lab-2



External network - external-net-2

```
mac% incus network show external-net-2
config:
  ipv4.address: 10.50.2.253/24
  ipv4.dhcp.ranges: 10.50.2.101-10.50.2.200
  ipv4.nat: "true"
  ipv6.address: none
description: ""
name: external-net-2
type: bridge
used_by: []
managed: true
status: Created
locations:
- none
project: default
```

Profile for external-net-2

```
mac% incus profile show external-net
config: {}
description: external-net profile for project Handson-lab-2
devices:
```

```
eth0:  
  name: eth0  
  network: external-net-2  
  type: nic  
name: external-net  
used_by: []  
project: Handson-lab-2
```

Internal network - internal-net-2

```
mac% incus network show internal-net-2  
config:  
  bridge.mtu: "1500"  
  ipv4.address: 172.16.10.253/24  
  ipv4.nat: "true"  
  ipv6.address: none  
  network: external-net-2  
  volatile.network.ipv4.address: 10.50.2.1  
description: routes via external-net-2  
name: internal-net-2  
type: ovn  
used_by:  
- /1.0/profiles/internal-net?project=Handson-lab-2  
managed: true  
status: Created  
locations:  
- none  
project: default
```

Profile for internal-net-2

```
mac% incus profile show internal-net  
config: {}  
description: internal network profile for project Handson-lab-2  
devices:  
  eth0:  
    name: eth0  
    network: internal-net-2  
    type: nic  
name: internal-net  
used_by: []  
project: Handson-lab-2
```

```
mac% incus profile list
```

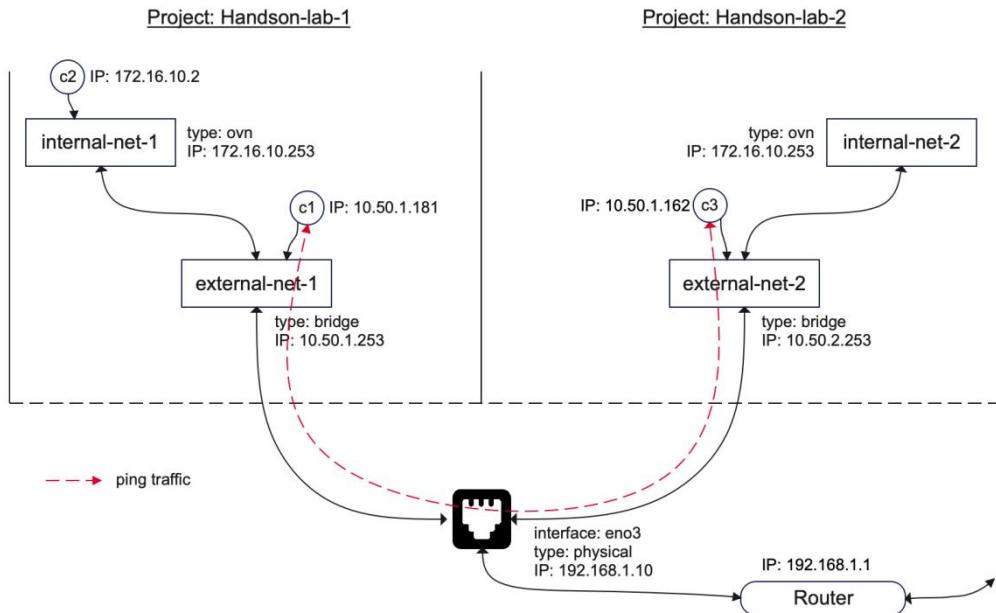
NAME	DESCRIPTION	USED BY
default	Default Incus profile for project Handson-lab-2	0
external-net	external-net profile for project Handson-lab-2	0
internal-net	internal network profile for project Handson-lab-2	0
storage	storage profile for project Handson-lab-2	0

```
mac% incus network list
```

NAME	TYPE	MANAGED	IPV4	IPV6	DESCRIPTION	USED BY	STATE
external-net-1	bridge	YES	10.50.1.253/24	none	uplink for Handson-lab-1	5	CREATED
external-net-2	bridge	YES	10.50.2.253/24	none	uplink for Handson-lab-2	2	CREATED
incusbr0	bridge	YES	10.1.1.253/24	none	routes traffic via eno3	1	CREATED
internal-net-1	ovn	YES	172.16.10.253/24	none	routes via external-net-1	4	CREATED
internal-net-2	ovn	YES	172.16.10.253/24	none	routes via external-net-2	1	CREATED

Task #27: Network test between projects - Take #1

Now that we have two projects in place namely Handson-lab-1 and Handson-lab-2, let us try to do a simple ping test across projects. We already have a few instances running in project Handson-lab-1. Let us launch a container instance in project Handson-lab-2 and see if we can ping a container instance in project Handson-lab-1.



Let us launch an ubuntu container (c3) connecting it to external-net-2

```
mac% incus launch images:ubuntu/24.04 c3 --profile=external-net --profile=storage -d root,size=5GiB
```

```
% incus list
+-----+-----+-----+-----+-----+
| NAME | STATE | IPV4 | IPV6 | TYPE | SNAPSHOTS |
+-----+-----+-----+-----+-----+
| c3 | RUNNING | 10.50.2.162 (eth0) | | CONTAINER | 0 |
+-----+-----+-----+-----+-----+
```

Let us get inside the container c3 and see if we can reach a host on the internet by pinging zabbix.com

```
mac% incus shell c3
root@c3:~# ping zabbix.com
PING zabbix.com (45.45.148.7) 56(84) bytes of data.
64 bytes from rproxy.dcmtl.stgraber.org (45.45.148.7): icmp_seq=1 ttl=49 time=84.3 ms
```

```
64 bytes from rproxy.dcctl.stgraber.org (45.45.148.7): icmp_seq=2 ttl=49 time=85.1 ms
64 bytes from rproxy.dcctl.stgraber.org (45.45.148.7): icmp_seq=3 ttl=49 time=81.6 ms
^C
--- zabbly.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 81.600/83.682/85.147/1.512 ms
root@c3:~#
```

We will ping from the container (c3, IP: 10.50.2.162) running inside project Handson-lab-2 to the container (c1, IP: 10.50.1.181) running inside project Handson-lab-1

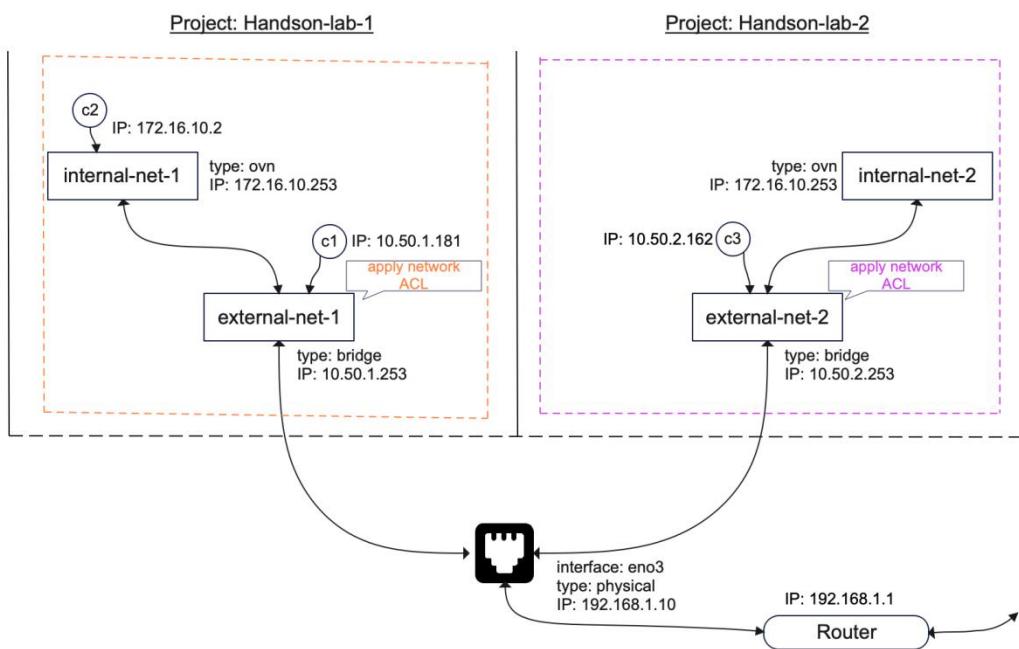
```
root@c3:~# ping 10.50.1.181
PING 10.50.1.181 (10.50.1.181) 56(84) bytes of data.
64 bytes from 10.50.1.181: icmp_seq=1 ttl=63 time=0.158 ms
64 bytes from 10.50.1.181: icmp_seq=2 ttl=63 time=0.100 ms
64 bytes from 10.50.1.181: icmp_seq=3 ttl=63 time=0.097 ms
^C
--- 10.50.1.181 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2045ms
rtt min/avg/max/mdev = 0.097/0.118/0.158/0.028 ms
root@c3:~#
```

Well, what it means is that there is NO network isolation among the Projects in our setup, which is a crucial aspect for multi-tenancy. So let us working on building network isolation in our next task.

Task #28: Setting up network isolation among projects

As we have seen before network traffic is not isolated between our projects, which is needed for a multi-tenancy setup. To accomplish this, we need to write network ACL (access control lists) rules and apply them onto appropriate networks based on the network architecture design. In our case, we need to apply ACL rules to the external networks of both the projects.

The color coding shown below defines the perimeter/boundary of network isolation per project defined by separate network ACL per project.



ACL rule for external-net-2 of project Handson-lab-2

Let us create the ACL rule

```
mac% incus project switch Handson-lab-2
```

```
mac% incus network acl list
+-----+-----+-----+
| NAME | DESCRIPTION | USED BY |
+-----+-----+-----+
```

```
mac% incus list
+-----+-----+-----+-----+-----+-----+
| NAME | STATE | IPV4 | IPV6 | TYPE | SNAPSHOTS |
+-----+-----+-----+-----+-----+-----+
```

```
| c3 | RUNNING | 10.50.2.162 (eth0) |      | CONTAINER |  0      |
+-----+-----+-----+-----+-----+
```

```
mac% incus network acl create isolate-lab-2
mac% incus network acl edit isolate-lab-2
name: isolate-lab-2
description: allow external-net-2 only. reject other external-nets
egress:
- action: allow
  state: enabled
- action: reject
  destination: 10.50.1.0/24
  state: enabled
- action: reject
  destination: 10.50.1.253/24
  state: enabled
ingress:
- action: allow
  state: enabled
- action: reject
  source: 10.50.1.0/24
  state: enabled
config: {}
used_by: []
project: default
```

We are defining the ingress and egress rules. For this project Handson-lab-2, we are mentioning that the external ip address range of project Handson-lab-1 to be explicitly rejected (10.50.1.0/24) and the rest are accepted. And we will do the reverse for project Handson-lab-1.

In this way, the external-net-2 that is assigned to project Handson-lab-2 will reject any traffic from external-net-1 assigned to Handson-lab-1.

So far we have only created the ACL rule. The next step is to apply this rule to the external-net-2.

```
mac% incus network edit external-net-2
config:
  ipv4.address: 10.50.2.253/24
  ipv4.dhcp.ranges: 10.50.2.101-10.50.2.200
  ipv4.nat: "true"
  ipv4.ovn.ranges: 10.50.2.1-10.50.2.25
  ipv4.routing: "true"
  ipv6.address: none
```

```

security.acls: isolate-lab-2 --> Added this line
description: uplink for Handson-lab-2
name: external-net-2
type: bridge
used_by:
- /1.0/instances/c3?project=Handson-lab-2
- /1.0/networks/internal-net-2
- /1.0/profiles/external-net?project=Handson-lab-2
managed: true
status: Created
locations:
- none
project: default

```

Let us take a look at how the newly created ACL rule shows up in web UI.

NAME	DESCRIPTION	USED BY
isolate-lab-2	allow external-net-2 only. reject other external-nets	1

Incus UI

Project: Handson-lab-2

Network ACLs / isolate-lab-2

Overview Configuration Ingress Egress

General

Name	isolate-lab-2
Description	allow external-net-2 only, reject other external-nets

Usage (1)

Networks (1)	default / external-net-2
Instances (0)	-
Profiles (0)	-

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

Version 6.12-ui-0.15

Delete ACL

Incus UI

Project: Handson-lab-2

Network ACLs / isolate-lab-2

4

Overview Configuration **Ingress** Egress

Create rule

ACTION	STATE	DESCRIPTION	SOURCE	DESTINATION
allow	enabled			
reject	enabled		10.50.1.0/24	

c88b6a7dd7e8fa5f... Documentation Discussion Report a bug

Version 6.12-ui-0.15

Delete ACL

ACTION	STATE	DESCRIPTION	SOURCE	DESTINATION
allow	enabled			
reject	enabled		10.50.1.0/24	
reject	enabled		10.50.1.253/24	

Now that we have completed applying a network ACL rule for external-net-2 that is backing project Handson-lab-2, let us do the same for external-net-1 that is backing project Handson-lab-1.

ACL rule for external-net-1 of project Handson-lab-1

Let us create the ACL rule

```
mac% incus project switch Handson-lab-1
```

```
mac% incus network acl list
```

NAME	DESCRIPTION	USED BY
isolate-lab-2	allow external-net-2 only. reject other external-nets	1

```
mac% incus network acl edit isolate-lab-1
```

```
name: isolate-lab-1
```

```
description: allow external-net-1 only. reject other external-nets
```

```
egress:
```

```
- action: allow
  state: enabled
- action: reject
  destination: 10.50.2.0/24
  state: enabled
- action: reject
  destination: 10.50.2.253/24
```

```

state: enabled
ingress:
- action: allow
  state: enabled
- action: reject
  source: 10.50.2.0/24
  state: enabled
config: {}
used_by: []
project: default

```

Let us apply this ACL onto external-net-1 which serves project Haandson-lab-1

```

mac% incus network edit external-net-1
config:
  ipv4.address: 10.50.1.253/24
  ipv4.dhcp.ranges: 10.50.1.101-10.50.1.200
  ipv4.nat: "true"
  ipv4.ovn.ranges: 10.50.1.1-10.50.1.25
  ipv4.routing: "true"
  ipv6.address: none
  security.acls: isolate-lab-1 --> Added this line
description: uplink for Handson-lab-1
name: external-net-1
type: bridge
used_by:
- /1.0/instances/c1?project=Handson-lab-1
- /1.0/instances/ubuntu-desktop-vm?project=Handson-lab-1
- /1.0/networks/internal-net-1
- /1.0/profiles/desktop?project=Handson-lab-1
- /1.0/profiles/external-net?project=Handson-lab-1
managed: true
status: Created
locations:
- none
project: default

```

```
mac% incus network acl list
```

NAME	DESCRIPTION	USED BY
isolate-lab-1	allow external-net-1 only. reject other external-nets	0
isolate-lab-2	allow external-net-2 only. reject other external-nets	1

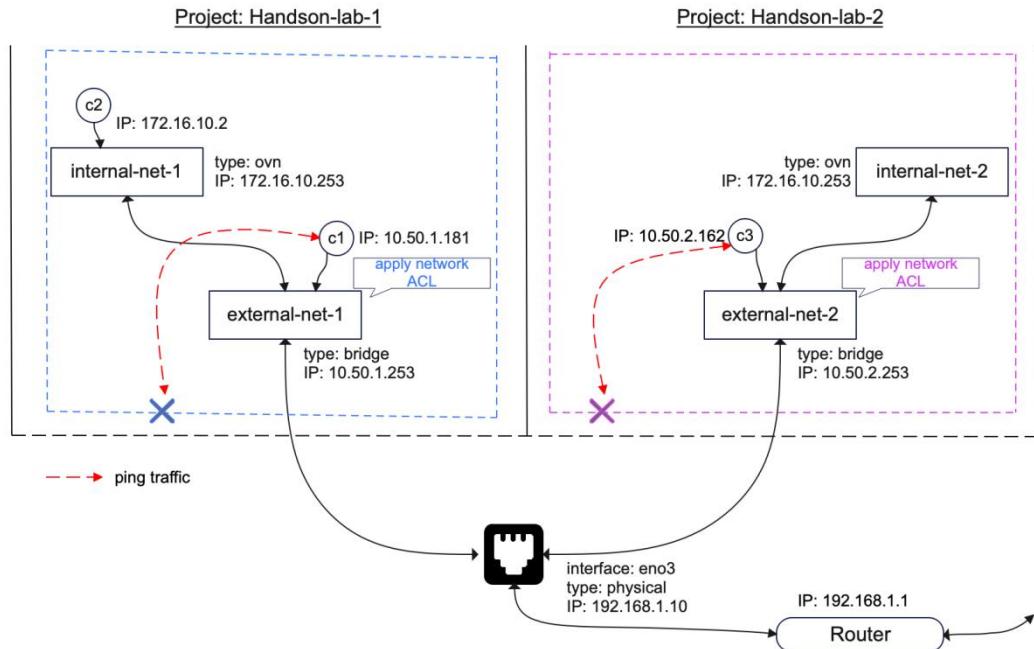
```
+-----+-----+-----+
```

Now that we have two network ACL rules in place, one (isolate-lab-1) is applied to external-net-1 and the other (isolate-lab-2) is to external-net-2.

Let us repeat the ping test again and see system's behavior.

Task #29: Network test between projects - Take #2

Let us see the outcome of the ping test once again, now that we have network ACLs in each project in place.



Let us access project Handson-lab-2 first.

```
mac% incus project switch Handson-lab-2
```

```
mac% incus list
```

NAME	STATE	IPV4	IPV6	TYPE	SNAPSHOTS
c3	RUNNING	10.50.2.162 (eth0)		CONTAINER	0

Let us get into the container (c3) and try to ping container (c1) running inside the project Handson-lab-1.

```
mac% incus shell c3
root@c3:~# ping 10.50.1.181
PING 10.50.1.181 (10.50.1.181) 56(84) bytes of data.
From 10.50.2.253 icmp_seq=1 Destination Port Unreachable
From 10.50.2.253 icmp_seq=2 Destination Port Unreachable
```

```

From 10.50.2.253 icmp_seq=3 Destination Port Unreachable
^C
--- 10.50.1.181 ping statistics ---
3 packets transmitted, 0 received, +3 errors, 100% packet loss, time 2027ms

```

We can clearly see the ping fails because of the network ACL in place in the project Handson-lab-1. Let us see if we can ping a host on the internet.

```

root@c3:~# ping zabbly.com
PING zabbly.com (45.45.148.7) 56(84) bytes of data.
64 bytes from rproxy.dcmtl.stgraber.org (45.45.148.7): icmp_seq=1 ttl=49 time=93.0 ms
64 bytes from rproxy.dcmtl.stgraber.org (45.45.148.7): icmp_seq=2 ttl=49 time=93.7 ms
64 bytes from rproxy.dcmtl.stgraber.org (45.45.148.7): icmp_seq=3 ttl=49 time=94.3 ms
^C
--- zabbly.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2001ms
rtt min/avg/max/mdev = 93.036/93.665/94.273/0.505 ms
root@c3:~#

```

This proves that the network isolation is only the context of reaching out any container or virtual machine inside Handson-lab-1.

Let us do the other way around as well. Let us access project Handson-lab-1

```
mac% incus project switch Handson-lab-1
```

```
mac% incus list
```

NAME	STATE	IPV4	IPV6	TYPE	SNAPSHOTS
c1	RUNNING	10.50.1.181 (eth0)		CONTAINER	0
c2	RUNNING	172.16.10.2 (eth0)		CONTAINER	0
ubuntu-desktop-vm	RUNNING	172.16.10.3 (enp6s0)		VIRTUAL-MACHINE	0
		10.50.1.110 (enp5s0)			

Let us get inside the container (c2) and try to ping container (c3) running inside project Handson-lab-2

```
% incus shell c2
root@c2:~# ping 10.50.2.162
```

```

PING 10.50.2.162 (10.50.2.162) 56(84) bytes of data.
From 10.50.1.253 icmp_seq=1 Destination Port Unreachable
From 10.50.1.253 icmp_seq=2 Destination Port Unreachable
From 10.50.1.253 icmp_seq=3 Destination Port Unreachable
^C
--- 10.50.2.162 ping statistics ---
3 packets transmitted, 0 received, +3 errors, 100% packet loss, time 2052ms

```

We see that the network traffic is blocked by the network ACL applied on external-net-2 (of project Handson-lab-2)

But if we try to ping another virtual machine running inside project Handson-lab-1 that is connected to the same external-net-1, the ping is successful.

```

root@c2:~# ping 10.50.1.110
PING 10.50.1.110 (10.50.1.110) 56(84) bytes of data.
64 bytes from 10.50.1.110: icmp_seq=1 ttl=63 time=7.22 ms
64 bytes from 10.50.1.110: icmp_seq=2 ttl=63 time=0.880 ms
64 bytes from 10.50.1.110: icmp_seq=3 ttl=63 time=0.311 ms
^C
--- 10.50.1.110 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2012ms
rtt min/avg/max/mdev = 0.311/2.803/7.218/3.130 ms

```

Finally let us make sure that we can ping a host on the internet, to prove that the network ACLs are applied only to deny each other's network traffic (from the individual projects point of view) and not to the outside world.

```

root@c2:~# ping zabbly.com
PING zabbly.com (45.45.148.7) 56(84) bytes of data.
64 bytes from rproxy.dcctl.stgraber.org (45.45.148.7): icmp_seq=1 ttl=48 time=95.9 ms
64 bytes from rproxy.dcctl.stgraber.org (45.45.148.7): icmp_seq=2 ttl=48 time=91.8 ms
64 bytes from rproxy.dcctl.stgraber.org (45.45.148.7): icmp_seq=3 ttl=48 time=90.5 ms
^C
--- zabbly.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 90.502/92.744/95.911/2.302 ms

```

Task #30: A fully built multi-tenant system with 6 tenants

In this task we will look at a fully built multi-tenant system with 6 tenants (Handson-lab-1 to Handson-lab-6), with an internal and external network each. Each having it's own storage pool (storage-1 to storage-6).

The way we need to build projects Handson-lab-1 to Handson-lab-6 is exactly how we built Handson-lab-1 and Handson-lab-2. So there is no value in repeating the same documentation again. Instead, we will take a look at how all the 6 project looks like at the end. This output of this can be used as a reference.

```
mac% incus project show Handson-lab-1
config:
  features.images: "true"
  features.networks: "false"
  features.networks.zones: "false"
  features.profiles: "true"
  features.storage.buckets: "true"
  features.storage.volumes: "true"
  restricted: "true"
  restricted.containers.lowlevel: allow
  restricted.containers.nesting: allow
  restricted.containers.privilege: allow
  restricted.devices.disk: managed
  restricted.devices.gpu: allow
  restricted.snapshots: allow
  restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-1
name: Handson-lab-1
used_by:
- /1.0/instances/c1?project=Handson-lab-1
- /1.0/instances/c2?project=Handson-lab-1
- /1.0/instances/ubuntu-desktop-vm?project=Handson-lab-1
-
/1.0/images/1684815eba8b8f9d41293b3784bd4b480b432c56cfab9436cdfdc82ce8ded254?project=Ha
ndson-lab-1
-
/1.0/images/4ead11cd2bb106abfea928d8d5d84cdd05f127fa5b3e657acfe1ec9e737d9093?project=Ha
ndson-lab-1
-
/1.0/images/7485692e4b5f05a3b15cef15a9db5e238e269bc9f1ed7fea01a873f479029443?project=Ha
ndson-lab-1
- /1.0/profiles/default?project=Handson-lab-1
```

```
- /1.0/profiles/desktop?project=Handson-lab-1
- /1.0/profiles/external-net?project=Handson-lab-1
- /1.0/profiles/internal-net?project=Handson-lab-1
- /1.0/profiles/storage?project=Handson-lab-1
```

```
mac% incus project show Handson-lab-2
```

```
config:
  features.images: "true"
  features.networks: "false"
  features.networks.zones: "false"
  features.profiles: "true"
  features.storage.buckets: "true"
  features.storage.volumes: "true"
  restricted: "true"
  restricted.containers.lowlevel: allow
  restricted.containers.nesting: allow
  restricted.containers.privilege: allow
  restricted.devices.disk: managed
  restricted.devices.gpu: allow
  restricted.snapshots: allow
  restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-2
name: Handson-lab-2
used_by:
- /1.0/instances/c3?project=Handson-lab-2
-
/1.0/images/7485692e4b5f05a3b15cef15a9db5e238e269bc9f1ed7fea01a873f479029443?project=Ha
ndson-lab-2
- /1.0/profiles/default?project=Handson-lab-2
- /1.0/profiles/external-net?project=Handson-lab-2
- /1.0/profiles/internal-net?project=Handson-lab-2
- /1.0/profiles/storage?project=Handson-lab-2
```

```
mac% incus project show Handson-lab-3
```

```
config:
  features.images: "true"
  features.networks: "false"
  features.networks.zones: "false"
  features.profiles: "true"
  features.storage.buckets: "true"
  features.storage.volumes: "true"
  restricted: "true"
  restricted.containers.lowlevel: allow
  restricted.containers.nesting: allow
```

```
restricted.containers.privilege: allow
restricted.devices.disk: managed
restricted.devices.gpu: allow
restricted.snapshots: allow
restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-3
name: Handson-lab-3
used_by:
- /1.0/profiles/default?project=Handson-lab-3
- /1.0/profiles/external-net?project=Handson-lab-3
- /1.0/profiles/internal-net?project=Handson-lab-3
- /1.0/profiles/storage?project=Handson-lab-3
```

```
mac% incus project show Handson-lab-4
config:
features.images: "true"
features.networks: "false"
features.networks.zones: "false"
features.profiles: "true"
features.storage.buckets: "true"
features.storage.volumes: "true"
restricted: "true"
restricted.containers.lowlevel: allow
restricted.containers.nesting: allow
restricted.containers.privilege: allow
restricted.devices.disk: managed
restricted.devices.gpu: allow
restricted.snapshots: allow
restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-4
name: Handson-lab-4
used_by:
- /1.0/profiles/default?project=Handson-lab-4
- /1.0/profiles/external-net?project=Handson-lab-4
- /1.0/profiles/internal-net?project=Handson-lab-4
- /1.0/profiles/storage?project=Handson-lab-4
```

```
mac% incus project show Handson-lab-5
config:
features.images: "true"
features.networks: "false"
features.networks.zones: "false"
features.profiles: "true"
features.storage.buckets: "true"
```

```

features.storage.volumes: "true"
restricted: "true"
restricted.containers.lowlevel: allow
restricted.containers.nesting: allow
restricted.containers.privilege: allow
restricted.devices.disk: managed
restricted.devices.gpu: allow
restricted.snapshots: allow
restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-5
name: Handson-lab-5
used_by:
- /1.0/profiles/default?project=Handson-lab-5
- /1.0/profiles/external-net?project=Handson-lab-5
- /1.0/profiles/internal-net?project=Handson-lab-5
- /1.0/profiles/storage?project=Handson-lab-5

```

```

mac% incus project show Handson-lab-6
config:
  features.images: "true"
  features.networks: "false"
  features.networks.zones: "false"
  features.profiles: "true"
  features.storage.buckets: "true"
  features.storage.volumes: "true"
  restricted: "true"
  restricted.containers.lowlevel: allow
  restricted.containers.nesting: allow
  restricted.containers.privilege: allow
  restricted.devices.disk: allow
  restricted.devices.gpu: allow
  restricted.snapshots: allow
  restricted.virtual-machines.lowlevel: allow
description: Project Handson-lab-6
name: Handson-lab-6
used_by:
- /1.0/profiles/default?project=Handson-lab-6
- /1.0/profiles/external-net?project=Handson-lab-6
- /1.0/profiles/internal-net?project=Handson-lab-6
- /1.0/profiles/storage?project=Handson-lab-6

```

```

mac% incus storage list -f compact
      NAME      DRIVER      DESCRIPTION      USED BY      STATE
  basepool      zfs      passed to incus admin init      1      CREATED

```

shared-disks	zfs	shared disks accross Projects	0	CREATED
shared-images	zfs	shared images accross Projects	1	CREATED
shared-iso	zfs	shared ISOs accross Projects	1	CREATED
storage-1	zfs	Storage for Containers, VMs of Handson-lab-1	9	CREATED
storage-2	zfs	Storage for Containers, VMs of Handson-lab-2	5	CREATED
storage-3	zfs	Storage for Containers, VMs of Handson-lab-3	2	CREATED
storage-4	zfs	Storage for Containers, VMs of Handson-lab-4	2	CREATED
storage-5	zfs	Storage for Containers, VMs of Handson-lab-5	2	CREATED
storage-6	zfs	Storage for Containers, VMs of Handson-lab-6	1	CREATED

mac% incus network list -f compact

NAME	TYPE	MANAGED	IPV4	IPV6	DESCRIPTION	USED BY	STATE
external-net-1	bridge	YES	10.50.1.253/24	none	uplink for Handson-lab-1	5	CREATED
external-net-2	bridge	YES	10.50.2.253/24	none	uplink for Handson-lab-2	3	CREATED
external-net-3	bridge	YES	10.50.3.253/24	none	uplink for Handson-lab-3	2	CREATED
external-net-4	bridge	YES	10.50.4.253/24	none	uplink for Handson-lab-4	2	CREATED
external-net-5	bridge	YES	10.50.5.253/24	none	uplink for Handson-lab-5	2	CREATED
external-net-6	bridge	YES	10.50.6.253/24	none	uplink for Handson-lab-6	2	CREATED
incusbr0	bridge	YES	10.1.1.253/24	none	routes traffic via eno3	1	CREATED
internal-net-1	ovn	YES	172.16.10.253/24	none	routes via external-net-1	4	CREATED
internal-net-2	ovn	YES	172.16.10.253/24	none	routes via external-net-2	1	CREATED
internal-net-3	ovn	YES	172.16.10.253/24	none	routes via external-net-3	1	CREATED
internal-net-4	ovn	YES	172.16.10.253/24	none	routes via external-net-4	1	CREATED
internal-net-5	ovn	YES	172.16.10.253/24	none	routes via external-net-5	1	CREATED
internal-net-6	ovn	YES	172.16.10.253/24	none	routes via external-net-6	1	CREATED

mac % incus network acl list -f compact

NAME	DESCRIPTION	USED BY
isolate-lab-1	allow external-net-1 only. reject other external-nets	1
isolate-lab-2	allow external-net-2 only. reject other external-nets	1
isolate-lab-3	allow external-net-3 only. reject other external-nets	1
isolate-lab-4	allow external-net-4 only. reject other external-nets	1
isolate-lab-5	allow external-net-5 only. reject other external-nets	1
isolate-lab-6	allow external-net-6 only. reject other external-nets	1

With web UI

The screenshot shows the Incus UI web interface. The top navigation bar has 'Incus UI' and 'Instances' with a refresh icon. A red circle with the number '1' is on the project dropdown menu. The sidebar on the left shows a list of projects: Handson-lab-1, Handson-lab-2, Handson-lab-3, Handson-lab-4, Handson-lab-5, and Handson-lab-6. Handson-lab-1 has 3 instances, Handson-lab-2 has 1 instance, and the others have 0 instances. Below the projects is a '+ Create project' button. The main content area says 'No instances found' and 'There are no instances in this project. Spin up your first instance!'. It includes a 'How to create instances' link and a green 'Create instance' button. At the bottom, there's a footer with links for documentation, discussion, and reporting bugs, and a version number 'Version 6.12-ui-0.15'.

The screenshot shows the Incus UI web interface. The top navigation bar has 'Incus UI' and 'Pools' with a refresh icon. A red circle with the number '2' is on the 'Storage' item in the sidebar. The sidebar also lists: Instances, Profiles, Networking (Networks, ACLs), Storage (Pools, Volumes, Custom ISOs), Images, Configuration, Operations, Warnings, and Settings. Below the sidebar is a list of storage pools: basepool, shared-disks, shared-images, shared-iso, storage-1, storage-2, storage-3, storage-4, storage-5, and storage-6. Each pool entry includes its name, driver (zfs), size, volumes (this project) count, volumes (all projects) count, and status (Created). A green '+ Create pool' button is at the top right of the table. At the bottom, there's a footer with links for documentation, discussion, and reporting bugs, and a version number 'Version 6.12-ui-0.15'.

Incus UI

Networks

[See network map](#) [+ Create network](#)

NAME	TYPE	MANAGED	IPV4	IPV6	DESCRIPTION	FORWARDS	USED BY	STATE
external-net-1	bridge	Yes	10.50.1.253/24	none	uplink for Handson-lab-1	0	5	Created
external-net-2	bridge	Yes	10.50.2.253/24	none	uplink for Handson-lab-2	0	3	Created
external-net-3	bridge	Yes	10.50.3.253/24	none	uplink for Handson-lab-3	0	2	Created
external-net-4	bridge	Yes	10.50.4.253/24	none	uplink for Handson-lab-4	0	2	Created
external-net-5	bridge	Yes	10.50.5.253/24	none	uplink for Handson-lab-5	0	2	Created
external-net-6	bridge	Yes	10.50.6.253/24	none	uplink for Handson-lab-6	0	2	Created
incusbr0	bridge	Yes	10.1.1.253/24	none	routes traffic via eno3	1	1	Created
internal-net-1	ovn	Yes	172.16.10.253/24	none	routes via external-net-1	0	4	Created
internal-net-2	ovn	Yes	172.16.10.253/24	none	routes via external-net-2	0	1	Created
internal-net-3	ovn	Yes	172.16.10.253/24	none	routes via external-net-3	0	1	Created
internal-net-4	ovn	Yes	172.16.10.253/24	none	routes via external-net-4	0	1	Created
internal-net-5	ovn	Yes	172.16.10.253/24	none	routes via external-net-5	0	1	Created
internal-net-6	ovn	Yes	172.16.10.253/24	none	routes via external-net-6	0	1	Created

Version 6.12-ui-0.15

Incus UI

ACLs

[+ Create ACL](#)

NAME	DESCRIPTION	USED BY
isolate-lab-2	allow external-net-2 only. reject other external-nets	1
isolate-lab-1	allow external-net-1 only. reject other external-nets	1
isolate-lab-3	allow external-net-3 only. reject other external-nets	1
isolate-lab-4	allow external-net-4 only. reject other external-nets	1
isolate-lab-5	allow external-net-5 only. reject other external-nets	1
isolate-lab-6	allow external-net-6 only. reject other external-nets	1

Version 6.12-ui-0.15

Credits and Thanks to my incus Gurus

Stephane Graber, Zabbly

<https://www.youtube.com/@TheZabbly/videos>

Trevor Sullivan

<https://www.youtube.com/@TrevorSullivan/videos>

Simos Xenitellis

<https://blog.simos.info/>

Scott at Scotti-BYTE

<https://www.youtube.com/@scottibyte/videos>

Incus community

<https://discuss.linuxcontainers.org>

Deepseek.com

<https://www.deepseek.com/en>

About the Author



Coordinates:

Name : Ananda Kammampati

Contact : ananda at fieldday dot io

Profile : <https://www.linkedin.com/in/anandakammampati>

Humble Brag:

1. Founder www.fieldday.io
2. Creator of StudentOS: <http://fieldday.io/StudentOS>
3. Conceptualized and designed 'Hands-on-Box' , a unique portable server appliance for the needs of field and training (<http://www.fieldday.io/hands-on-box/>)
4. Published articles:
<https://www.linkedin.com/in/anandakammampati/recent-activity/articles/>
5. More articles in fieldday.io website
6. Repository:
<https://github.com/fieldday-io/books>



- Education:

Bachelors in Computer Technology, Government College of Technology, Coimbatore, TN, India

- Many hats:

Technical Product Manager, Senior Solutions Architect, BizDev Manager, Technical Lead, R & D Engineer, Computer Operator Trainee

- Learned at:

1. VMWare (CA, USA)
2. Sun Microsystems (CA, USA)
3. HCL Consulting (CA, USA)
4. BFL Software Ltd (Bengaluru)
5. Tata Elxsi India Ltd (Bengaluru)
6. Supercomputer Education & Research Center (IISc, Bengaluru)